

Boosting Symbolic Execution for Vulnerability Detection

Dissertation Defense

Haoxin Tu

Committee Members: Prof. Lingxiao Jiang, Prof. Xuhua Ding, Prof. David Lo, Prof. Marcel Böhme

May 9th, 2025

Outline

Outline

□ Background

- Software vulnerabilities
- Symbolic execution

Outline

□ Background

- Software vulnerabilities
- Symbolic execution

□ Motivation

- What limitations prevent current symbolic execution for vulnerability detection?
- Thesis statement and research objectives

Outline

□ Background

- Software vulnerabilities
- Symbolic execution

□ Motivation

- What limitations prevent current symbolic execution for vulnerability detection?
- Thesis statement and research objectives

□ Methodology

- **SymLoc**: A new memory model for symbolic execution
- **FastKLEE and Vital**: Two new path exploration for symbolic execution
- **Cottontail**: A new structured input generation for symbolic execution

Outline

□ Background

- Software vulnerabilities
- Symbolic execution

□ Motivation

- What limitations prevent current symbolic execution for vulnerability detection?
- Thesis statement and research objectives

□ Methodology

- **SymLoc**: A new memory model for symbolic execution
- **FastKLEE and Vital**: Two new path exploration for symbolic execution
- **Cottontail**: A new structured input generation for symbolic execution

□ Future work

- Extension of current solutions
- Combine program analysis with GenAI

Outline

□ Background

- Software vulnerabilities
- Symbolic execution

□ Motivation

- What limitations prevent current symbolic execution for vulnerability detection?
- Thesis statement and research objectives

□ Methodology

- **SymLoc**: A new memory model for symbolic execution
- **FastKLEE and Vital**: Two new path exploration for symbolic execution
- **Cottontail**: A new structured input generation for symbolic execution

□ Future work

- Extension of current solutions
- Combine program analysis with GenAI

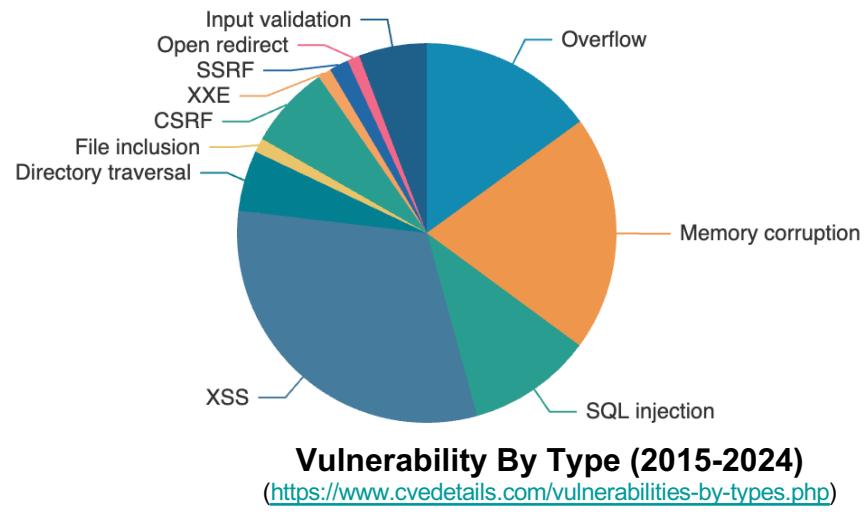
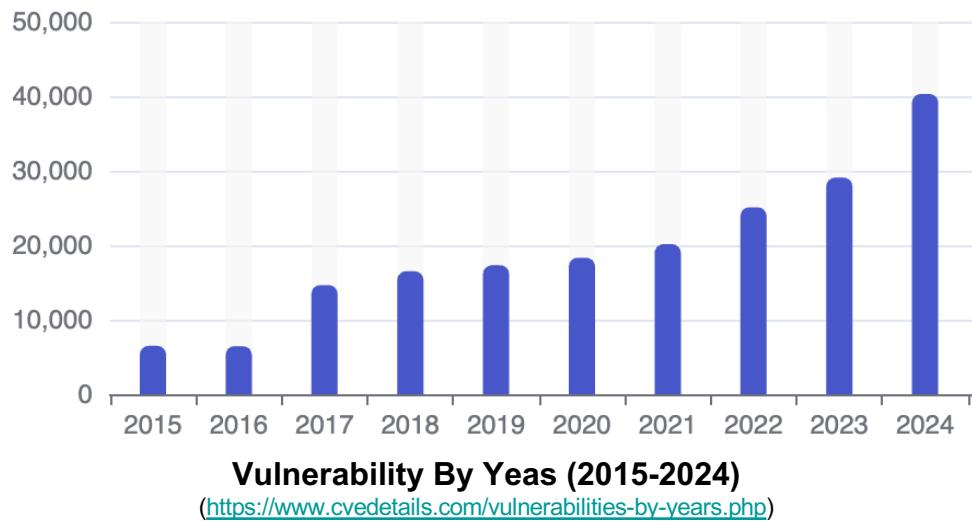
□ Conclusion and Acknowledgment

Background: software are everywhere

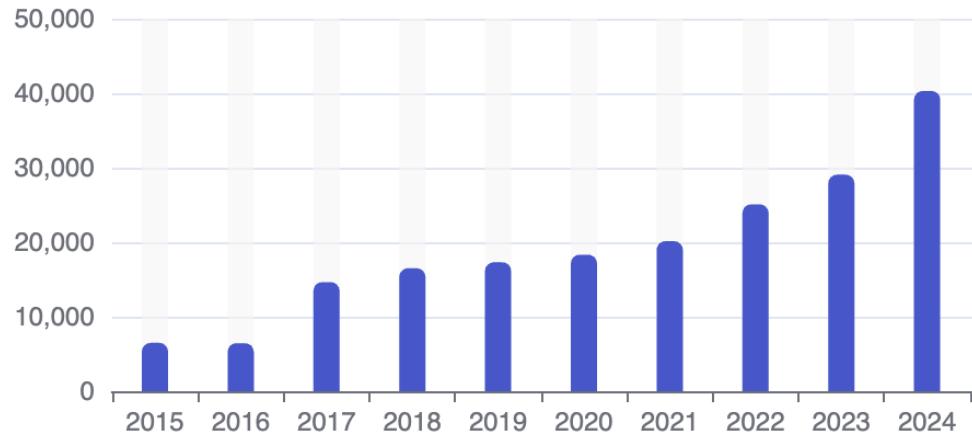


<https://www.coderus.com/software-101-a-complete-guide-to-the-different-types-of-software/>

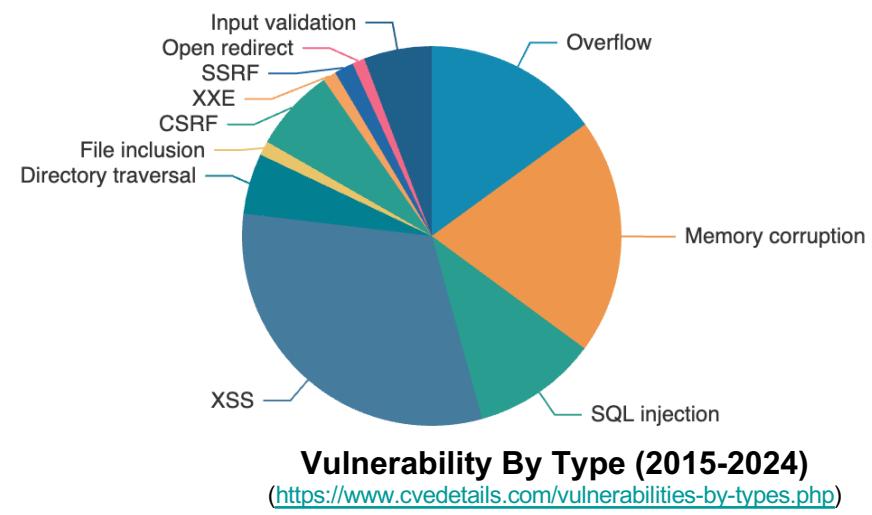
Background: inevitable software defects



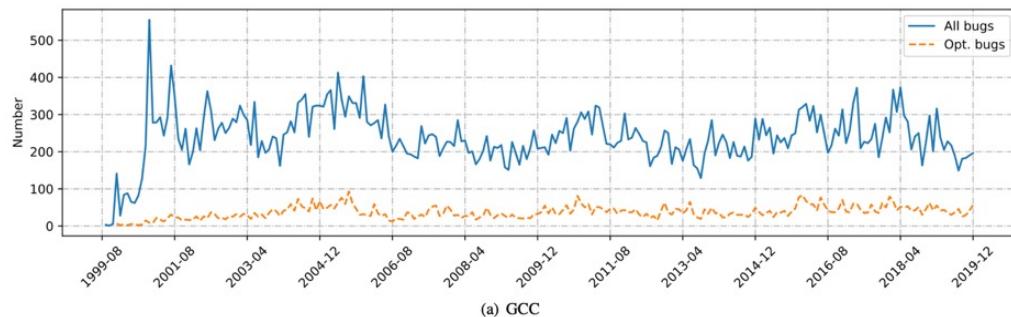
Background: inevitable software defects



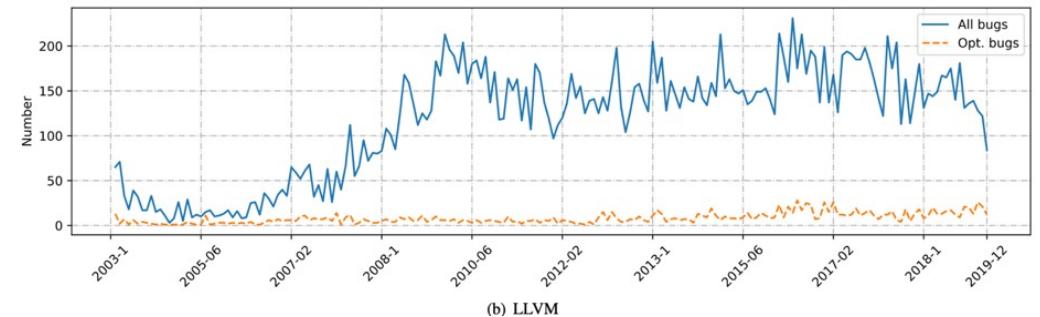
Vulnerability By Years (2015-2024)
(<https://www.cvedetails.com/vulnerabilities-by-years.php>)



Vulnerability By Type (2015-2024)
(<https://www.cvedetails.com/vulnerabilities-by-types.php>)



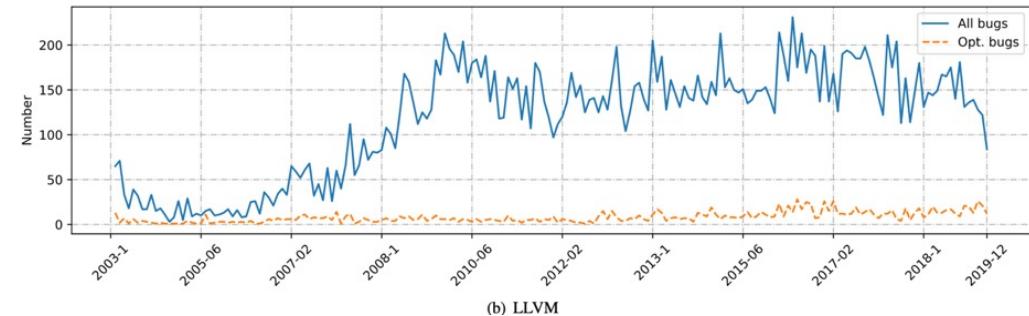
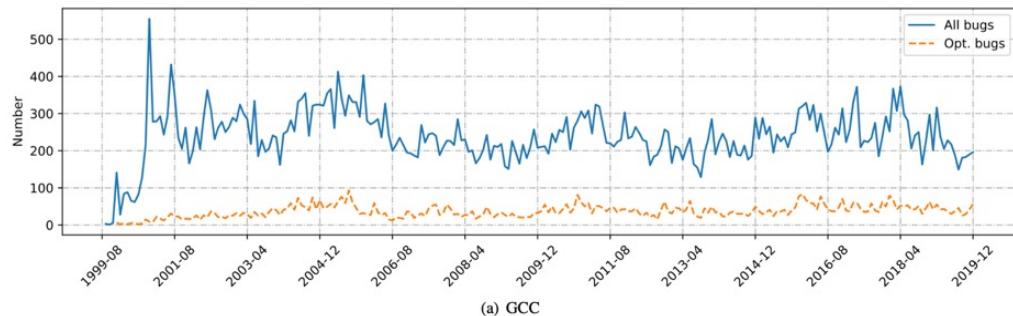
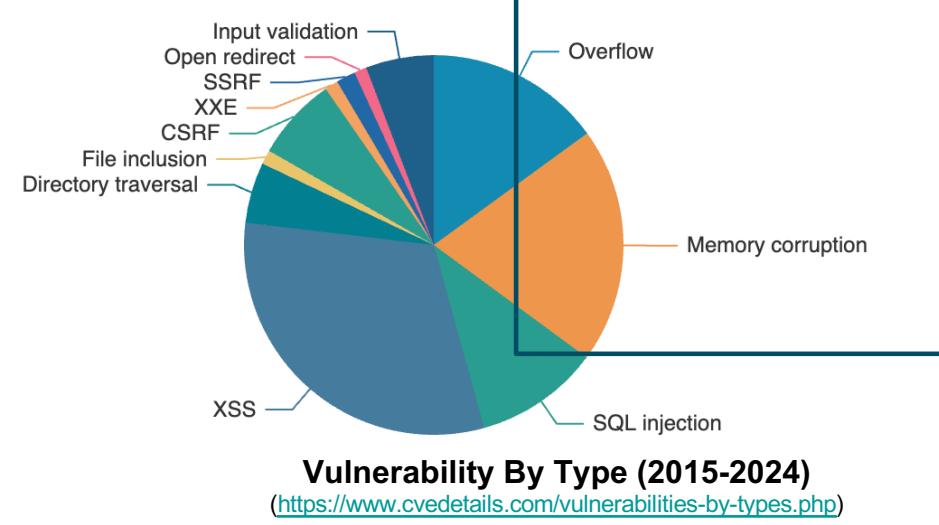
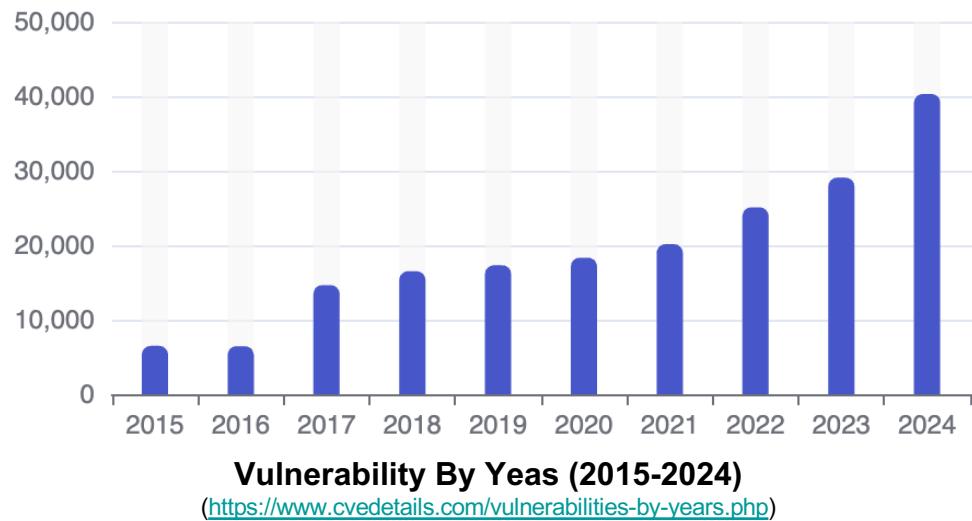
(a) GCC



(b) LLVM

Bug trends on two mainstream compilers (GCC and LLVM)

Background: inevitable software defects



Bug trends on two mainstream compilers (GCC and LLVM)

Background: symbolic execution (1/3)

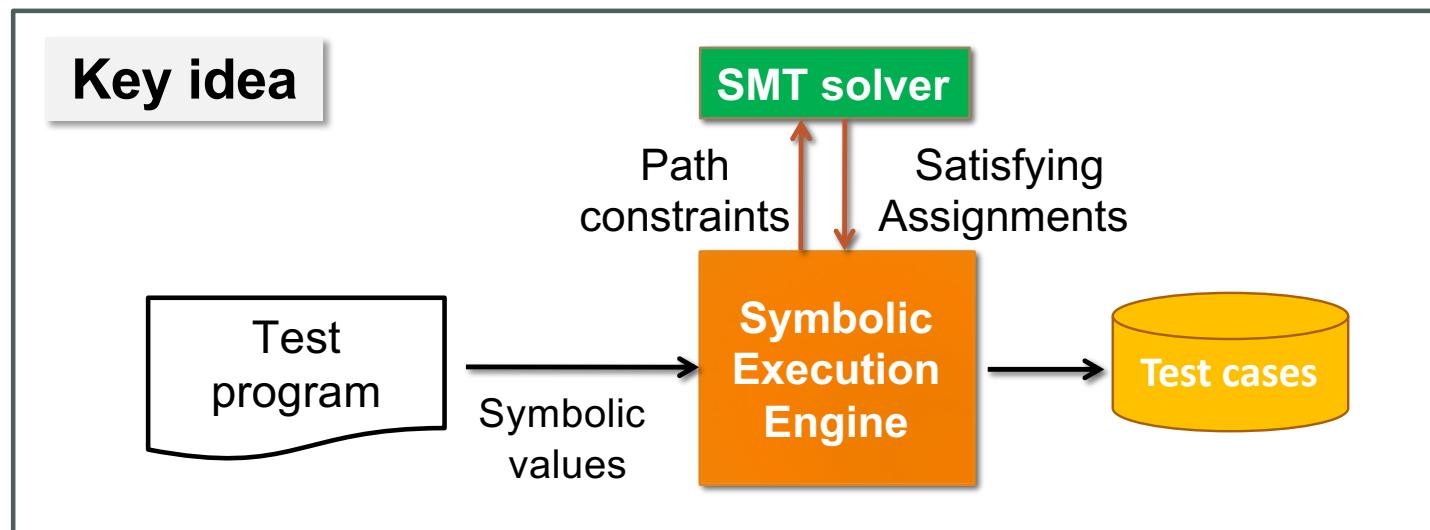
□ What is symbolic execution?

- Proposed in 1976 [1], one of the most popular program analysis techniques, which scales for many **software testing** and **computer security** applications

Background: symbolic execution (1/3)

□ What is symbolic execution?

- Proposed in 1976 [1], one of the most popular program analysis techniques, which scales for many **software testing** and **computer security** applications



Background: symbolic execution (2/3)

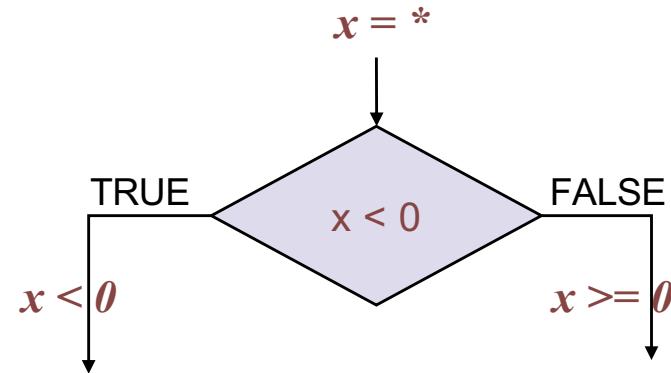
□ A toy example

```
int bad_abs(int x)
{
    if (x < 0)
        return -x;
    if (x == 1234)
        return -x;
    return x;
}
```

Background: symbolic execution (2/3)

□ A toy example

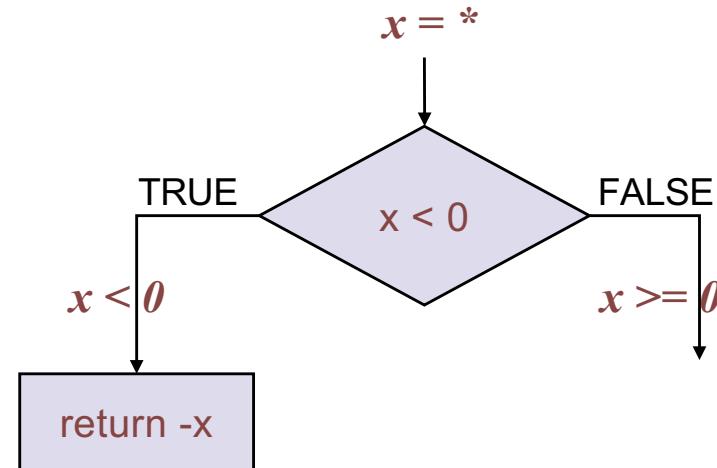
```
int bad_abs(int x)
{
    if (x < 0)
        return -x;
    if (x == 1234)
        return -x;
    return x;
}
```



Background: symbolic execution (2/3)

□ A toy example

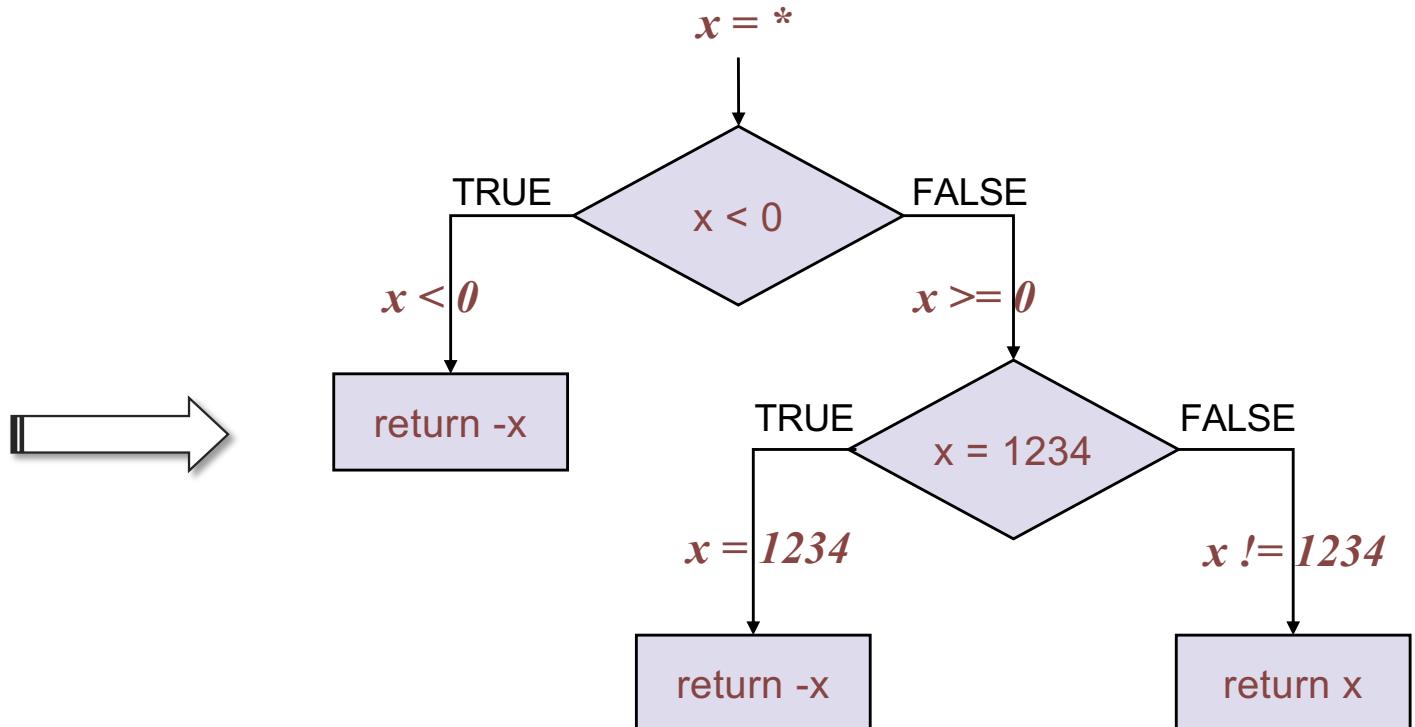
```
int bad_abs(int x)
{
    if (x < 0)
        return -x;
    if (x == 1234)
        return -x;
    return x;
}
```



Background: symbolic execution (2/3)

□ A toy example

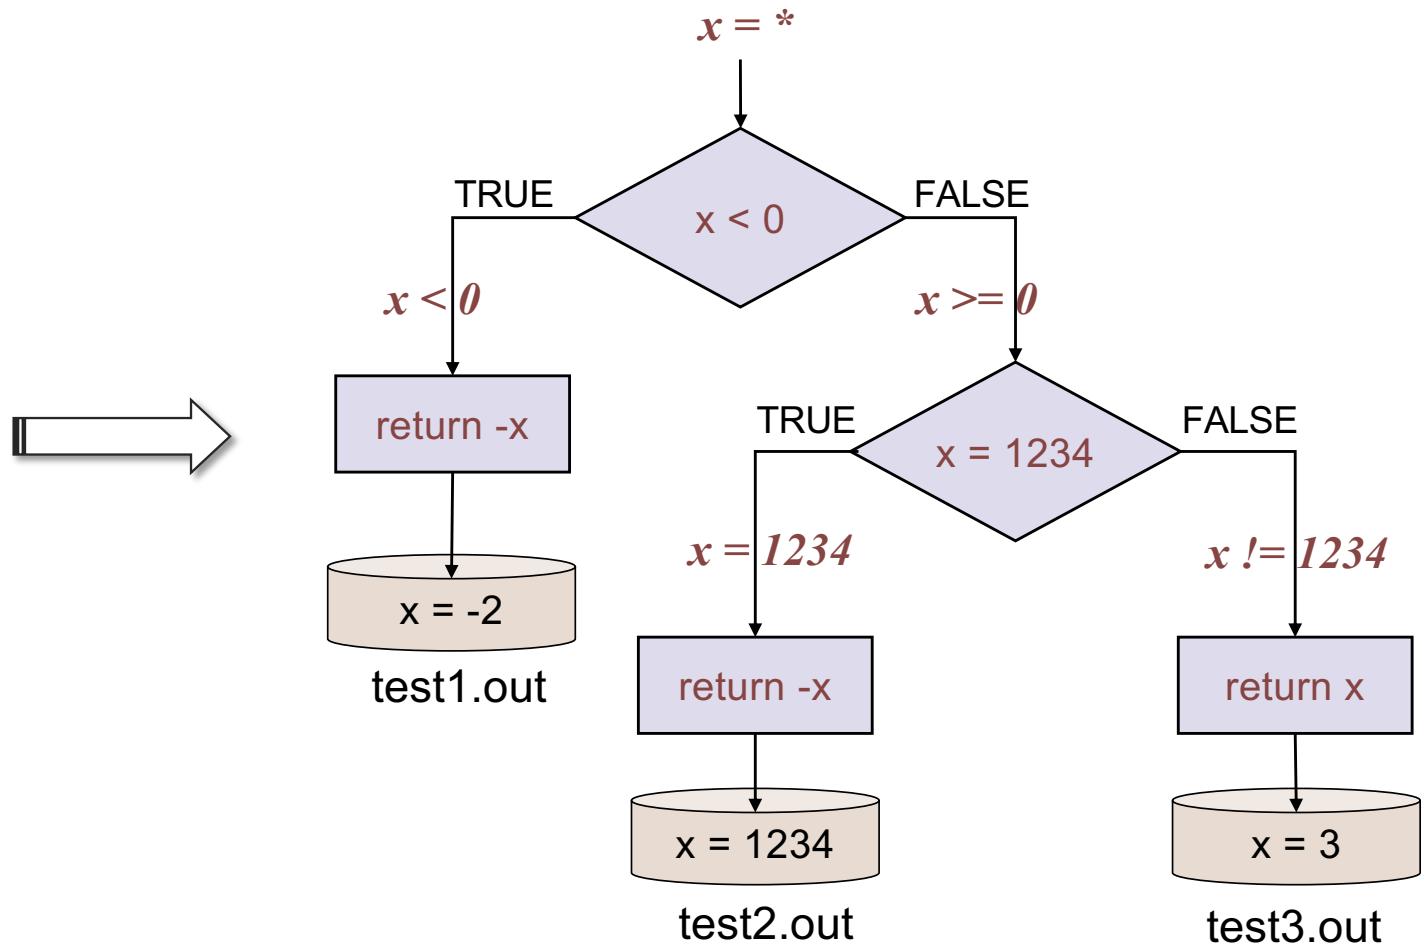
```
int bad_abs(int x)
{
    if (x < 0)
        return -x;
    if (x == 1234)
        return -x;
    return x;
}
```



Background: symbolic execution (2/3)

□ A toy example

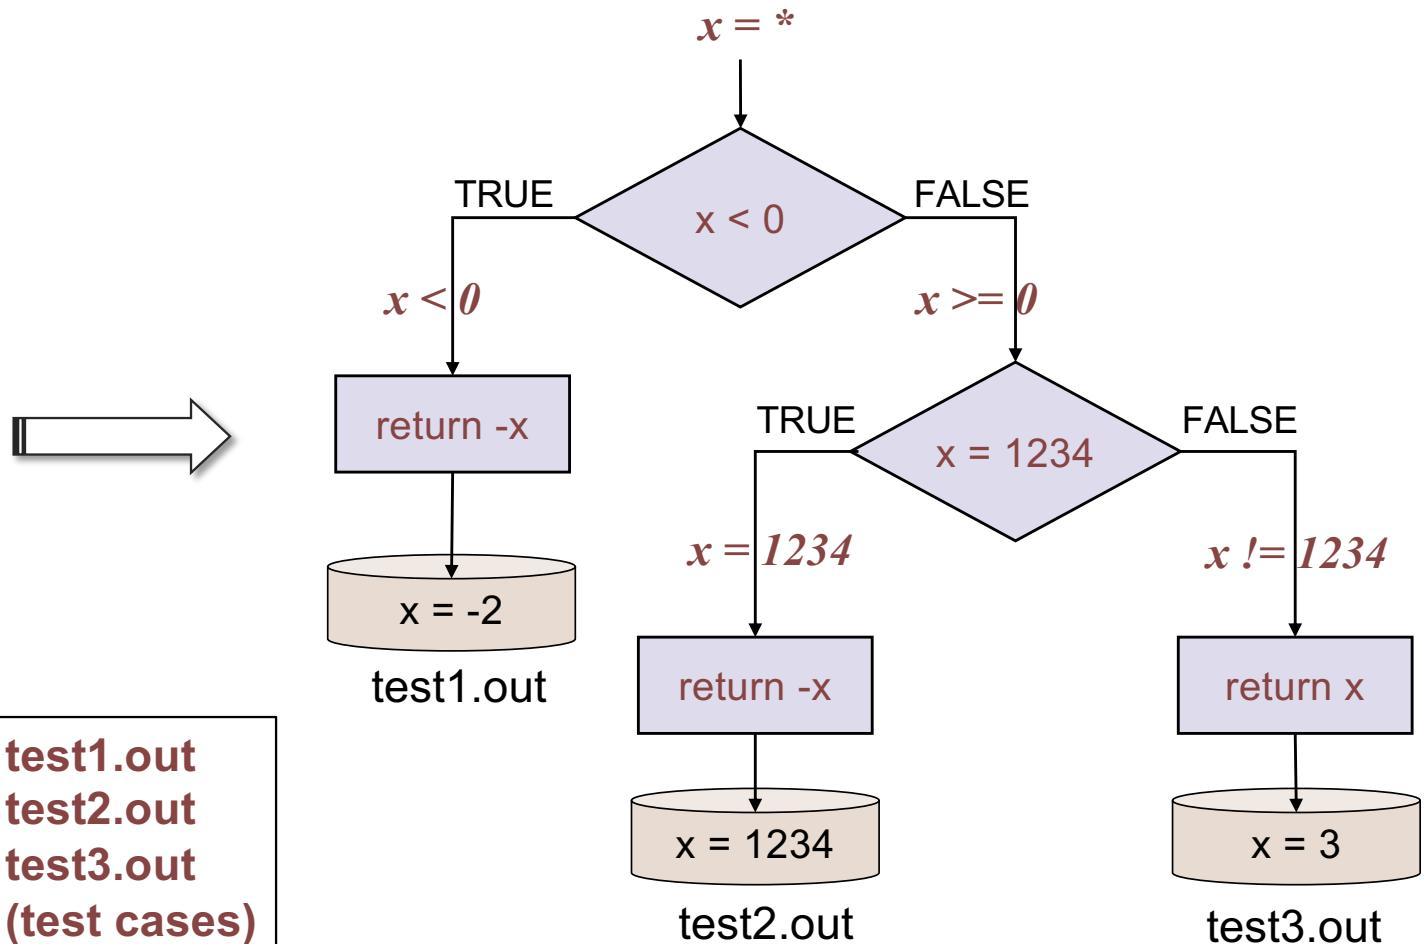
```
int bad_abs(int x)
{
    if (x < 0)
        return -x;
    if (x == 1234)
        return -x;
    return x;
}
```



Background: symbolic execution (2/3)

□ A toy example

```
int bad_abs(int x)
{
    if (x < 0)
        return -x;
    if (x == 1234)
        return -x;
    return x;
}
```



$x < 0;$
 $x \geq 0 \&& x = 1234;$
 $x \geq 0 \&& x \neq 1234;$
(path constraints)

`test1.out`
`test2.out`
`test3.out`
(test cases)

Background: symbolic execution (3/3)



School of
Computing and
Information Systems

Background: symbolic execution (3/3)

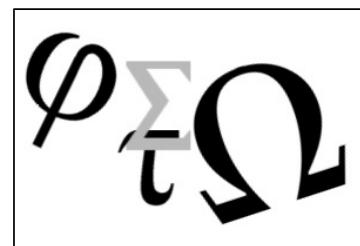
□ Why symbolic execution could work?

- Execute the program with symbolic inputs
- Represent **equivalent execution paths** with **path constraints**
- **Solve path constraints** to obtain one representative input that exercises the program to go down that specific path

Background: symbolic execution (3/3)

□ Why symbolic execution could work?

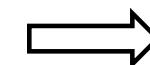
- Execute the program with symbolic inputs
- Represent **equivalent execution paths** with **path constraints**
- **Solve path constraints** to obtain one representative input that exercises the program to go down that specific path



Path constraints



Constraint Solver

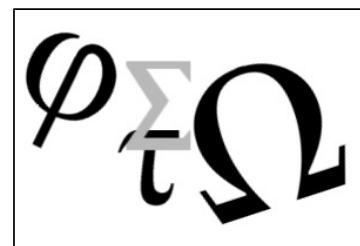


Worked!

Background: symbolic execution (3/3)

□ Why symbolic execution could work?

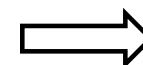
- Execute the program with symbolic inputs
- Represent **equivalent execution paths** with **path constraints**
- **Solve path constraints** to obtain one representative input that exercises the program to go down that specific path



Path constraints



Constraint Solver



Worked!

- Many applications
 - high-coverage test generation, automated debugging, automated program repair, exploit generation, wireless sensor networks, online gaming, ...

Background: symbolic execution (3/3)



School of
Computing and
Information Systems

□ Why symbolic execution could work?

- Execute the program with symbolic inputs

*Is current symbolic execution effective enough
to detect vulnerabilities in practice?*

Path constraints

- Many applications
 - high-coverage test generation, automated debugging, automated program repair, exploit generation, wireless sensor networks, online gaming, ...

Constraint Solver

Background: symbolic execution (3/3)



School of
Computing and
Information Systems

□ Why symbolic execution could work?

- Execute the program with symbolic inputs

*Is current symbolic execution effective enough
to detect vulnerabilities in practice?*

NO, and why?

Path constraints

- Many applications
 - high-coverage test generation, automated debugging, automated program repair, exploit generation, wireless sensor networks, online gaming, ...

Constraint Solver

Outline

□ Background

- Software vulnerabilities
- Symbolic execution

□ Motivation

- What limitations prevent current symbolic execution for vulnerability detection?
- Thesis statement and research objectives

□ Methodology

- **SymLoc**: A new memory model for symbolic execution
- **FastKLEE and Vital**: Two new path exploration for symbolic execution
- **Cottontail**: A new structured input generation for symbolic execution

□ Future work

- Extension of current solutions
- Combine program analysis with GenAI

□ Conclusion and Acknowledgment

Motivation Example

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```

Motivation Example

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```

Motivation Example

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```

Motivation Example

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```

Motivation Example

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```

Can existing symbolic execution effectively detect the three vulnerabilities?

Motivation Example

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```

Can existing symbolic execution effectively detect the three vulnerabilities?

NO

Limitation 1: Limited memory modeling



School of
Computing and
Information Systems

- How does the engine handle dynamic memory allocations?

Limitation 1: Limited memory modeling

□ How does the engine handle dynamic memory allocations?

1. int* array = malloc (100 * sizeof (int));
2. array [i] = 10; *// i is symbolic*
3. assert(array[j] != 0); *// j is symbolic*

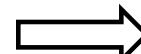
Limitation 1: Limited memory modeling

□ How does the engine handle dynamic memory allocations?

1. `int* array = malloc (100 * sizeof (int));`
2. `array [i] = 10; // i is symbolic`
3. `assert(array[j] != 0); // j is symbolic`

➤ Possible solutions

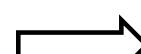
1. Fully symbolic (ASE'17)
 - consider **any possible** outcome
2. Fully concrete (OSDI'08)
 - consider **one possible** outcome
3. Partial symbolic and concrete (S&P'12)
 - concretize writes,
 - Possible ranges to read



N states
accurate but not scale



1 state
scale but not accurate



K states
scale but (in) accurate

Limitation 1: Limited memory modeling

□ How does the engine handle dynamic memory allocations?

1. `int* array = malloc (100 * sizeof (int));`
2. `array [i] = 10; // i is symbolic`
3. `assert(array[j] != 0); // j is symbolic`

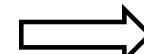


Fixed addresses

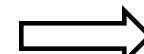
Problematic read/write

➤ Possible solutions

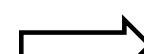
1. Fully symbolic (ASE'17)
 - consider **any possible** outcome
2. Fully concrete (OSDI'08)
 - consider **one possible** outcome
3. Partial symbolic and concrete (S&P'12)
 - concretize writes,
 - Possible ranges to read



N states
accurate but not scale



1 state
scale but not accurate



K states
scale but (in) accurate

Limitation 1: Limited memory modeling

□ How does the engine handle dynamic memory allocations?

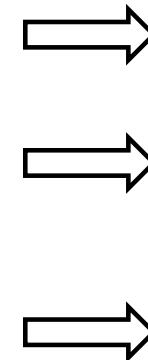
1. `int* array = malloc (100 * sizeof (int));`
2. `array [i] = 10; // i is symbolic`
3. `assert(array[j] != 0); // j is symbolic`



Fixed addresses
Problematic read/write

➤ Possible solutions

1. Fully symbolic (ASE'17)
 - consider **any possible** outcome
2. Fully concrete (OSDI'08)
 - consider **one possible** outcome
3. Partial symbolic and concrete (S&P'12)
 - concretize writes,
 - Possible ranges to read



N states
accurate but not scale

1 state
scale but not accurate

K states
scale but (in) accurate

We need a more complete and precise memory modeling !

Motivation Example

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```



Motivation Example

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```

Can existing symbolic execution effectively detect the three vulnerabilities?

Motivation Example

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```



Can existing symbolic execution effectively detect the three vulnerabilities?

NO, two remaining

Limitation 2: path explosion



School of
Computing and
Information Systems

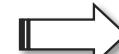
- How does symbolic execution deal with path explosion?

Limitation 2: path explosion

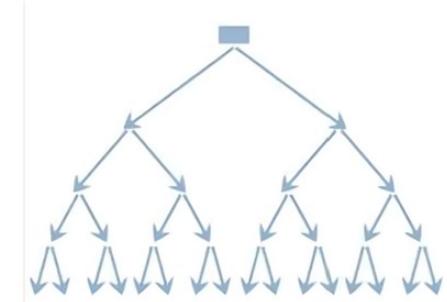
□ How does symbolic execution deal with path explosion?

```
void process(char input[3]) {  
    int counter = 0;  
    if (input[0] == 'a') counter++;  
    if (input[1] == 'b') counter++;  
    if (input[2] == 'c') counter++;  
    if (counter >= 3) success();  
    error();  
}
```

- Exponentially many execution paths



4 conditional nodes



16 (2^4) execution paths

Limitation 2: path explosion

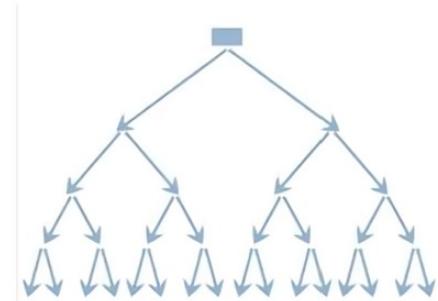
□ How does symbolic execution deal with path explosion?

```
void process(char input[3]) {  
    int counter = 0;  
    if (input[0] == 'a') counter++;  
    if (input[1] == 'b') counter++;  
    if (input[2] == 'c') counter++;  
    if (counter >= 3) success();  
    error();  
}
```

- Exponentially many execution paths



4 conditional nodes



16 (2^4) execution paths

➤ Efficient solutions

- With specialized optimizations (HotOS'13)
- With code transformation (ECOOP'18)

Limitation 2: path explosion

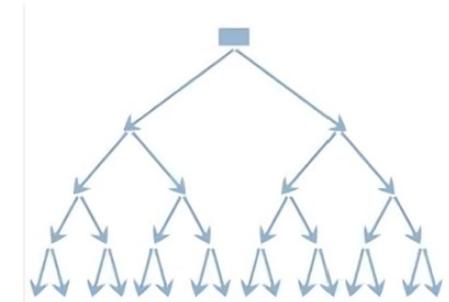
□ How does symbolic execution deal with path explosion?

```
void process(char input[3]) {  
    int counter = 0;  
    if (input[0] == 'a') counter++;  
    if (input[1] == 'b') counter++;  
    if (input[2] == 'c') counter++;  
    if (counter >= 3) success();  
    error();  
}
```

- Exponentially many execution paths



4 conditional nodes



16 (2^4) execution paths

➤ Efficient solutions

- With specialized optimizations (HotOS'13)
- With code transformation (ECOOP'18)

➤ Effective solutions

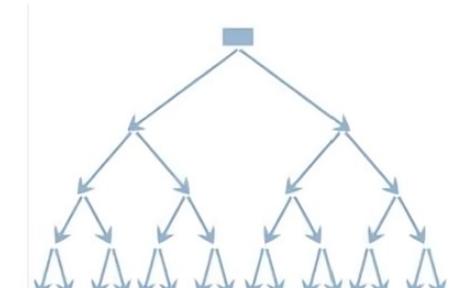
- Random search (DFS and BFS)
- Heuristic guided search (OSDI'08)
- coverage, instruction, etc.

Limitation 2: path explosion

□ How does symbolic execution deal with path explosion?

```
void process(char input[3]) {  
    int counter = 0;  
    if (input[0] == 'a') counter++;  
    if (input[1] == 'b') counter++;  
    if (input[2] == 'c') counter++;  
    if (counter >= 3) success();  
    error();  
}
```

- Exponentially many execution paths



➤ Efficient solutions

- With specialized optimizations (HotOS'13)
- With code transformation (ECOOP'18)

➤ Effective solutions

- Random search (DFS and BFS)
- Heuristic guided search (OSDI'08)
 - coverage, instruction, etc.



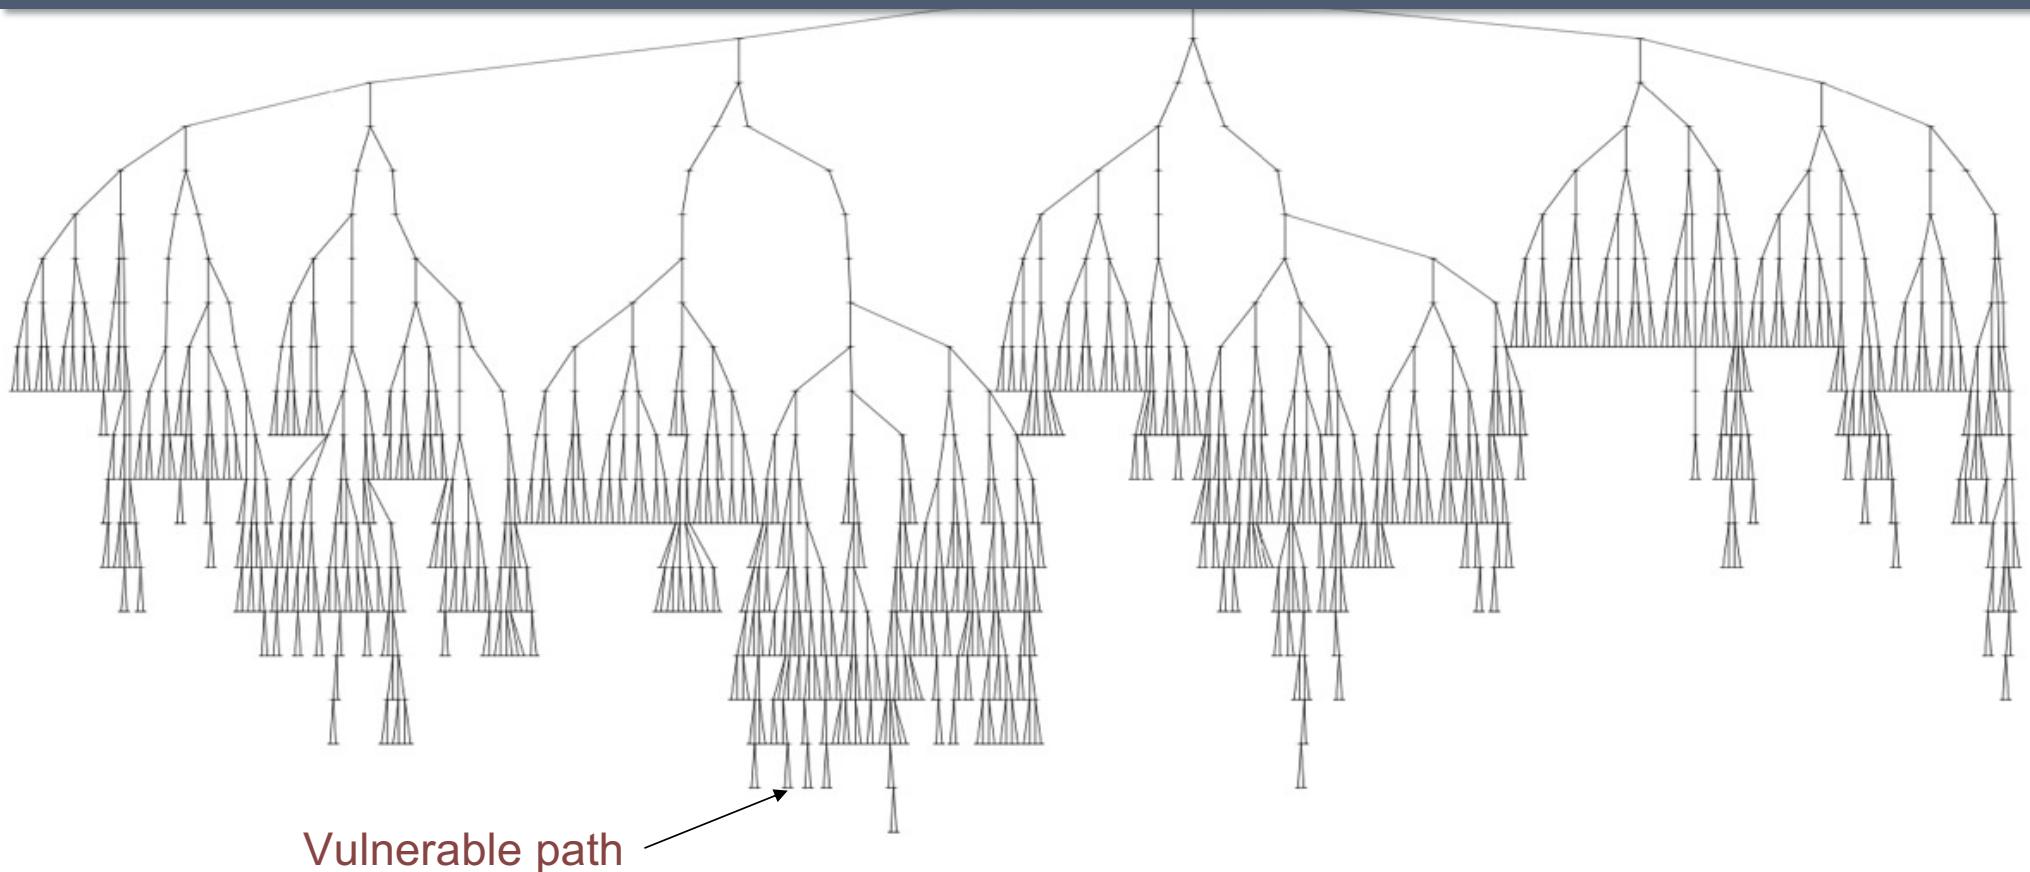
Inefficient path search

No unified targets

Ineffective path search

Path explosion visualization

How to search vulnerable paths in rather complex software systems?



Path explosion visualization

How to search vulnerable paths in rather complex software systems?

We need a more efficient and effective path exploration !

Vulnerable path



Motivation Example

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```

Motivation Example

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```

Can existing symbolic execution effectively detect the three vulnerabilities?

Motivation Example

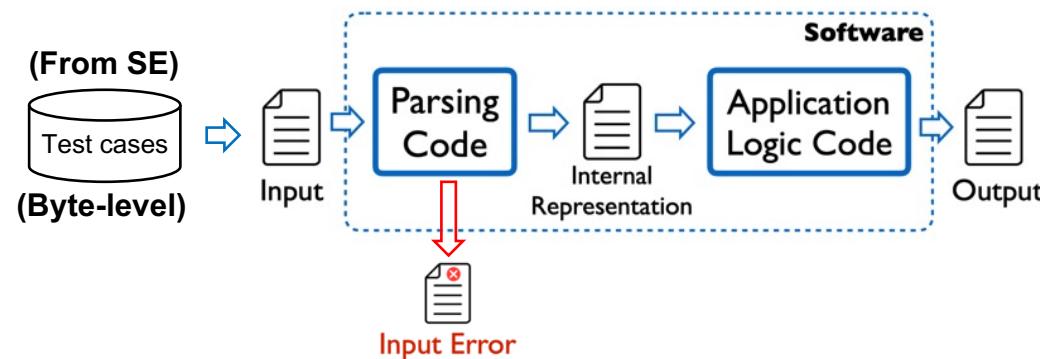
```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```

Can existing symbolic execution effectively detect the three vulnerabilities?

NO, one remaining still

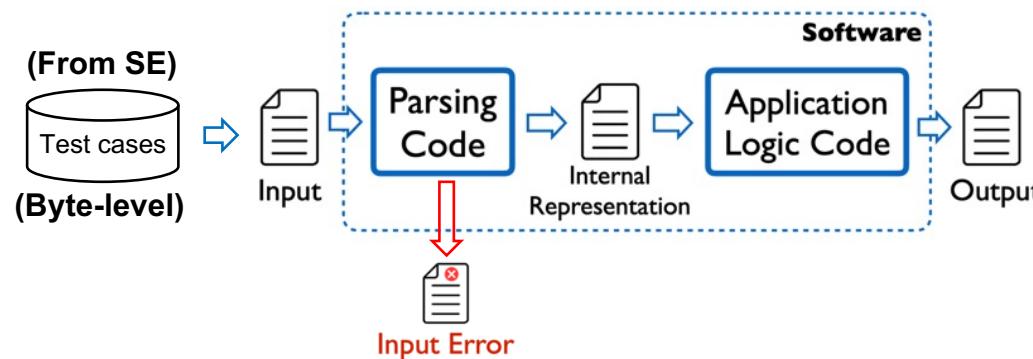
Limitation 3: structured input generation

□ How does symbolic execution generate structured test cases?



Limitation 3: structured input generation

□ How does symbolic execution generate structured test cases?

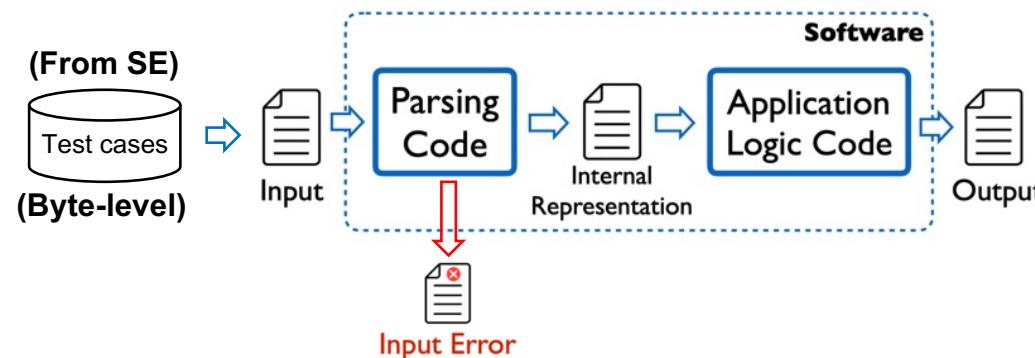


➤ Possible solutions

- Grammar-based generation (PLDI'08)
 - Use grammar specifications to guide generation
- Token-level symbolization (ISSTA'21)
 - Treating tokens rather than bytes as symbolic data

Limitation 3: structured input generation

□ How does symbolic execution generate structured test cases?



➤ Possible solutions

- Grammar-based generation (PLDI'08)
 - Use grammar specifications to guide generation
- Token-level symbolization (ISSTA'21)
 - Treating tokens rather than bytes as symbolic data

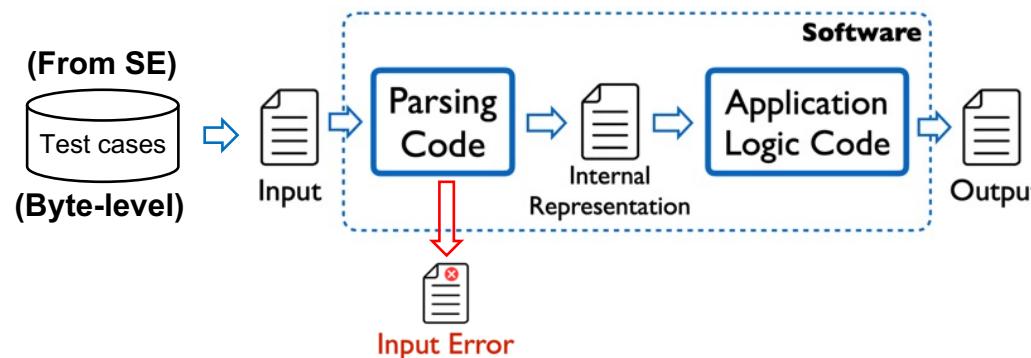


Structure-agnostic
path selection

Ineffective
constraint solving

Limitation 3: structured input generation

□ How does symbolic execution generate structured test cases?



➤ Possible solutions

- Grammar-based generation (PLDI'08)
 - Use grammar specifications to guide generation
- Token-level symbolization (ISSTA'21)
 - Treating tokens rather than bytes as symbolic data



Structure-agnostic
path selection

Ineffective
constraint solving

We need a new structured test case generation strategy !

Thesis statement and research objectives



School of
Computing and
Information Systems

- **Statement:** This thesis aims to **boost symbolic execution** by designing new solutions to alleviate three key limitations in **memory modeling/path exploration/test input generation**, **for efficient and effective automatic vulnerability detection**.

Summary of Key Research Problems

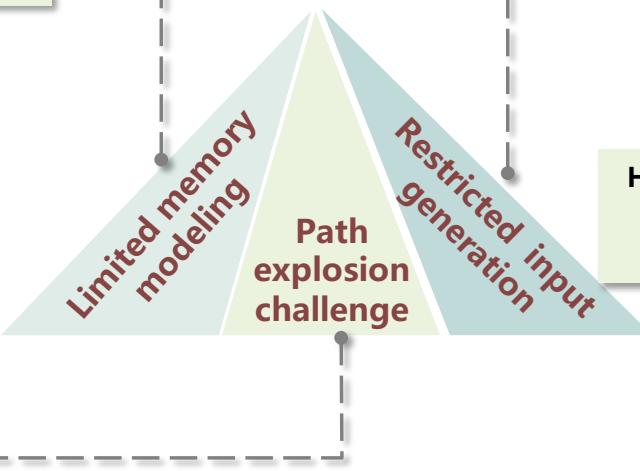
How can we support a more complete memory model
for dynamic memory allocations ?



Fixed addresses



Problematic W/R

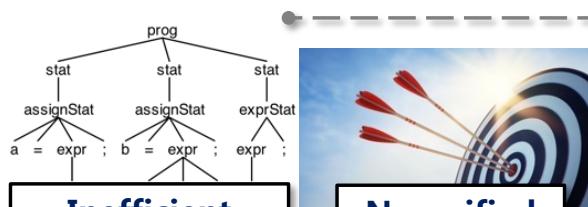


Structure-agnostic
path selection



Ineffective
constraint solving

How can we perform concolic execution to generate
highly structured test inputs for systematically
testing parsing programs?



Inefficient
path search



No unified
targets

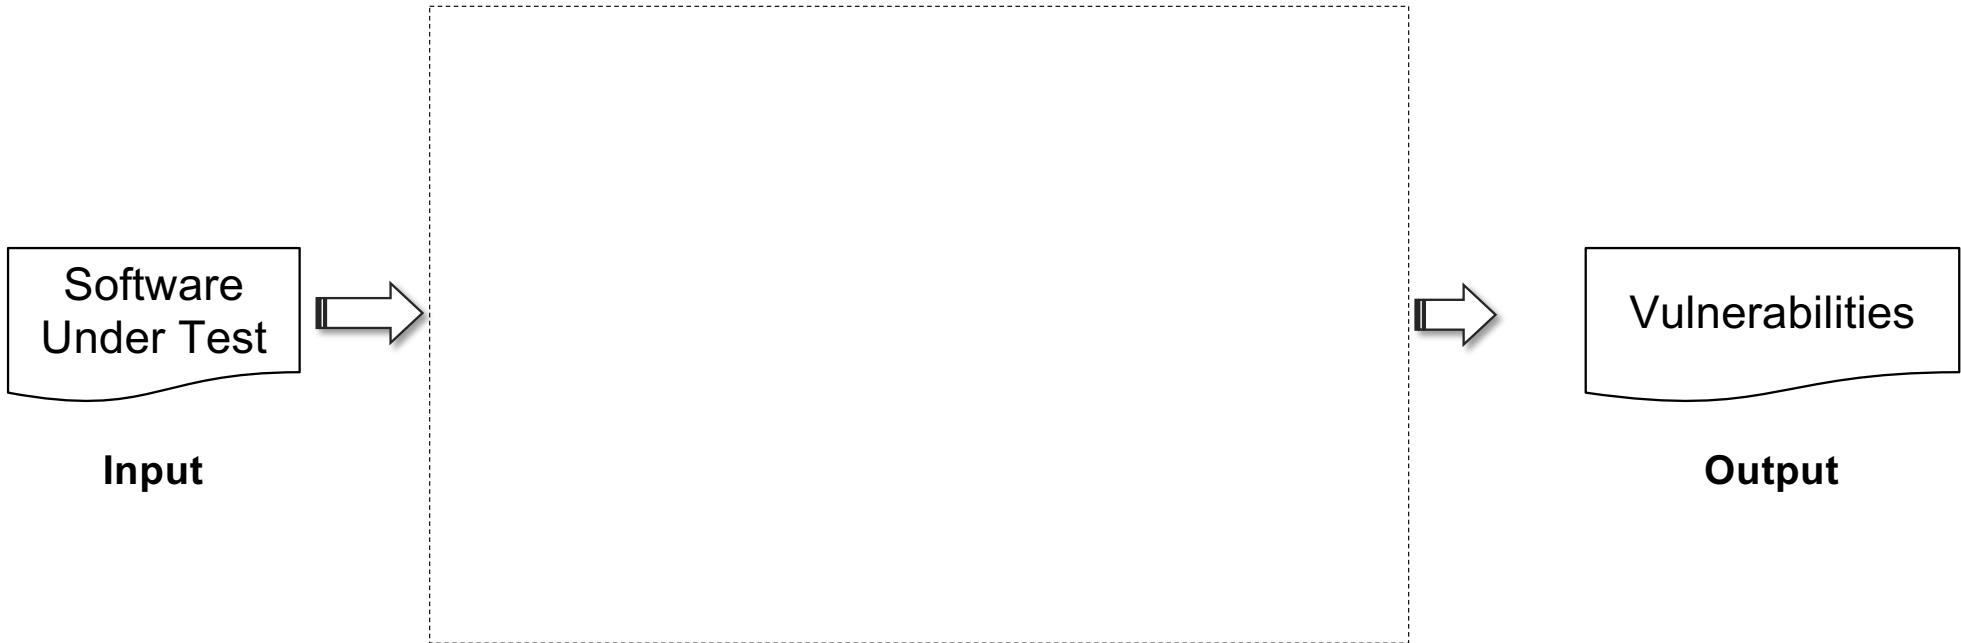


Ineffective path
search

How can we efficiently and effectively search vulnerable paths in
rather complex software systems?

Thesis statement and research objectives

- **Statement:** This thesis aims to **boost symbolic execution** by designing new solutions to alleviate three key limitations in **memory modeling/path exploration/test input generation**, **for efficient and effective automatic vulnerability detection.**



[1] Haixin Tu, Lingxiao Jiang, and et.al, "Concretely Mapped Symbolic Memory Locations for Memory Error Detection" (IEEE TSE).

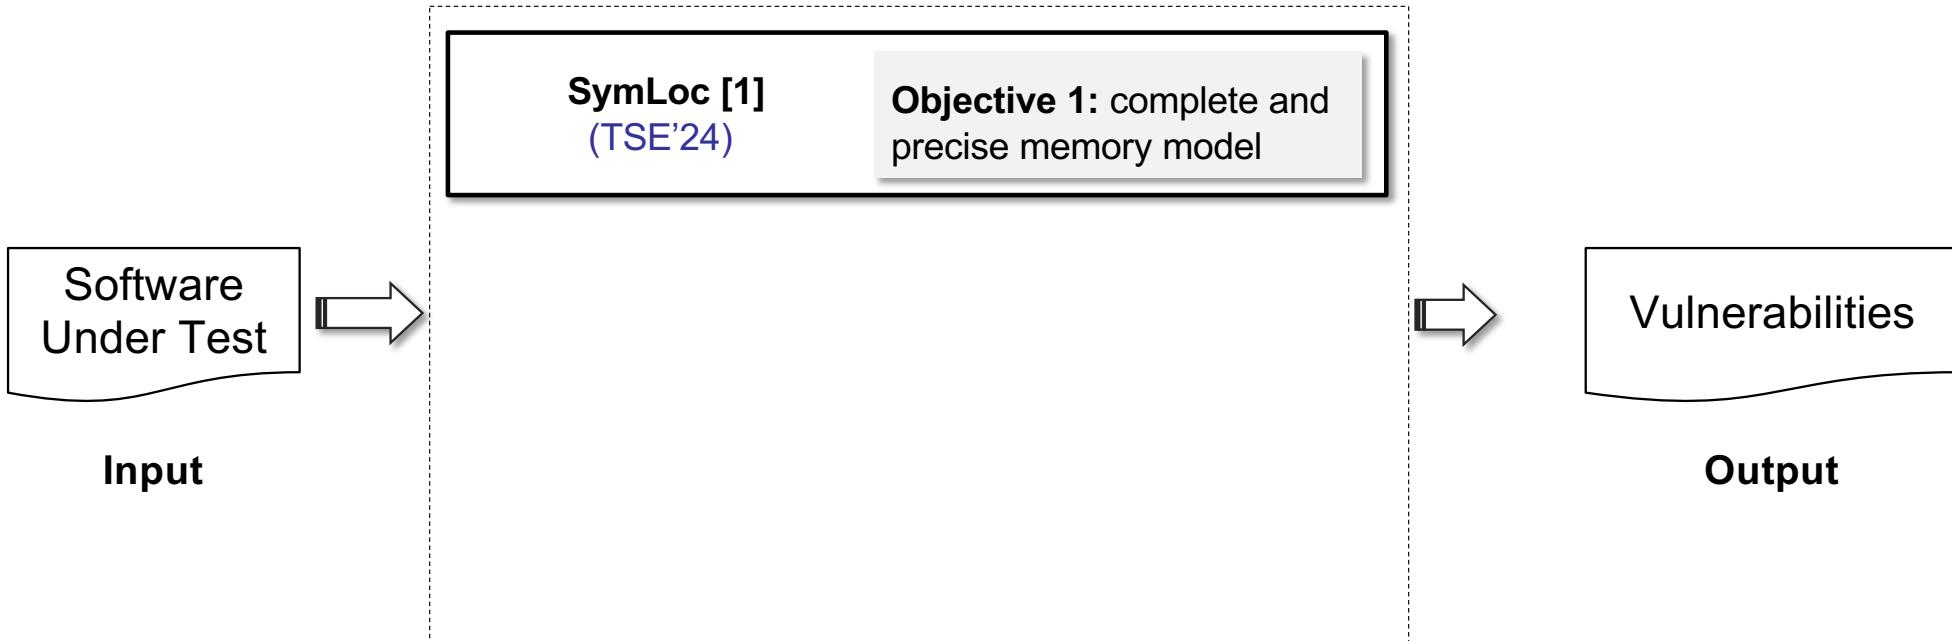
[2] Haixin Tu, Lingxiao Jiang, and et.al., "FastKLEE: Faster Symbolic Execution via Redundant Bound Checking of Type-Safe Pointers", in Tool Demonstrations Track of FSE 2022.

[3] Haixin Tu, Lingxiao Jiang and et.al., "Vital: Vulnerability-Oriented Symbolic Execution via Type-Unsafe Pointer-Guided Monte Carlo Tree Search." arXiv:2408.08772 (2024).

[4] Haixin Tu, Seongmin Lee, and et.al. "Large Language Model-Driven Concolic Execution for Highly Structured Test Input Generation." arXiv:2504.17542 (2025).

Thesis statement and research objectives

- **Statement:** This thesis aims to **boost symbolic execution** by designing new solutions to alleviate three key limitations in **memory modeling/path exploration/test input generation**, for efficient and effective automatic **vulnerability detection**.



[1] Haixin Tu, Lingxiao Jiang, and et.al. "Concretely Mapped Symbolic Memory Locations for Memory Error Detection" (IEEE TSE).

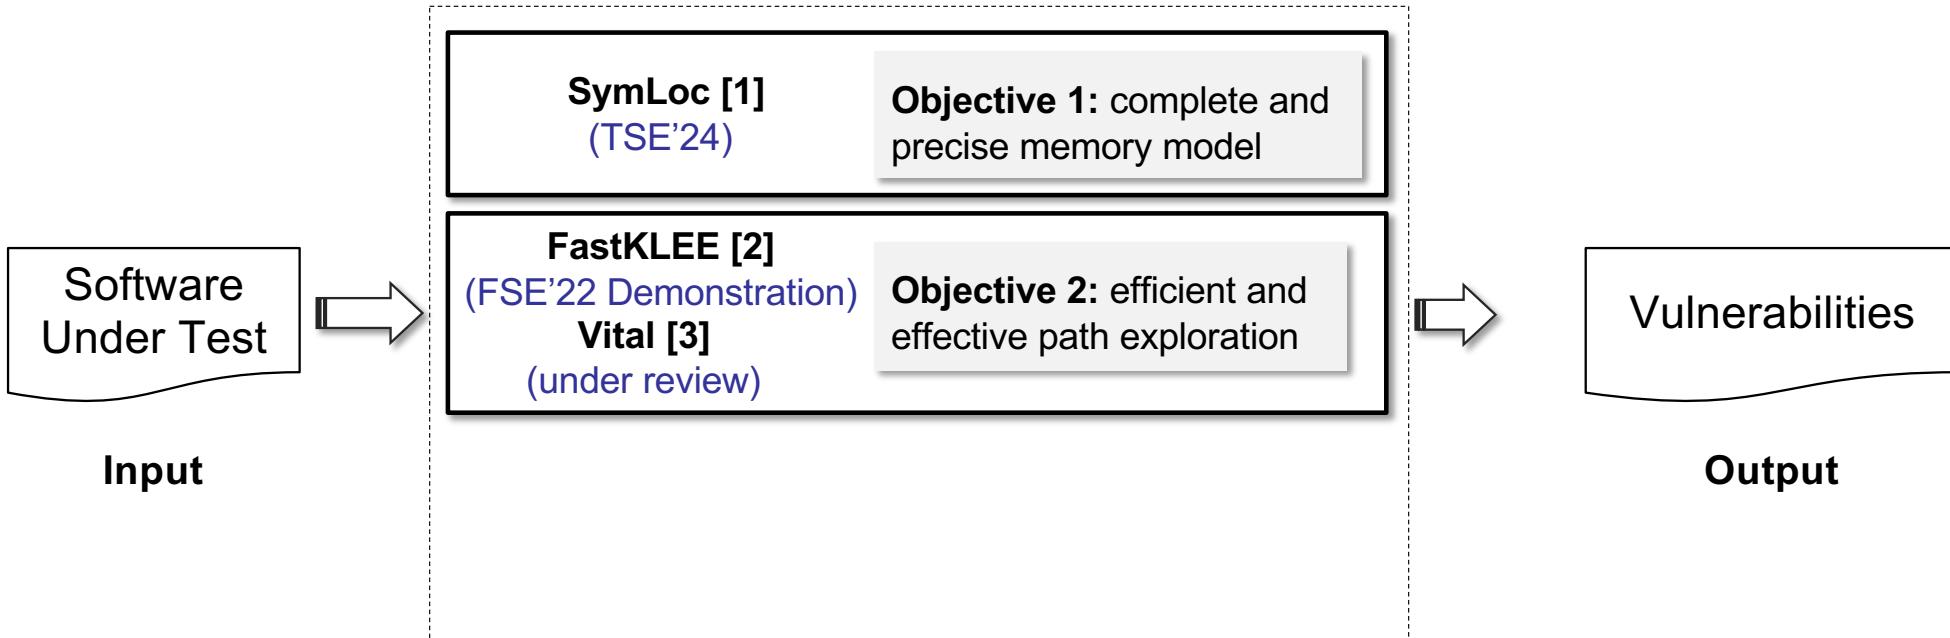
[2] Haixin Tu, Lingxiao Jiang, and et.al., "FastKLEE: Faster Symbolic Execution via Redundant Bound Checking of Type-Safe Pointers", in Tool Demonstrations Track of FSE 2022.

[3] Haixin Tu, Lingxiao Jiang and et.al., "Vital: Vulnerability-Oriented Symbolic Execution via Type-Unsafe Pointer-Guided Monte Carlo Tree Search." arXiv:2408.08772 (2024).

[4] Haixin Tu, Seongmin Lee, and et.al. "Large Language Model-Driven Concolic Execution for Highly Structured Test Input Generation." arXiv:2504.17542 (2025).

Thesis statement and research objectives

- **Statement:** This thesis aims to **boost symbolic execution** by designing new solutions to alleviate three key limitations in **memory modeling/path exploration/test input generation**, for efficient and effective automatic vulnerability detection.



[1] Haixin Tu, Lingxiao Jiang, and et.al, "Concretely Mapped Symbolic Memory Locations for Memory Error Detection" (IEEE TSE).

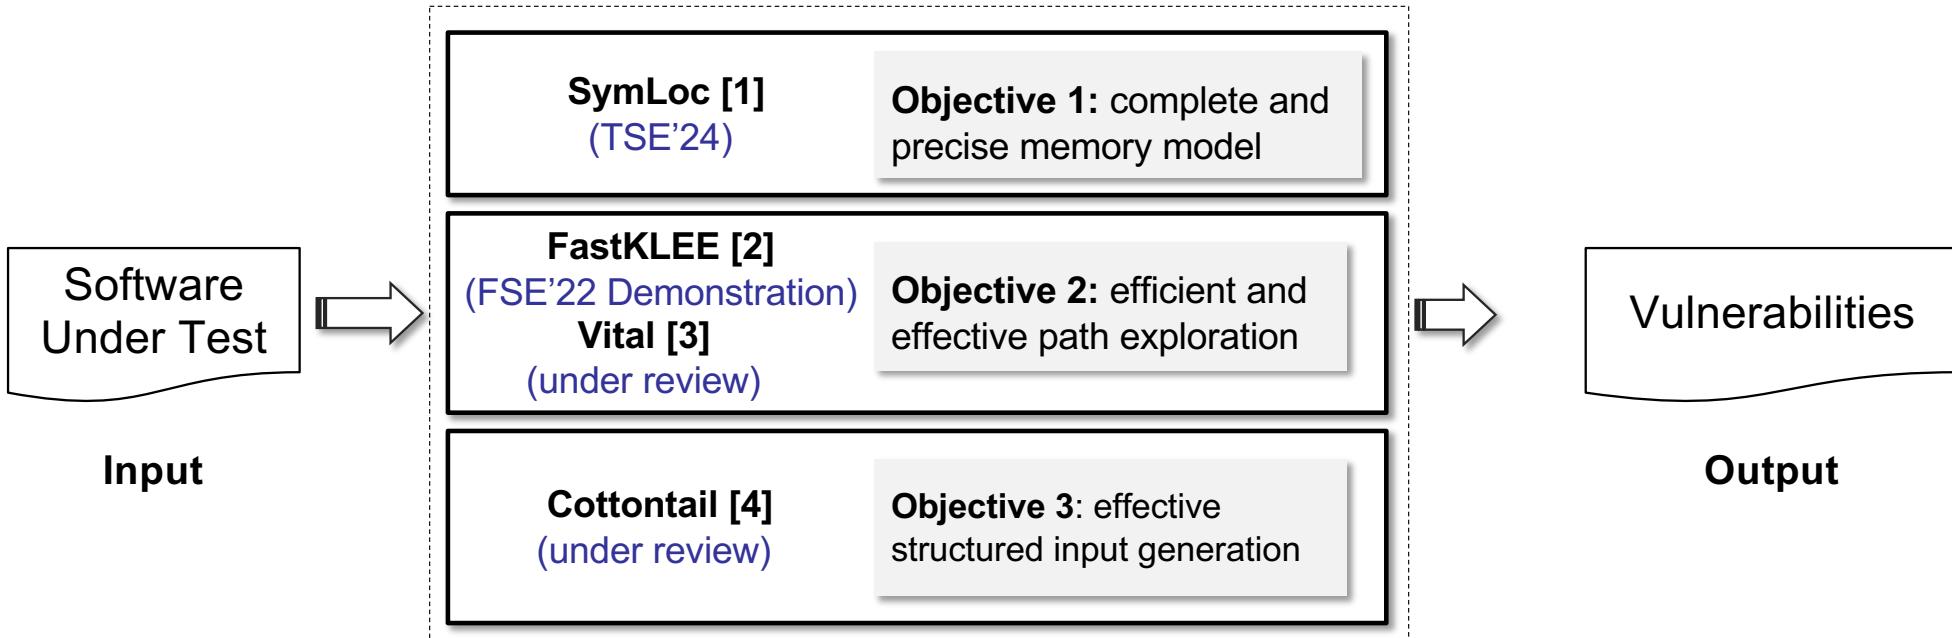
[2] Haixin Tu, Lingxiao Jiang, and et.al., "FastKLEE: Faster Symbolic Execution via Redundant Bound Checking of Type-Safe Pointers", in Tool Demonstrations Track of FSE 2022.

[3] Haixin Tu, Lingxiao Jiang and et.al., "Vital: Vulnerability-Oriented Symbolic Execution via Type-Unsafe Pointer-Guided Monte Carlo Tree Search." arXiv:2408.08772 (2024).

[4] Haixin Tu, Seongmin Lee, and et.al. "Large Language Model-Driven Concolic Execution for Highly Structured Test Input Generation." arXiv:2504.17542 (2025).

Thesis statement and research objectives

- **Statement:** This thesis aims to **boost symbolic execution** by designing new solutions to alleviate three key limitations in **memory modeling/path exploration/test input generation**, for efficient and effective automatic vulnerability detection.



[1] Haixin Tu, Lingxiao Jiang, and et.al, "Concretely Mapped Symbolic Memory Locations for Memory Error Detection" (IEEE TSE).

[2] Haixin Tu, Lingxiao Jiang, and et.al., "FastKLEE: Faster Symbolic Execution via Redundant Bound Checking of Type-Safe Pointers", in Tool Demonstrations Track of FSE 2022.

[3] Haixin Tu, Lingxiao Jiang and et.al., "Vital: Vulnerability-Oriented Symbolic Execution via Type-Unsafe Pointer-Guided Monte Carlo Tree Search." arXiv:2408.08772 (2024).

[4] Haixin Tu, Seongmin Lee, and et.al. "Large Language Model-Driven Concolic Execution for Highly Structured Test Input Generation." arXiv:2504.17542 (2025).

Motivation Example (with our solutions)

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```



Motivation Example (with our solutions)

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```

*Can our boosted symbolic execution
effectively detect the three vulnerabilities?*

YES

Outline

□ Background

- Software vulnerabilities
- Symbolic execution

□ Motivation

- What limitations prevent current symbolic execution for vulnerability detection?
- Thesis statement and research objects

□ Methodology

- **SymLoc**: A new memory model for symbolic execution
- **FastKLEE and Vital**: Two new path exploration for symbolic execution
- **Cottontail**: A new structured input generation for symbolic execution

□ Future work

- Extension of current solutions
- Combine Program Analysis with GenAI

□ Conclusion and Acknowledgment

A New Memory Model: insights

A New Memory Model: insights

- Three fundamental designs are required



A New Memory Model: insights

- **Three fundamental designs are required**
 - A. Symbolization of addresses and modeling them into path constraints



A New Memory Model: insights

➤ Three fundamental designs are required

- A. Symbolization of addresses and modeling them into path constraints
- B. Practical read/write operation from/to symbolic addresses



A New Memory Model: insights

➤ Three fundamental designs are required

- A. Symbolization of addresses and modeling them into path constraints
- B. Practical read/write operation from/to symbolic addresses
- C. Effectively tracking the uses of symbolic addresses



A New Memory Model: insights

➤ Three fundamental designs are required

- A. Symbolization of addresses and modeling them into path constraints
- B. Practical read/write operation from/to symbolic addresses
- C. Effectively tracking the uses of symbolic addresses



➤ Existing approaches are difficult to satisfy all above requirements

- KLEE and Symsize (FSE'21): none of the requirements can be satisfied
- RAM (ICSE'18): satisfies requirements #B and partially #A but not #C
- MemSight (ASE'17): satisfies requirements #A and #B but not #C

Rethinking Pointer Reasoning
in Symbolic Execution

Emilio Coppa, Daniele Cenno D'Elia, and Camil Demetrescu
Department of Computer, Control, and Management Engineering

Relocatable Addressing Model for Symbolic Execution

David Trabish
Tel Aviv University
Israel
davivtra@post.tau.ac.il

Noam Rinotzky
Tel Aviv University
Israel
maon@cs.tau.ac.il

A Bounded Symbolic-Size Model for Symbolic Execution

David Trabish
Tel Aviv University
Israel
davivtra@post.tau.ac.il

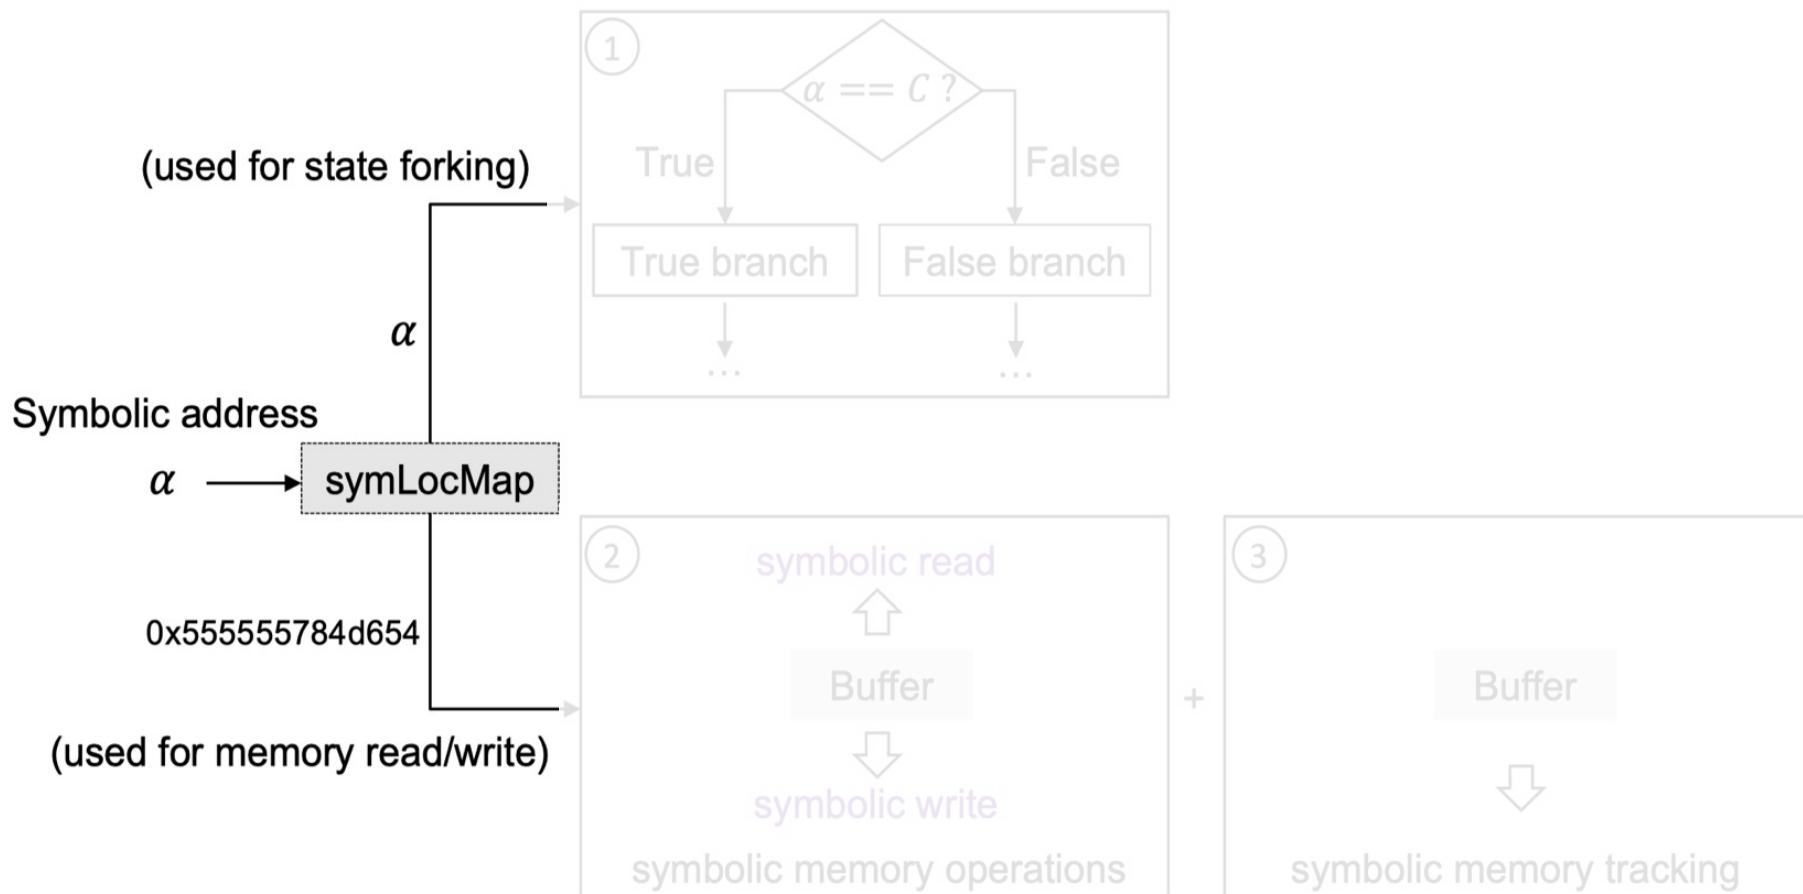
Shachar Itzhaky
Technion
Israel
shachari@cs.technion.ac.il

Noam Rinotzky
Tel Aviv University
Israel
maon@cs.tau.ac.il

Solution: SymLoc (1/3)

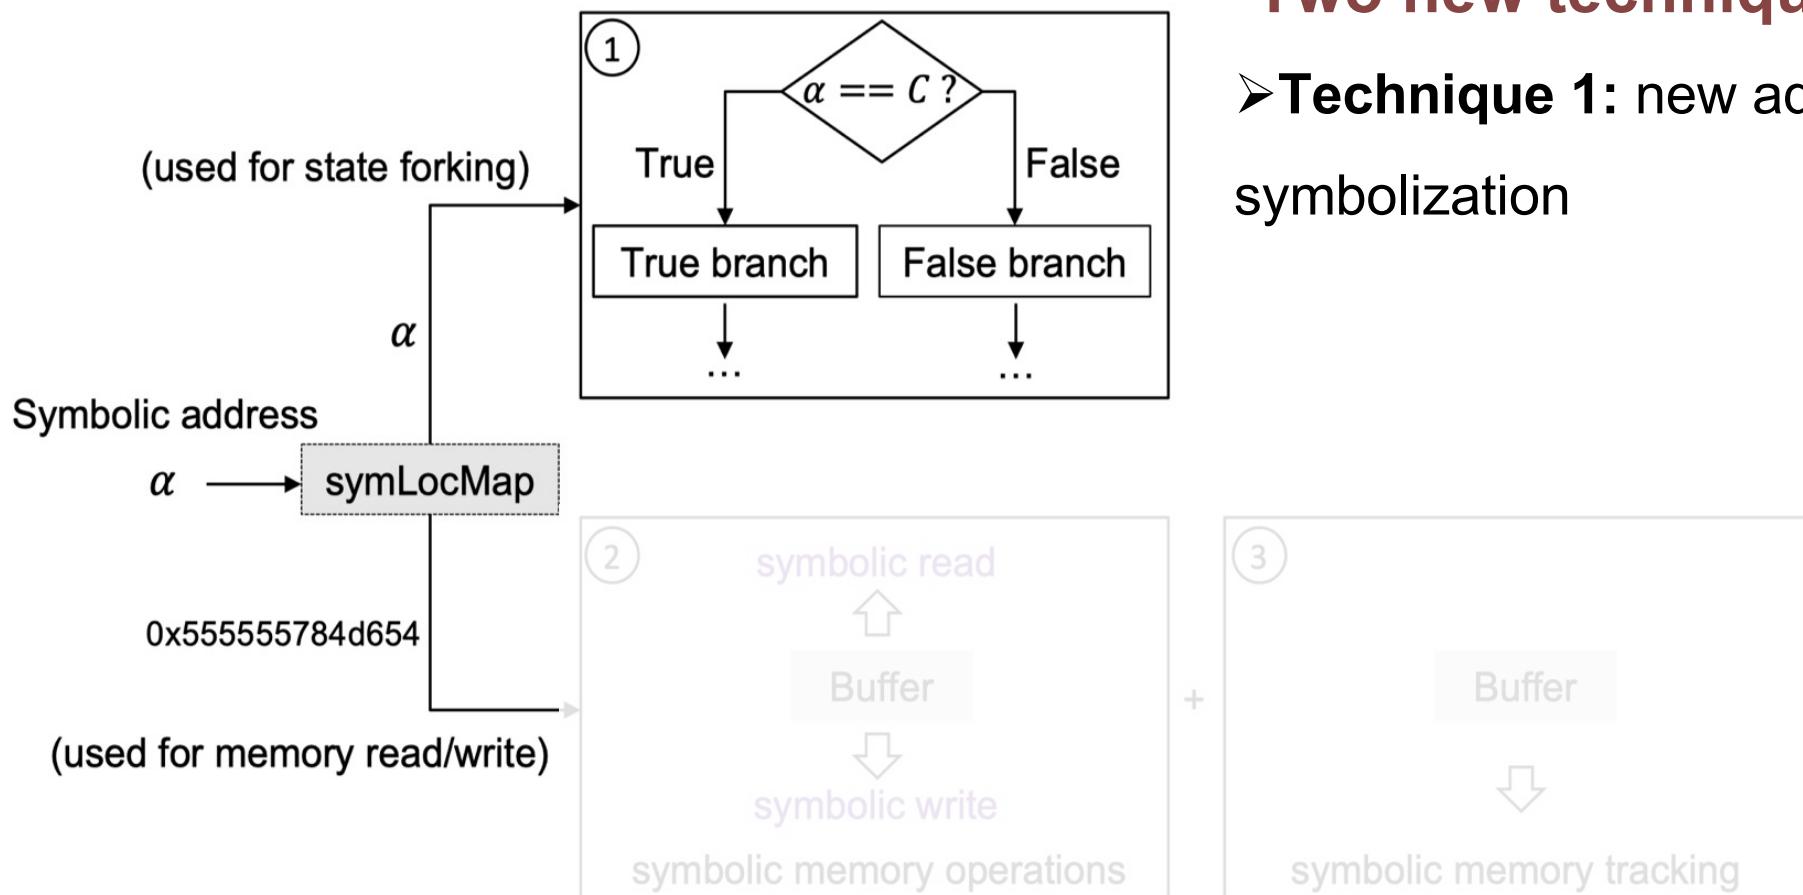
Solution: SymLoc (1/3)

High-level Idea



Solution: SymLoc (1/3)

High-level Idea

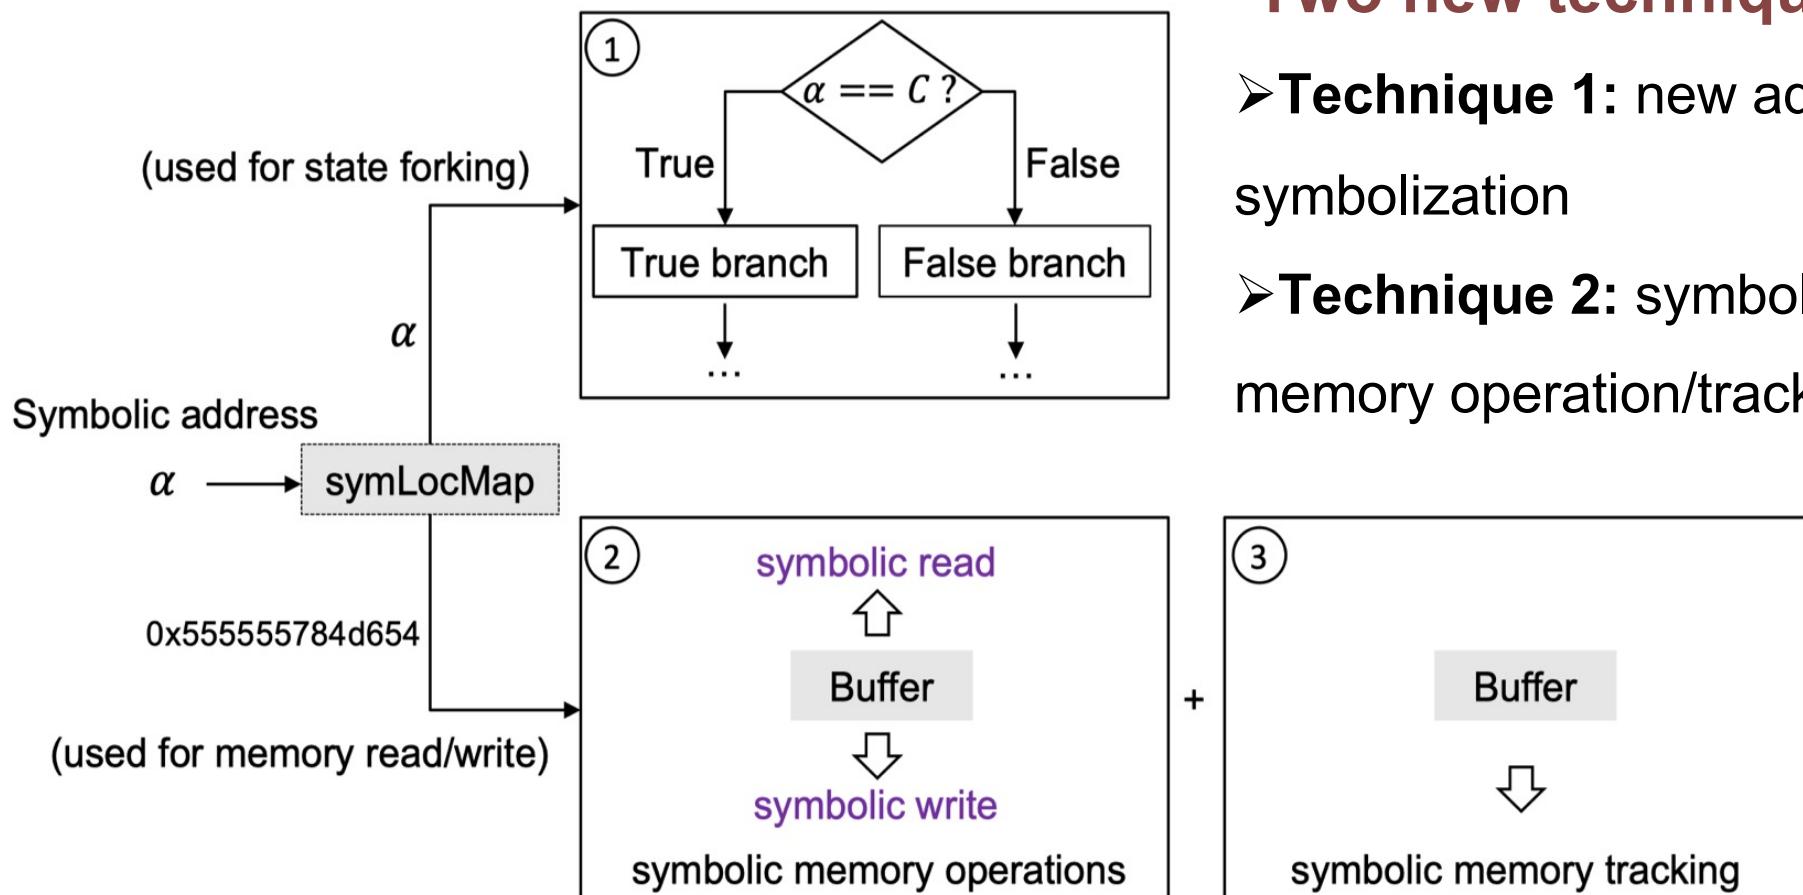


Two new techniques

- **Technique 1: new address symbolization**

Solution: SymLoc (1/3)

High-level Idea

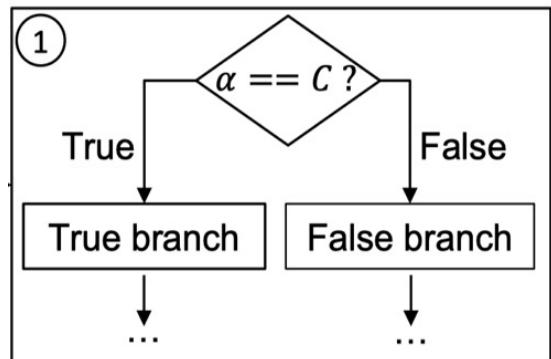


Two new techniques

- **Technique 1:** new address symbolization
- **Technique 2:** symbolic memory operation/tracking

Solution: SymLoc (2/3)

Solution: SymLoc (2/3)

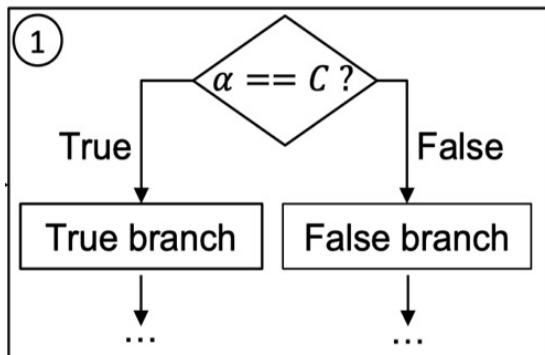


Solution: SymLoc (2/3)

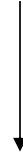
□ New address symbolization

■ Symbolic addressing model

- Encoding the symbolic address into path constraints



Existing: $(addr, size, arry) \in N^+ \times N^+ \times A$



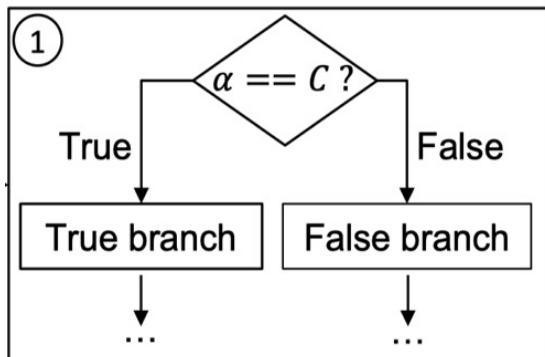
Ours: $(symAddr, size, arry) \in N^+ \times N^+ \times A$

Solution: SymLoc (2/3)

□ New address symbolization

■ Symbolic addressing model

- Encoding the symbolic address into path constraints



Existing: $(addr, size, arry) \in N^+ \times N^+ \times A$

Ours: $(symAddr, size, arry) \in N^+ \times N^+ \times A$

■ Input

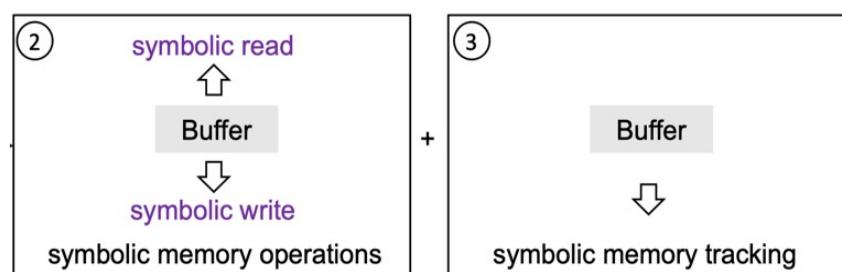
- A set of variables to return from malloc function

■ Output

- A symbolic-concrete memory map (symLocMap)
- Will be used in the latter phase

Solution: SymLoc (3/3)

Solution: SymLoc (3/3)



Solution: SymLoc (3/3)

□ Symbolic memory operation and tracking

Algorithm 1: Symbolic memory operations and tracking

Input: the map symLocMap, a symbolic expression symExpr,
and a function func being executed
Output: a concrete or symbolic expression, or an error

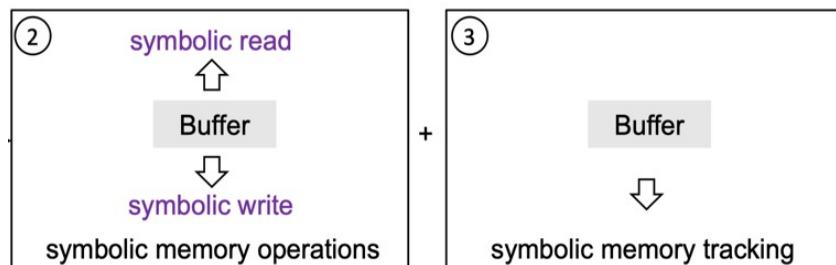
```
1 conExpr ← ∅ // initialize a concrete expression
2 FreeList ← ∅ // initialize a list to store freed objects
3 Function SymAddrRes (symLocMap, symExpr, func):
...
    ↑
```

■ Input

- symLocMap, a symbolic expression, and a function

■ Output

- A concrete address or normal symbolic variable or a bug



Solution: SymLoc (3/3)

□ Symbolic memory operation and tracking

Algorithm 1: Symbolic memory operations and tracking

Input: the map symLocMap, a symbolic expression symExpr, and a function func being executed
Output: a concrete or symbolic expression, or an error

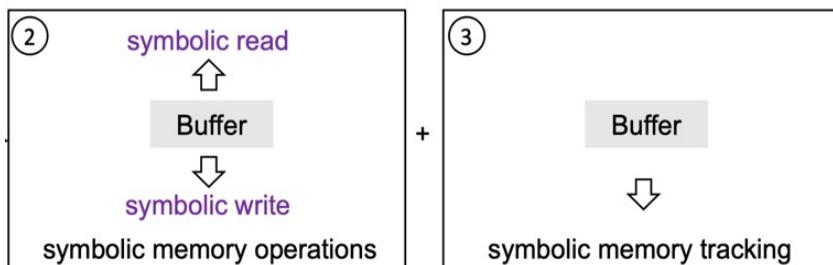
```
1 conExpr ← ∅ // initialize a concrete expression
2 FreeList ← ∅ // initialize a list to store freed objects
3 Function SymAddrRes (symLocMap, symExpr, func):
    ...
    ↑
```

■ Input

- symLocMap, a symbolic expression, and a function

■ Output

- A concrete address or normal symbolic variable or a bug



■ Tracking example

- A memory address is symbolized as “sym_a”
- If the freed object is “sym_a” or “sym_a + 100”
 - Indicating UAF bugs

SymLoc: Evaluation Results & Takeaways



School of
Computing and
Information Systems

SymLoc: Evaluation Results & Takeaways

RQ1: How does SymLoc perform in detecting spatial memory errors?

- SymLoc could cover **15%** and **48%** more unique lines of code on average than the two baseline approaches.
- SymLoc could detect **169%** and **218%** more spatial memory errors than the two baseline approaches.
- **Two** new vulnerabilities have been detected

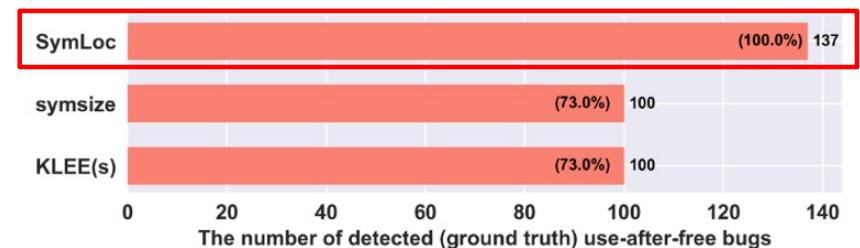
SymLoc: Evaluation Results & Takeaways

RQ1: How does SymLoc perform in detecting spatial memory errors?

- SymLoc could cover **15%** and **48%** more unique lines of code on average than the two baseline approaches.
- SymLoc could detect **169%** and **218%** more spatial memory errors than the two baseline approaches.
- **Two** new vulnerabilities have been detected

RQ2: How does SymLoc perform in detecting temporal memory errors?

- SymLoc has an **overall better** temporal memory error detection capability for detecting UAF and DoF errors than static, dynamic, and symbolic execution-based approaches.



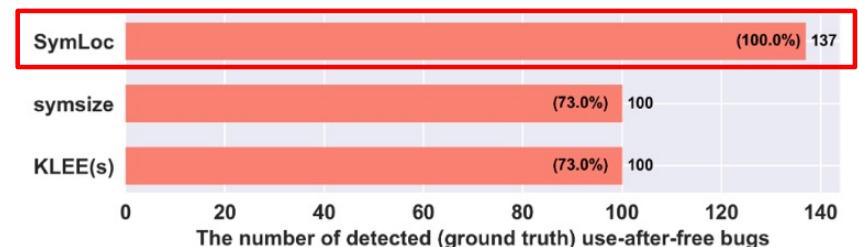
SymLoc: Evaluation Results & Takeaways

RQ1: How does SymLoc perform in detecting spatial memory errors?

- SymLoc could cover **15%** and **48%** more unique lines of code on average than the two baseline approaches.
- SymLoc could detect **169%** and **218%** more spatial memory errors than the two baseline approaches.
- **Two** new vulnerabilities have been detected

RQ2: How does SymLoc perform in detecting temporal memory errors?

- SymLoc has an **overall better** temporal memory error detection capability for detecting UAF and DoF errors than static, dynamic, and symbolic execution-based approaches.



• Takeaways

- A more complete memory model could help detect tricky vulnerabilities

Outline

□ Background

- Software vulnerabilities
- Symbolic execution

□ Motivation

- What limitations prevent current symbolic execution for vulnerability detection?
- Thesis statement and research objects

□ Methodology

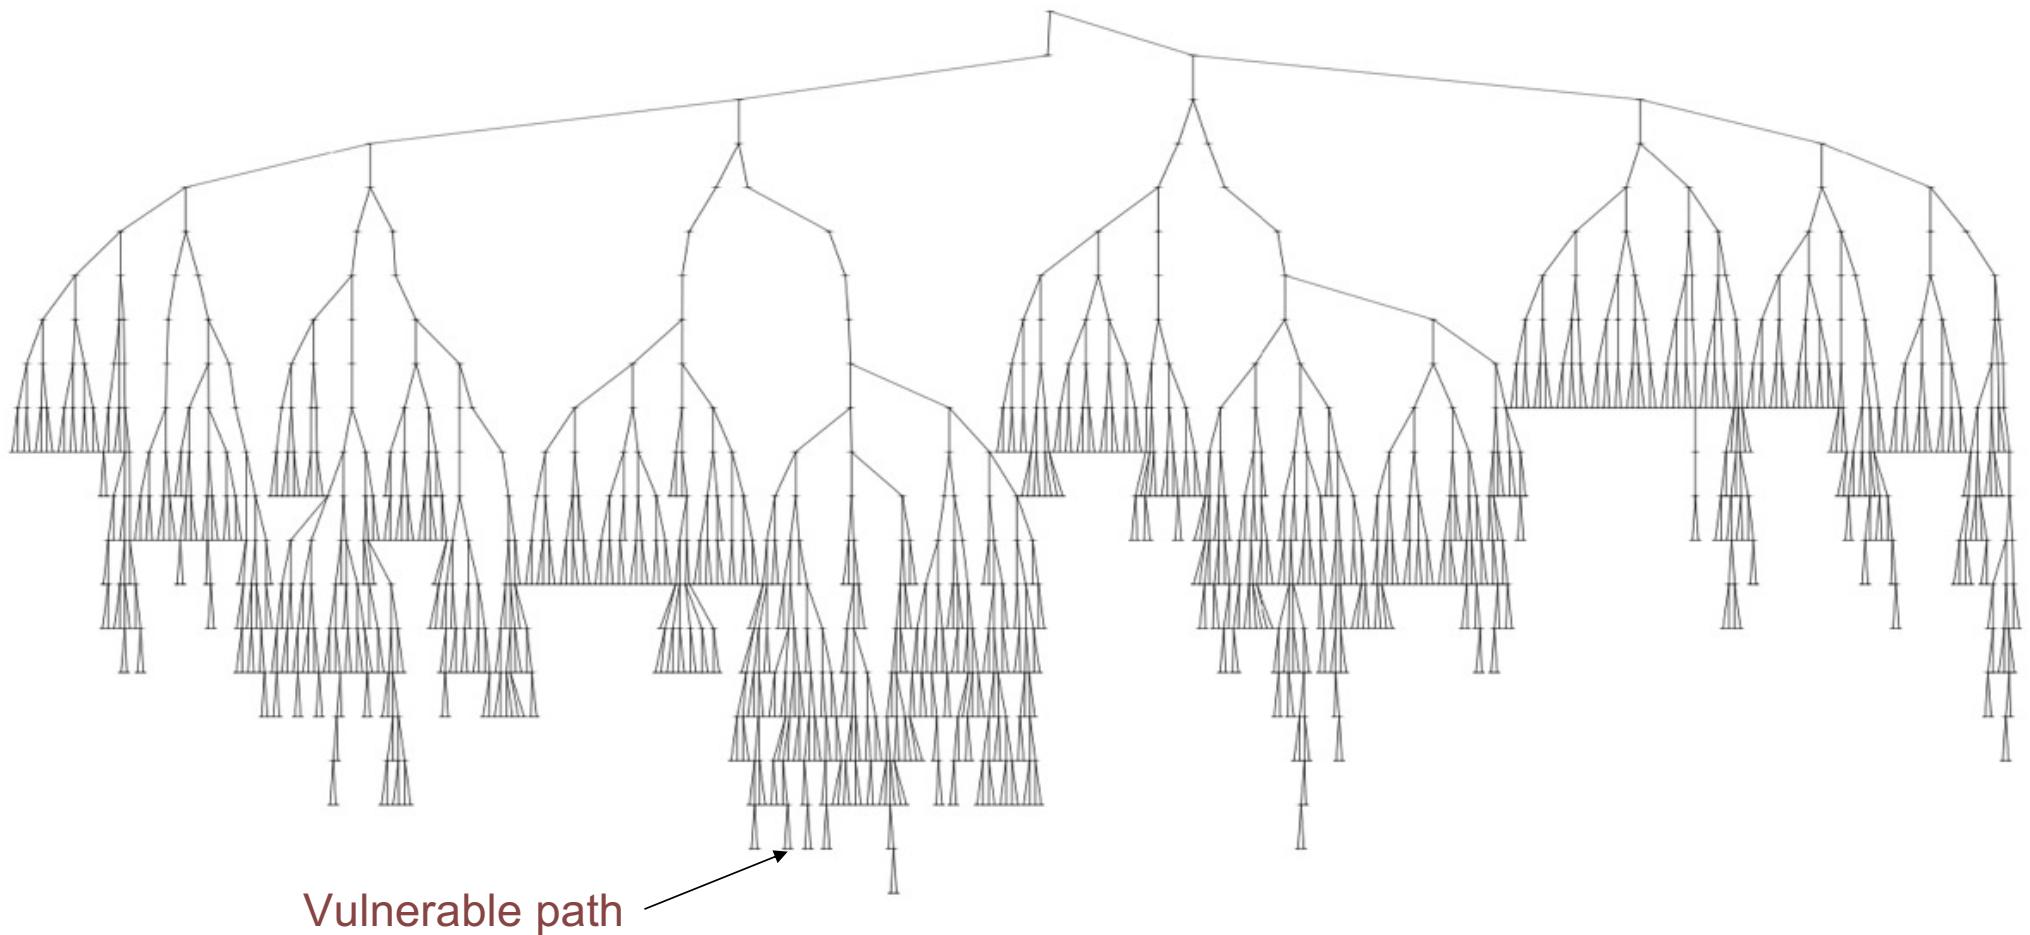
- **SymLoc**: A new memory model for symbolic execution
- **FastKLEE and Vital**: Two new path exploration for symbolic execution
- **Cottontail**: A new structured input generation for symbolic execution

□ Future work

- Extension of current solutions
- Combine Program Analysis with GenAI

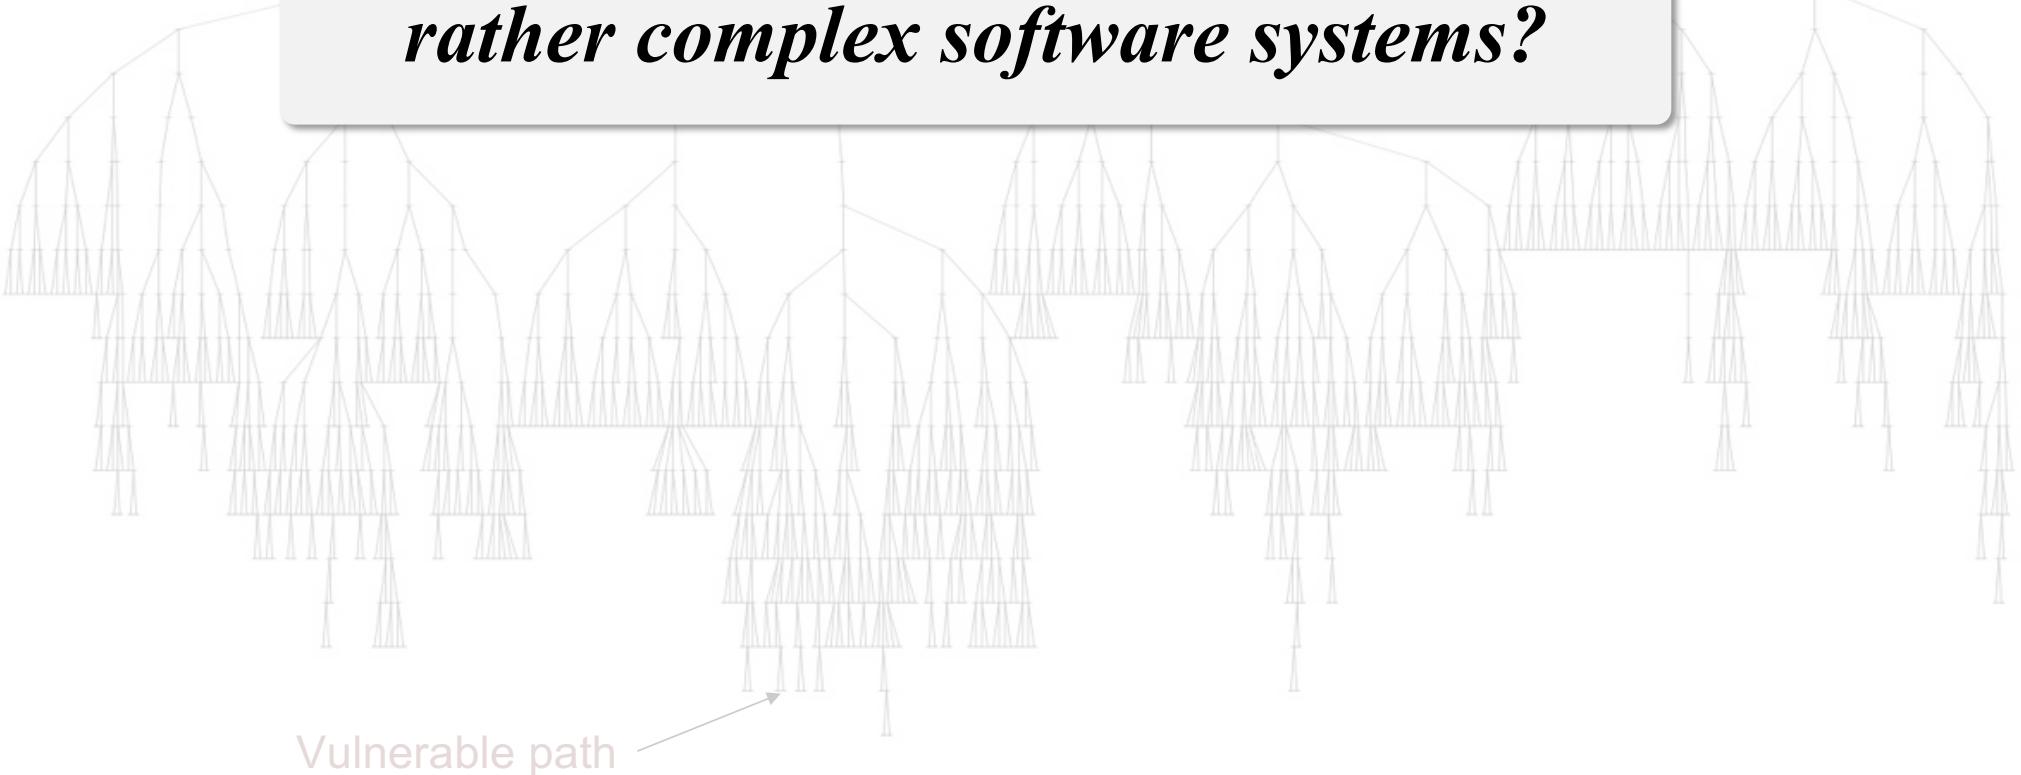
□ Conclusion and Acknowledgment

Two New Path Explorations (1/2)



Two New Path Explorations (1/2)

*How to search vulnerable paths in
rather complex software systems?*



Two New Path Explorations (1/2)

*How to search vulnerable paths in
rather complex software systems?*

*Direction 1: Can
we do the path
search faster?
(FastKLEE)*

Vulnerable path

*Direction 2: Can
we search vuln-
erable paths first?
(Vital)*

Two New Path Explorations (1/2)

*How to search vulnerable paths in
rather complex software systems?*

*Direction 1: Can
we do the path
search faster?
(FastKLEE)*

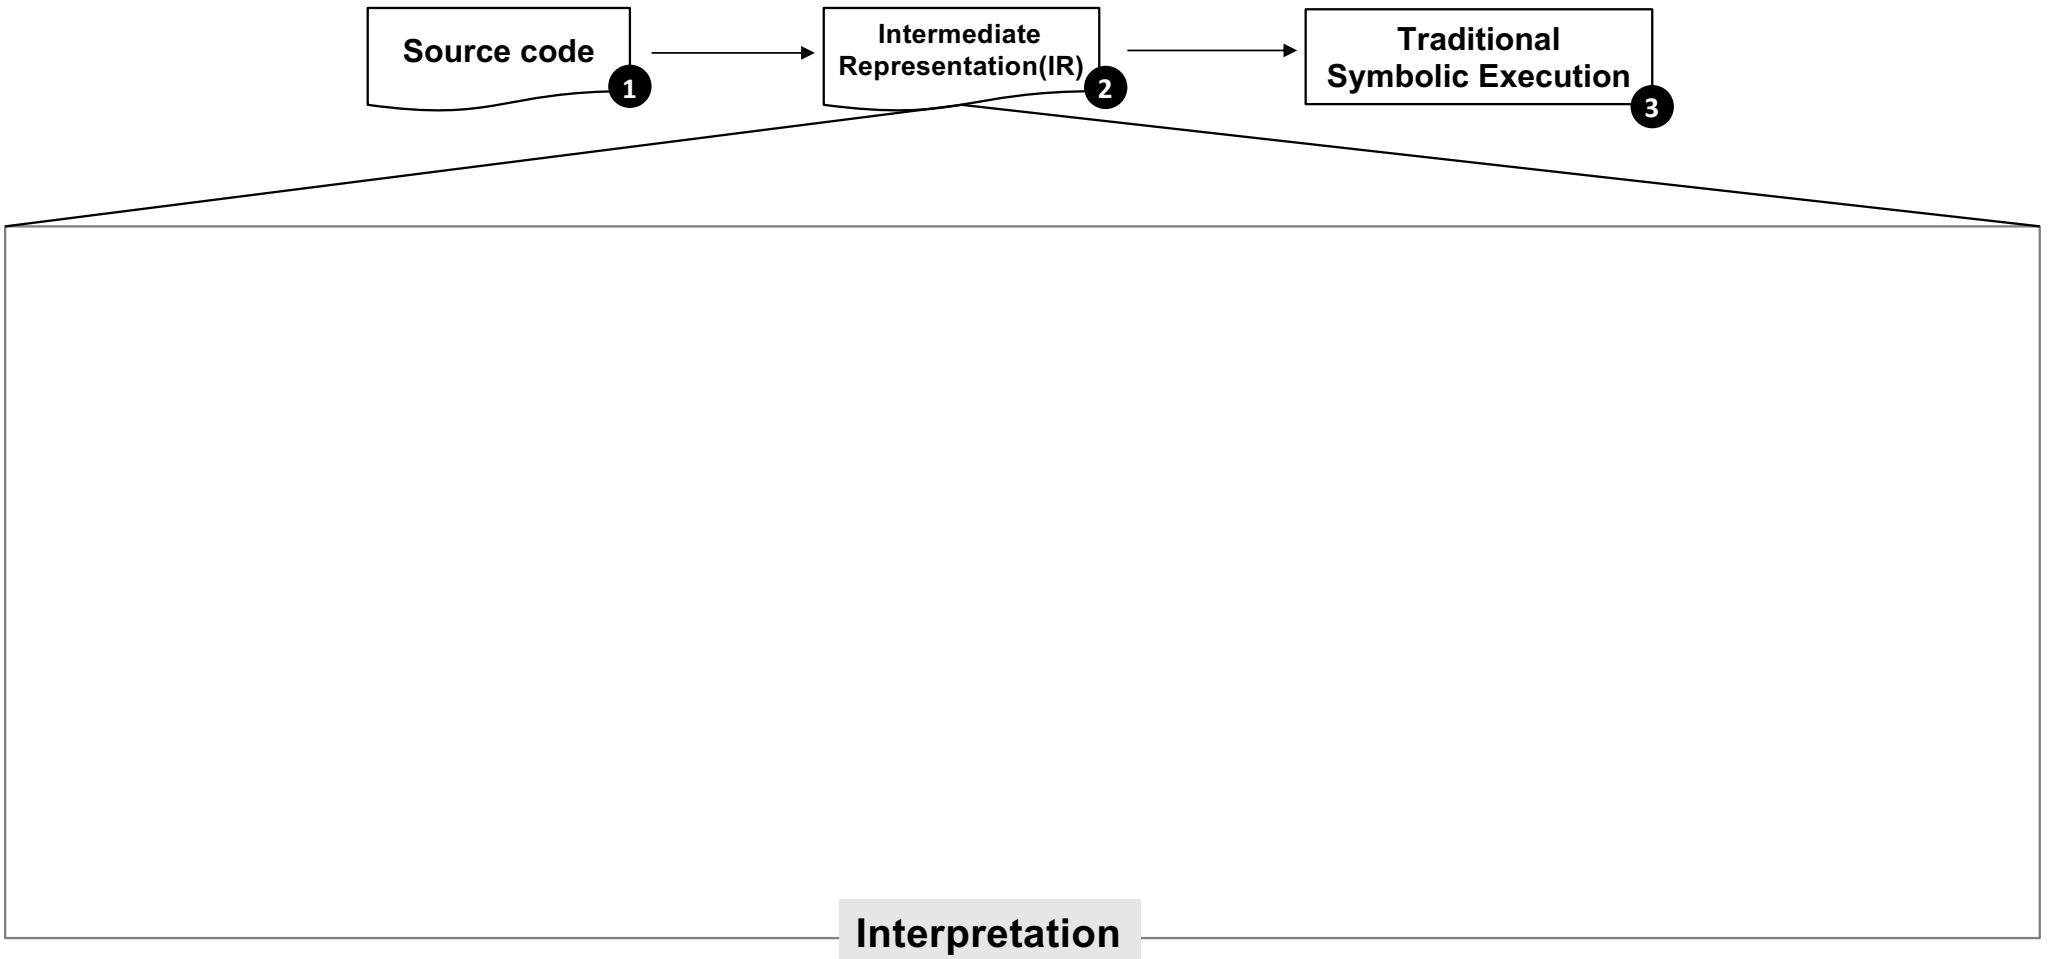
*Direction 2: Can
we search vuln-
erable paths first?
(Vital)*

Vulnerable path

Efficient Path Exploration: Rethinking



Efficient Path Exploration: Rethinking



Efficient Path Exploration: Rethinking



(1) Observation

- The number of interpreted instructions tends to be **huge** (several billion only in one hour run)

```
Elapsed: 01:00:04
KLEE: done: explored paths = 125017
KLEE: done: avg. constructs per query = 74
KLEE: done: total queries = 8859
KLEE: done: valid queries = 6226
KLEE: done: invalid queries = 2633
KLEE: done: query cex = 8859

KLEE: done: total instructions = 605113213
KLEE: done: completed paths = 125017
KLEE: done: generated tests = 65
```

Interpretation

Efficient Path Exploration: Rethinking



(1) Observation

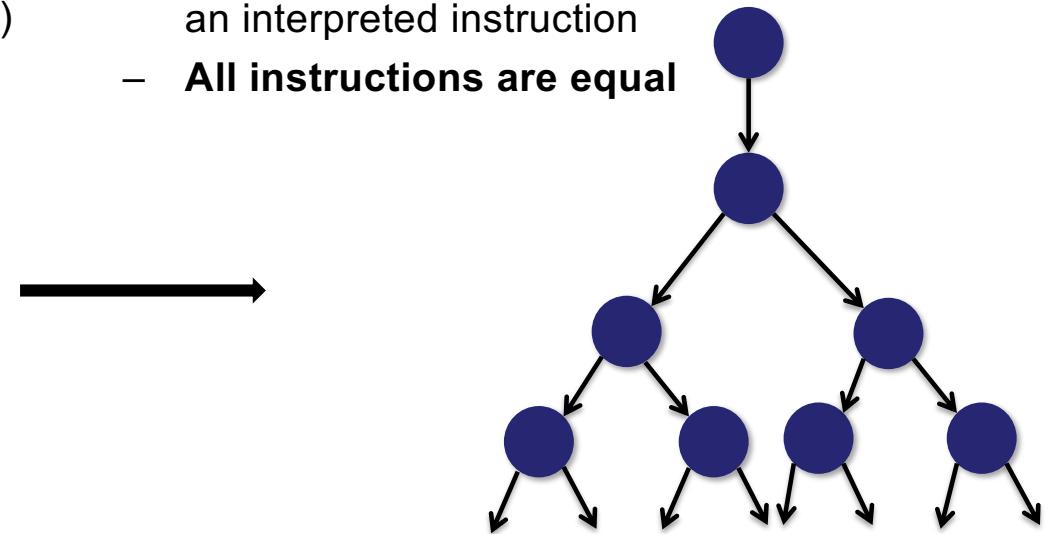
- The number of interpreted instructions tends to be **huge** (several billion only in one hour run)

```
Elapsed: 01:00:04
KLEE: done: explored paths = 125017
KLEE: done: avg. constructs per query = 74
KLEE: done: total queries = 8859
KLEE: done: valid queries = 6226
KLEE: done: invalid queries = 2633
KLEE: done: query cex = 8859

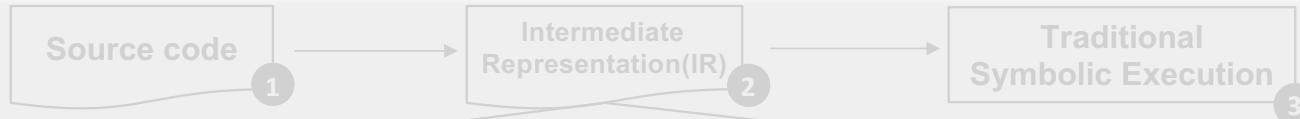
KLEE: done: total instructions = 605113213
KLEE: done: completed paths = 125017
KLEE: done: generated tests = 65
```

(2) Overheads in current symbolic execution

- The color depth represents the overheads of an interpreted instruction
- **All instructions are equal**



Efficient Path Exploration: Rethinking



(1) Observation

- The number of interpreted instructions tends to

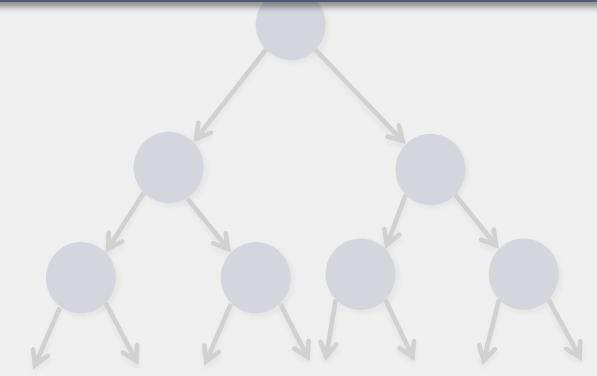
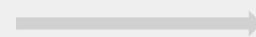
(2) Overheads in current symbolic execution

- The color depth represents the overheads of

Can we reduce the overhead of interpreted instructions for faster symbolic execution?

```
KLEE: done: avg. constructs per query = 74
KLEE: done: total queries = 8859
KLEE: done: valid queries = 6226
KLEE: done: invalid queries = 2633
KLEE: done: query cex = 8859

KLEE: done: total instructions = 605113213
KLEE: done: completed paths = 125017
KLEE: done: generated tests = 65
```



Interpretation

Solution: FastKLEE (1/2)

- Key insights

Solution: FastKLEE (1/2)

- **Key insights**
 - Only a small portion of memory-related instructions need bound checking

Solution: FastKLEE (1/2)

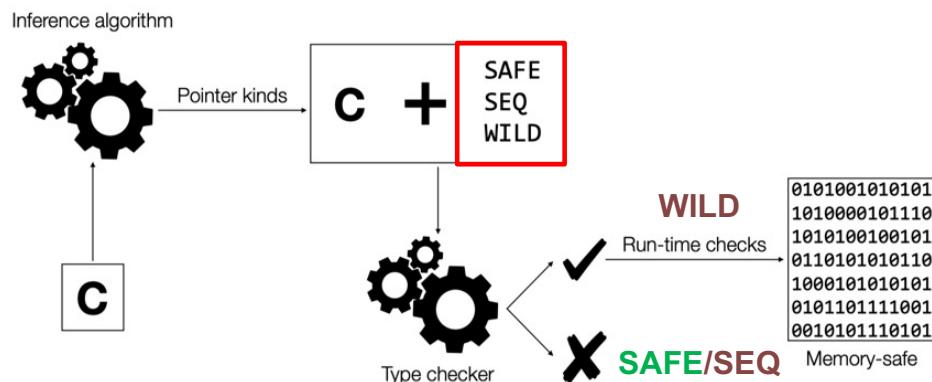
- **Key insights**

- Only a small portion of memory-related instructions need bound checking
- Reduce the interpreting overhead of the most frequently interpreted ones (i.e., load/store instructions)

Solution: FastKLEE (1/2)

- **Key insights**

- Only a small portion of memory-related instructions need bound checking
- Reduce the interpreting overhead of the most frequently interpreted ones (i.e., load/store instructions)
- Inspired by *Type Inference* system [1]

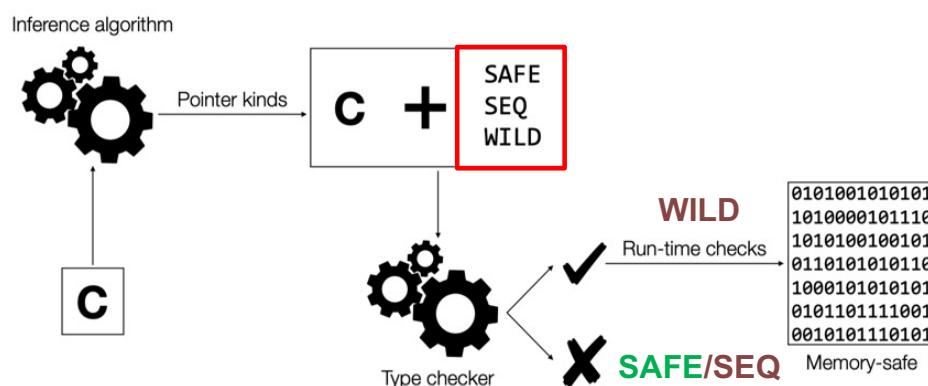


[1] CCured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. CCured: type-safe retrofitting of legacy software. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Solution: FastKLEE (1/2)

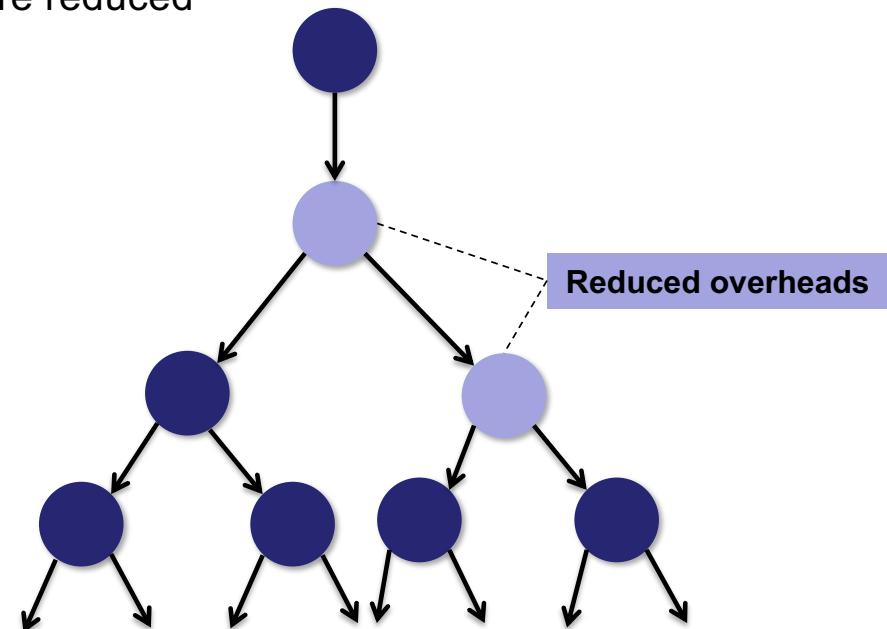
- **Key insights**

- Only a small portion of memory-related instructions need bound checking
- Reduce the interpreting overhead of the most frequently interpreted ones (i.e., load/store instructions)
- Inspired by *Type Inference* system [1]



- **Advantage: overheads in FastKLEE**

- Interpretation overheads for some instructions are reduced



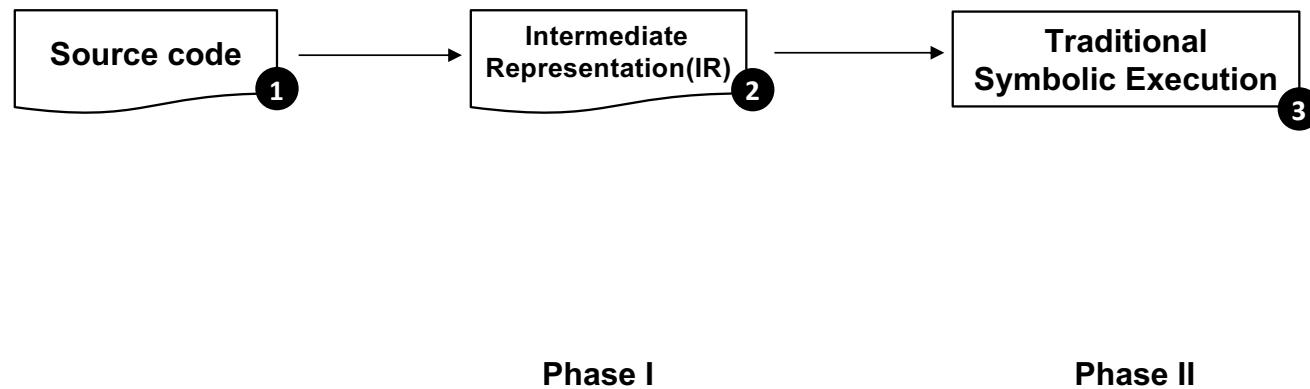
[1] CCured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. CCured: type-safe retrofitting of legacy software. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Solution: FastKLEE (2/2)

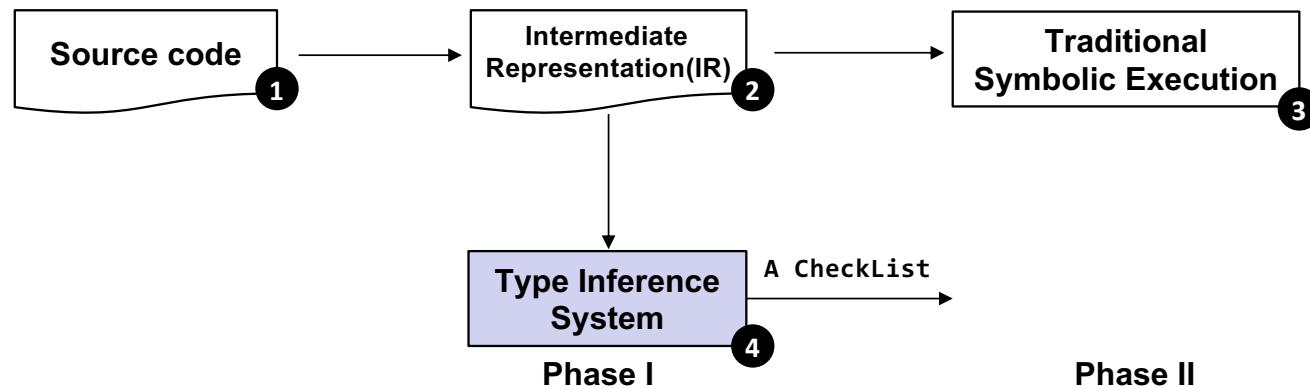
Solution: FastKLEE (2/2)



Solution: FastKLEE (2/2)

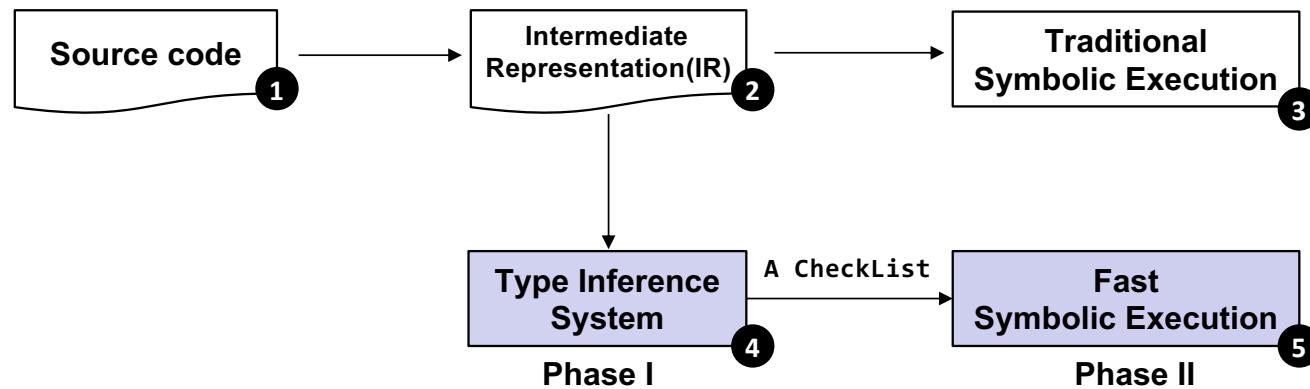


Solution: FastKLEE (2/2)



- ④ • **Phase I:** Introduce a **Type Inference System** to classify memory-related instruction types
 - **Unsafe** memory instructions will be stored in **CheckList**

Solution: FastKLEE (2/2)



- ④ • **Phase I:** Introduce a **Type Inference System** to classify memory-related instruction types
 - **Unsafe** memory instructions will be stored in **CheckList**
- ⑤ • **Phase II:** Conduct **Customized Memory Operation** in Fast symbolic execution
 - Only perform checking for **Unsafe** memory instructions during interpretation

Two New Path Explorations (2/2)

*How to search vulnerable paths in
rather complex software systems?*

*Direction 1: Can
we do the path
search faster?
(FastKLEE)*

Vulnerable path

*Direction 2: Can
we search vuln-
erable paths first?
(Vital)*

Insight 1: approximate a vulnerable path

[1] Ccured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. **CCured: type-safe retrofitting of legacy software**. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Insight 1: approximate a vulnerable path

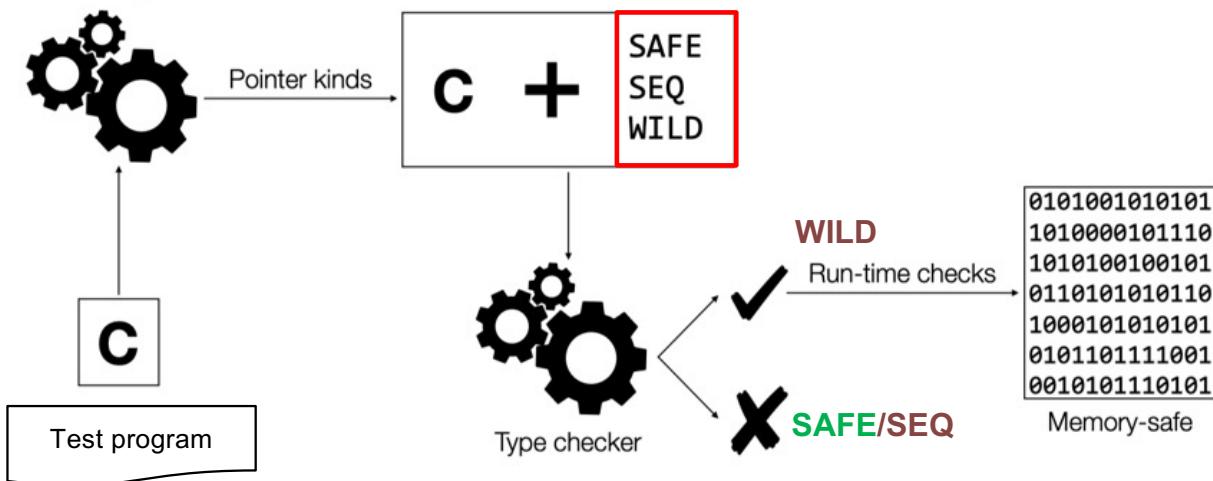
- Revisit type inference (Ccured [1])

[1] Ccured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. **CCured: type-safe retrofitting of legacy software**. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Insight 1: approximate a vulnerable path

- Revisit type inference (Ccured [1])

Inference algorithm

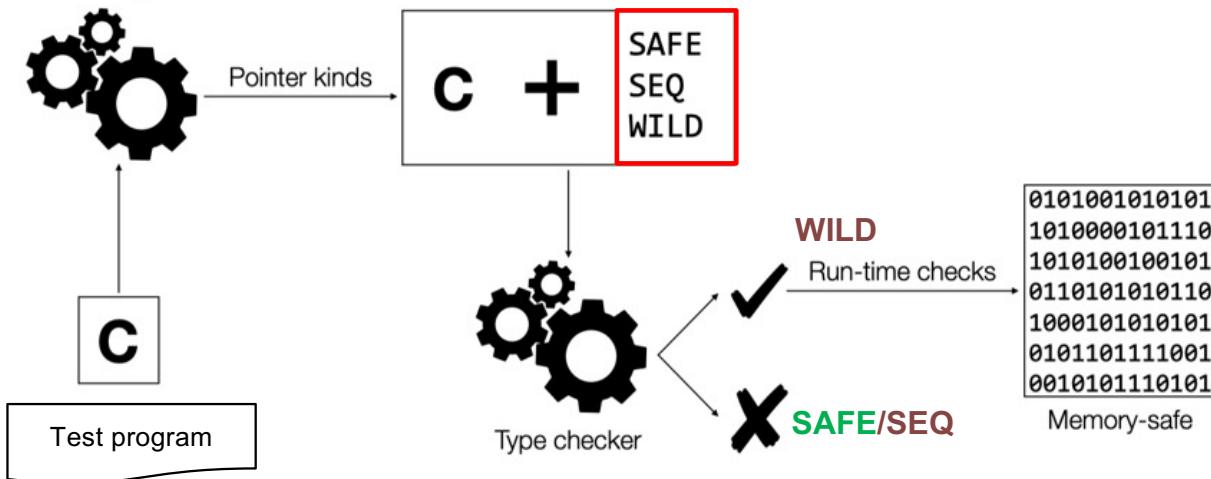


[1] Ccured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. CCured: type-safe retrofitting of legacy software. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Insight 1: approximate a vulnerable path

- Revisit type inference (Ccured [1])

Inference algorithm



Observation: Memory errors can only happen in **unsafe** pointer operations (SEQ/WILD)!

Insight 1: approximate a vulnerable path

- Revisit type inference (Ccure [1])

Inference algorithm



Pointer kinds



Test program



```
int * p = malloc (100);  
  
*(p + "input") = 1; // unsafe pointer operation
```

0101001010101
1010000101110
1010100100101
0110101010110
1000101010101
0101101111001
0010101110101

Memory-safe

Observation: Memory errors can only happen in **unsafe** pointer operations (SEQ/WILD)!

Insight 1: approximate a vulnerable path

- Revisit type inference (Ccurd [1])

Inference algorithm



Pointer kinds



Test program



```
int * p = malloc (100);  
  
*(p + "input") = 1; // unsafe pointer operation
```

0101001010101
1010000101110
1010100100101
0110101010110
1000101010101
0101101111001
0010101110101

Memory-safe

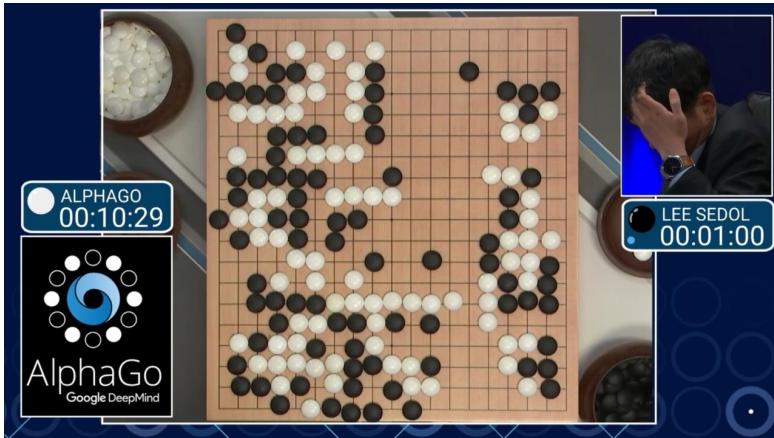
Observation: Memory errors can only happen in **unsafe** pointer operations (SEQ/WILD)!

We can exploit such information from type inference to guide the path search!

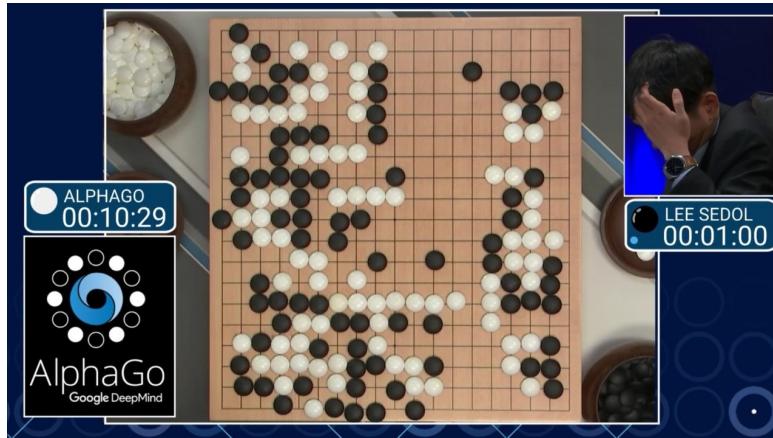
[1] Ccured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. CCured: type-safe retrofitting of legacy software. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Insight 2: search the path smartly

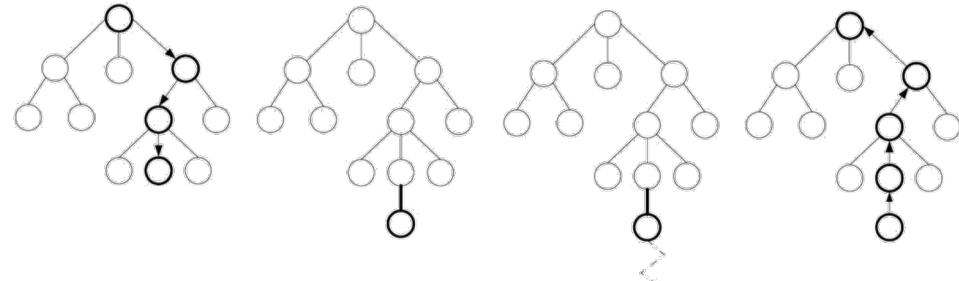
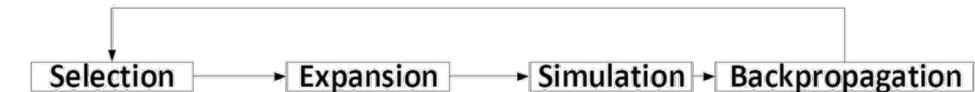
Insight 2: search the path smartly



Insight 2: search the path smartly

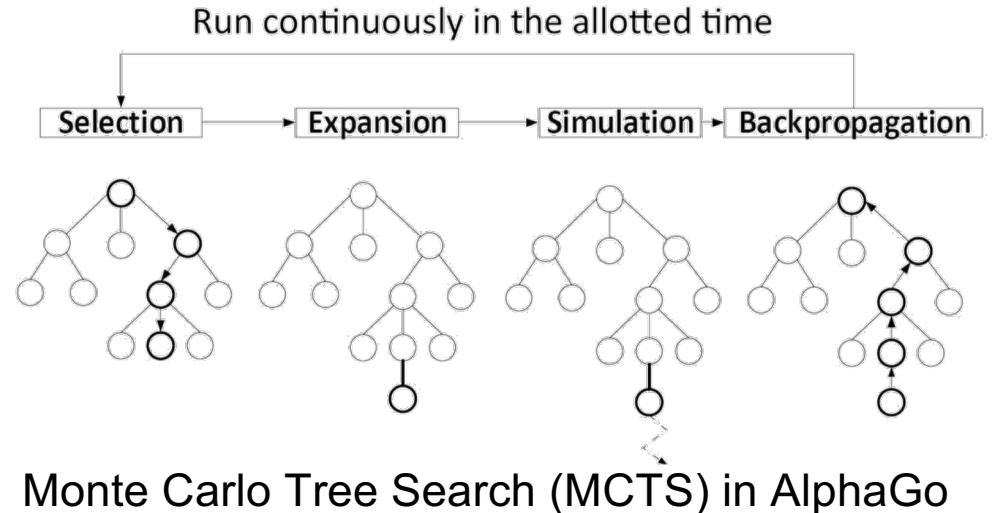


Run continuously in the allotted time



Monte Carlo Tree Search (MCTS) in AlphaGo

Insight 2: search the path smartly

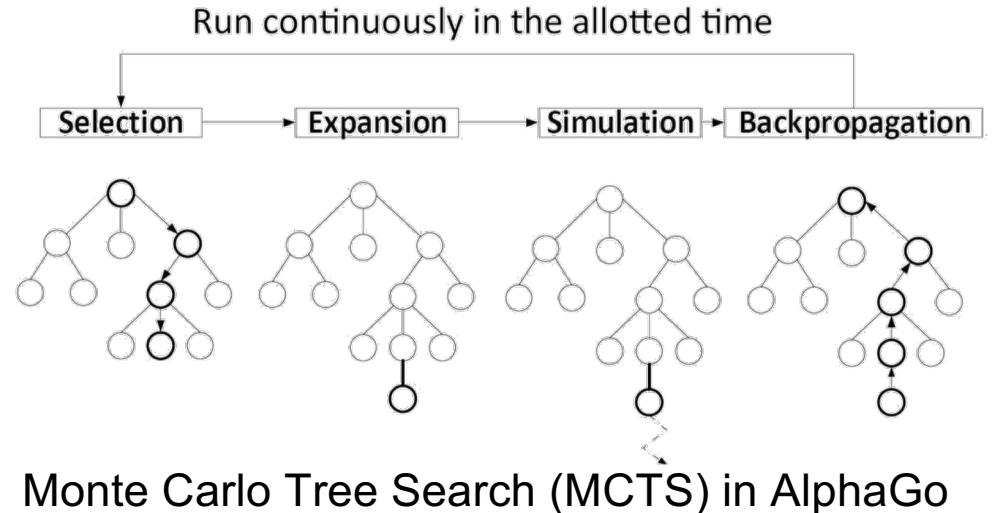
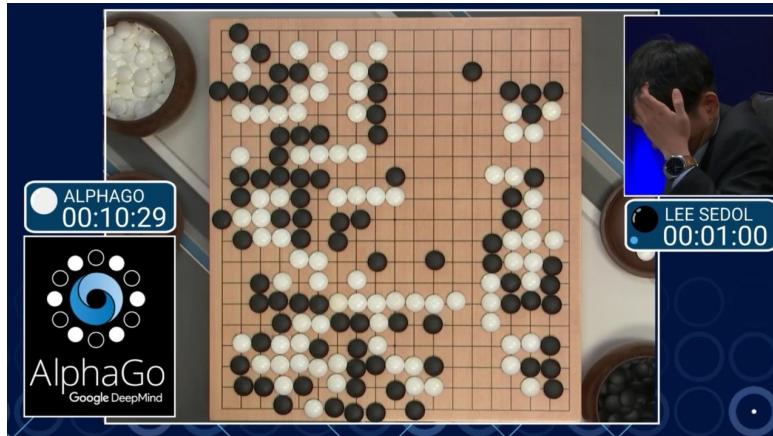


- **Analogy: Game tree (win) **VS** Execution tree (a vulnerable path)**

We can use MCTS to guide the path search towards the vulnerable paths!



Insight 2: search the path smartly



- **Analogy: Game tree (win) **VS** Execution tree (a vulnerable path)**

We can use MCTS to guide the path search towards the vulnerable paths!

Key question 1: How to select/expand tree nodes?

Key question 2: How to evaluate rewards?



Solution: Vital (vulnerability-oriented MCTS)



School of
Computing and
Information Systems

[1] Ccured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. **CCured: type-safe retrofitting of legacy software**. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Solution: Vital (vulnerability-oriented MCTS)



School of
Computing and
Information Systems



How to select/expand tree nodes?

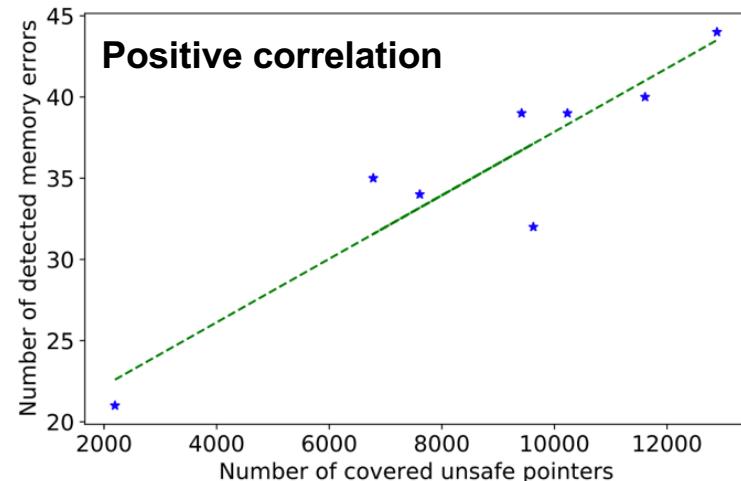
- Use the number of unsafe pointers
- Fact: vulnerabilities always happen on *type-unsafe* pointers

Solution: Vital (vulnerability-oriented MCTS)



How to select/expand tree nodes?

- Use the number of unsafe pointers
- Fact: vulnerabilities always happen on *type-unsafe* pointers



Solution: Vital (vulnerability-oriented MCTS)



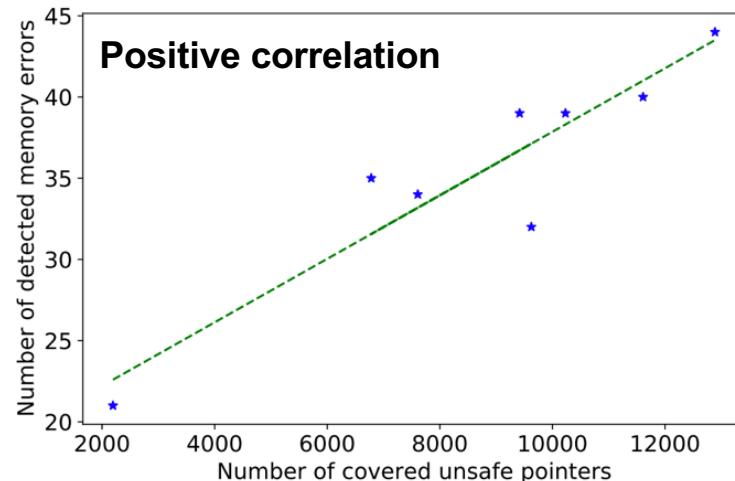
How to select/expand tree nodes?

- Use the number of unsafe pointers
- Fact: vulnerabilities always happen on *type-unsafe* pointers



How to evaluate rewards?

- Use state simulation (+backpropagation)
- Reward: number of unsafe pointers

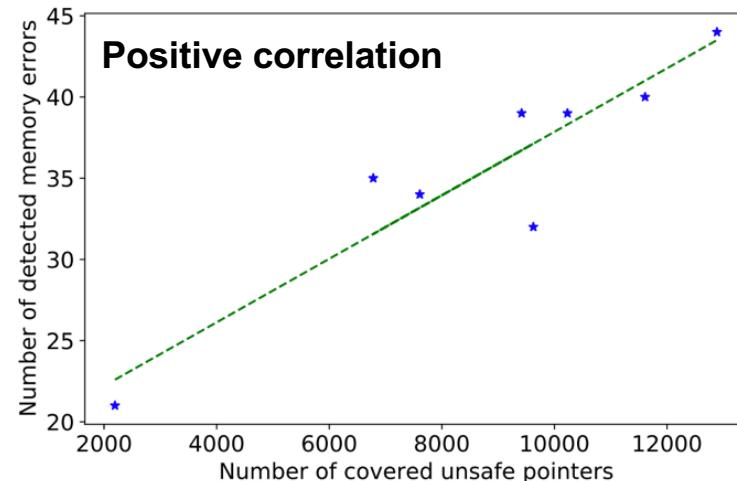


Solution: Vital (vulnerability-oriented MCTS)



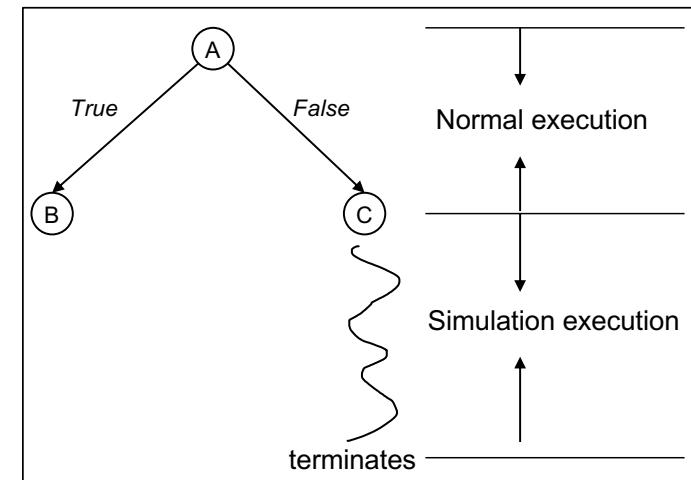
How to select/expand tree nodes?

- Use the number of unsafe pointers
- Fact: vulnerabilities always happen on type-unsafe pointers

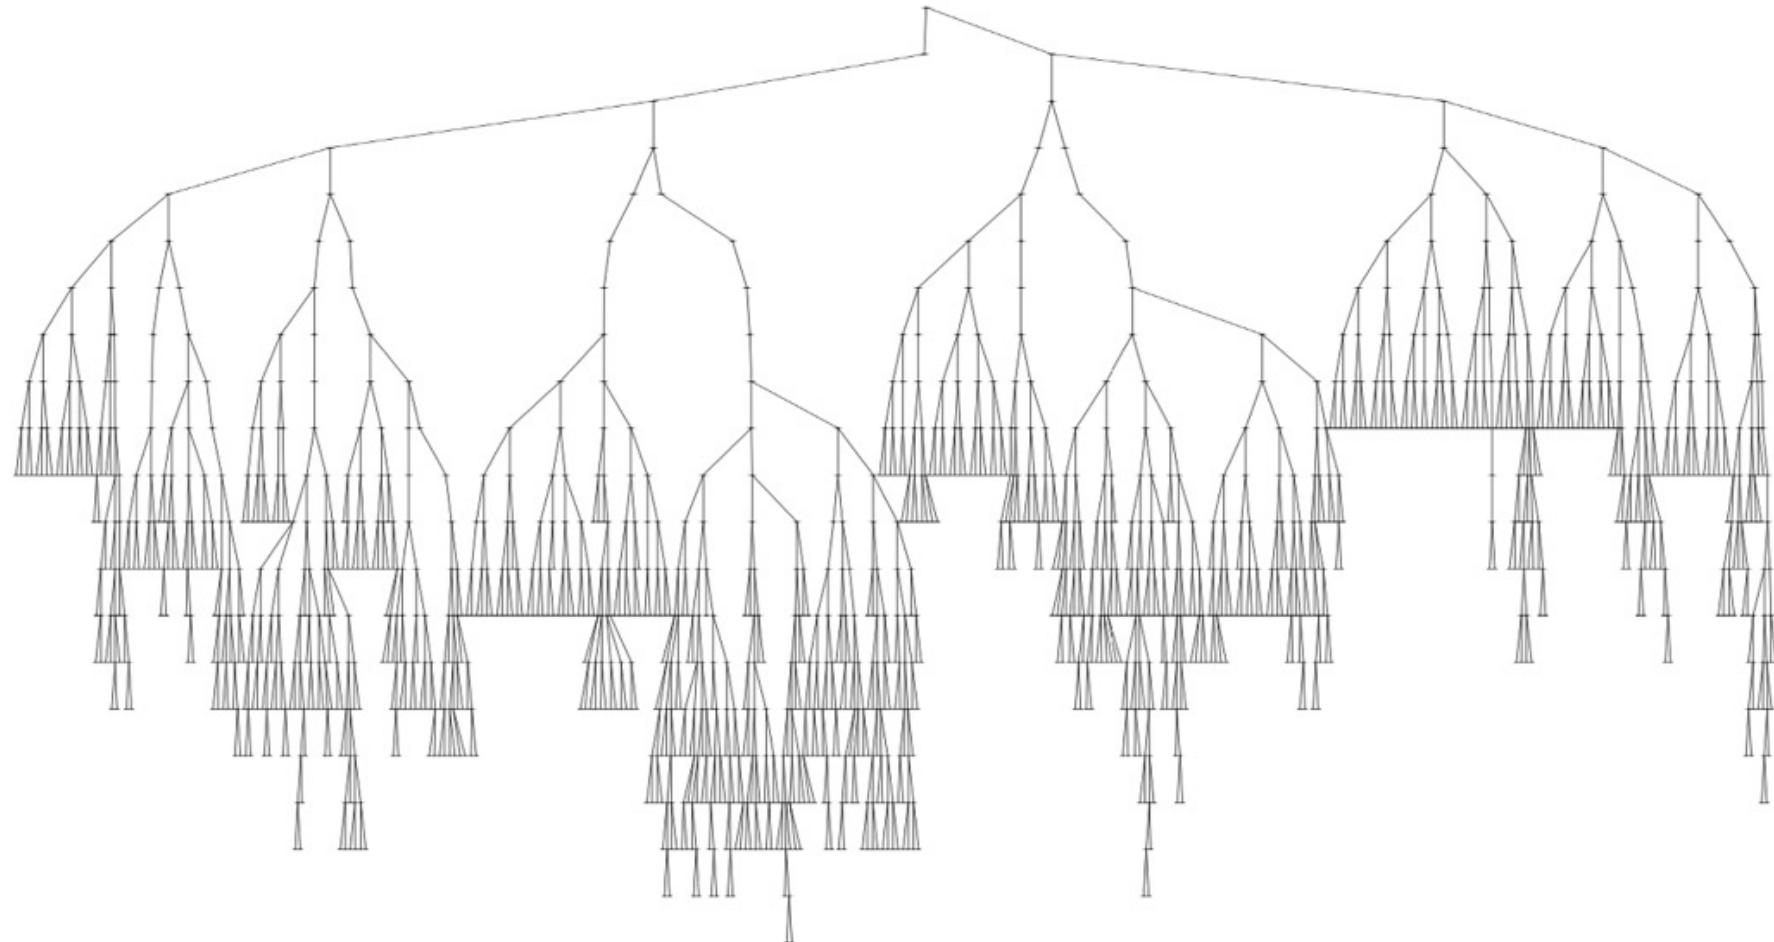


How to evaluate rewards?

- Use state simulation (+backpropagation)
- Reward: number of unsafe pointers

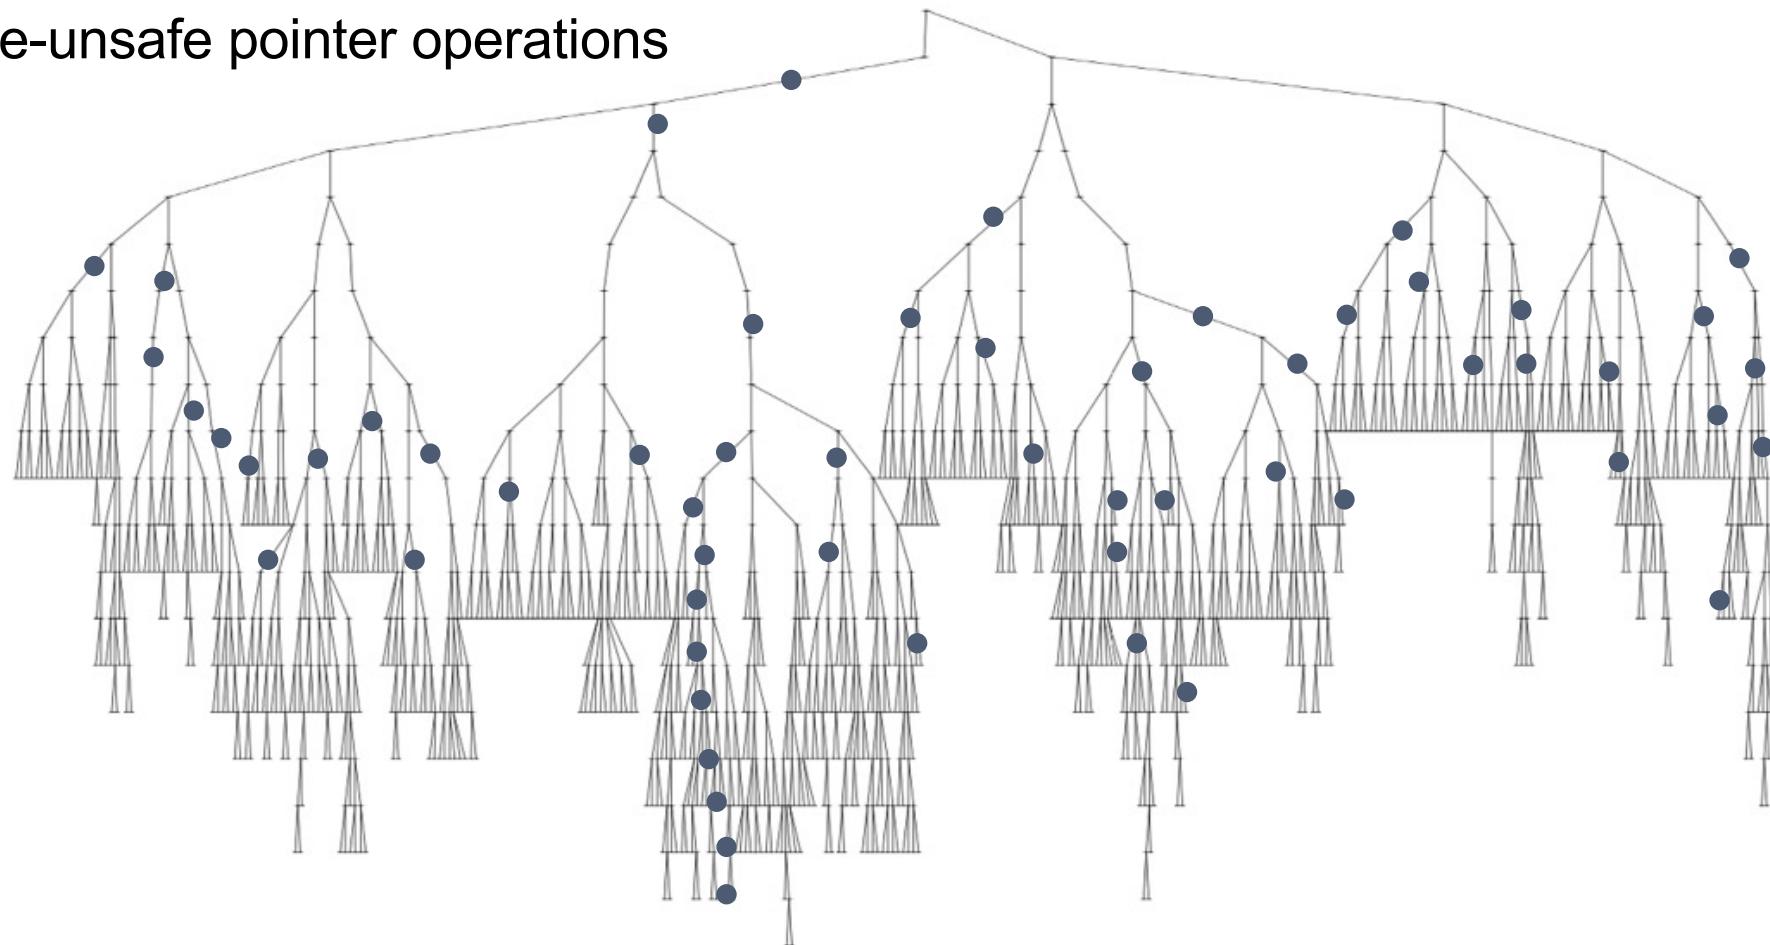


Effective Path Exploration: outcome



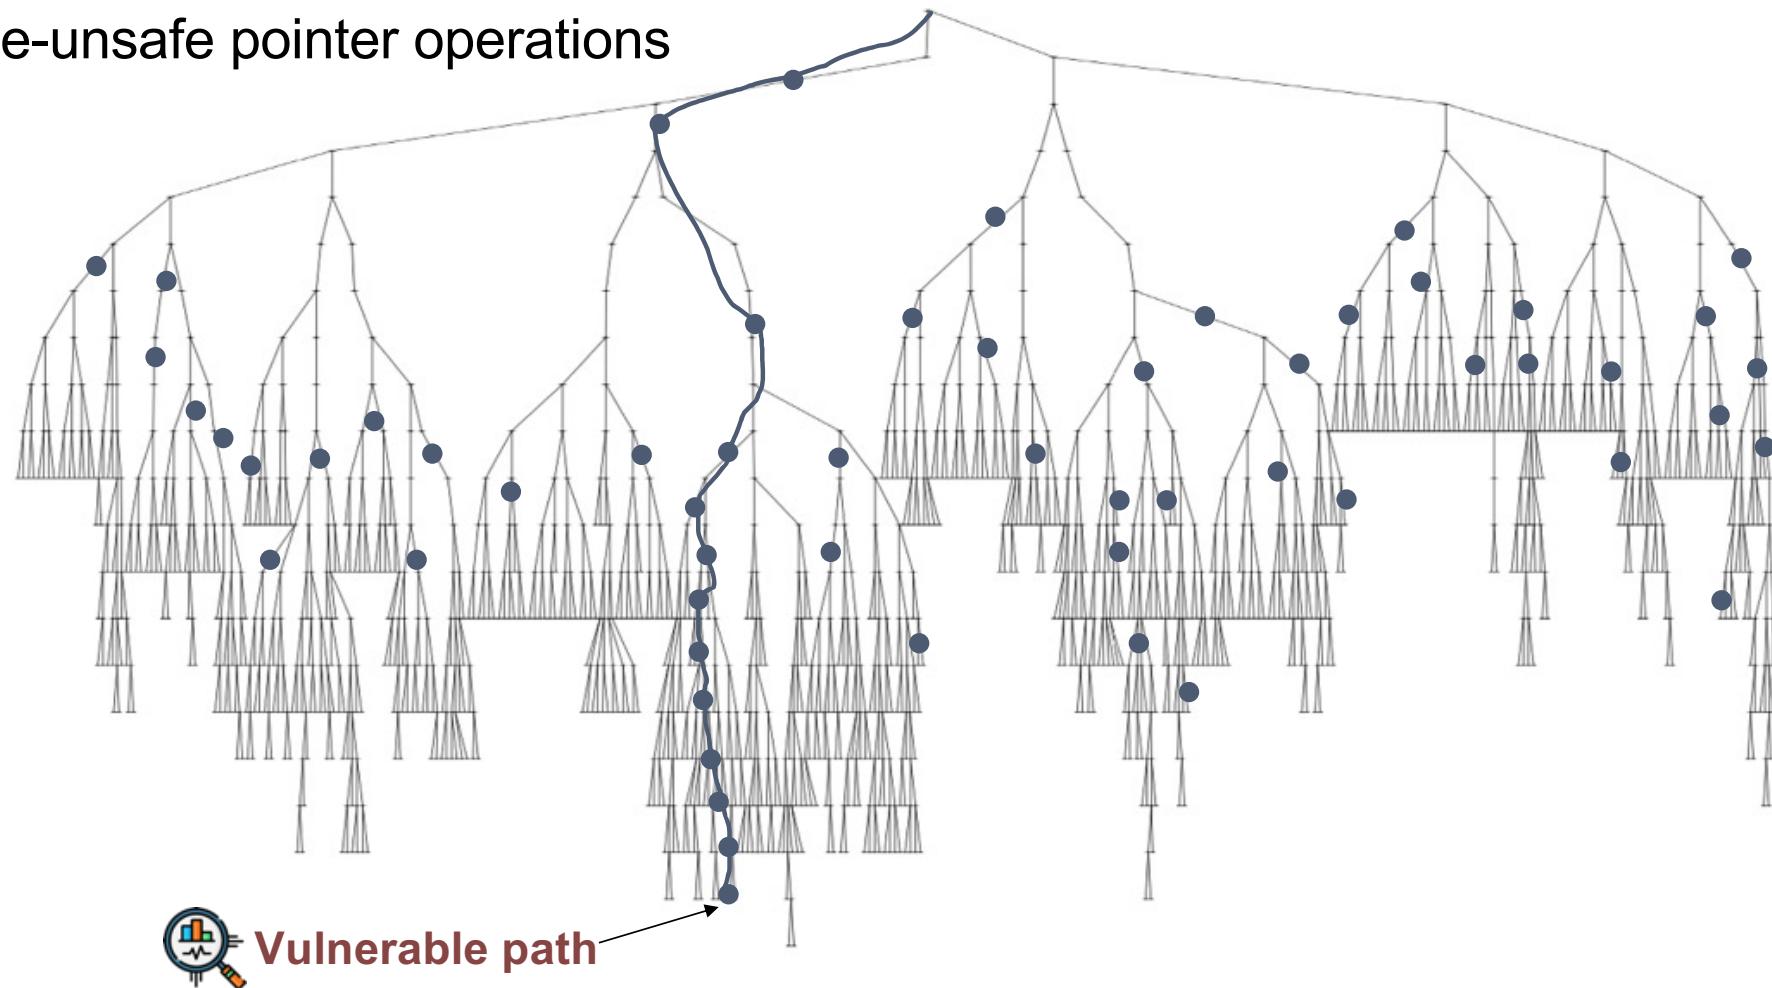
Effective Path Exploration: outcome

- Type-unsafe pointer operations



Effective Path Exploration: outcome

- Type-unsafe pointer operations

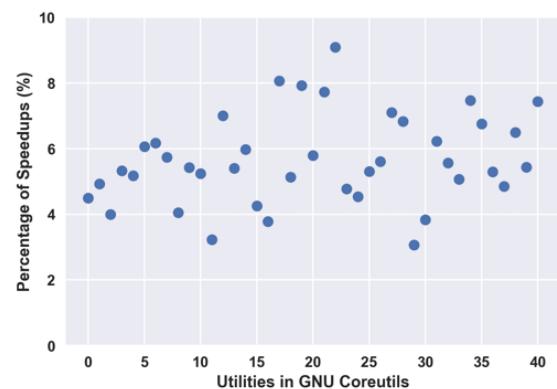


Evaluation Results & Takeaways

Evaluation Results & Takeaways

Results for Direction 1: Can we do the path search faster?

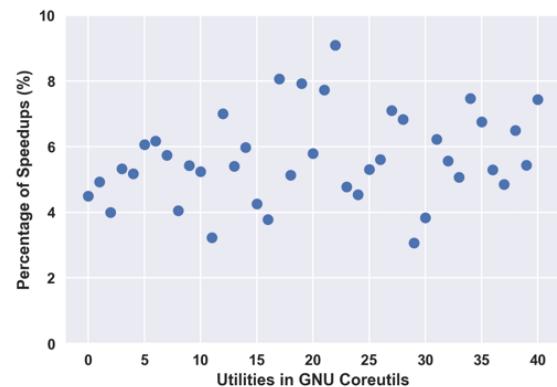
- **FastKLEE** can reduce by up to **9.1%** time compared with the state-of-the-art approach (i.e., KLEE)



Evaluation Results & Takeaways

Results for Direction 1: Can we do the path search faster?

- **FastKLEE** can reduce by up to **9.1%** time compared with the state-of-the-art approach (i.e., KLEE)



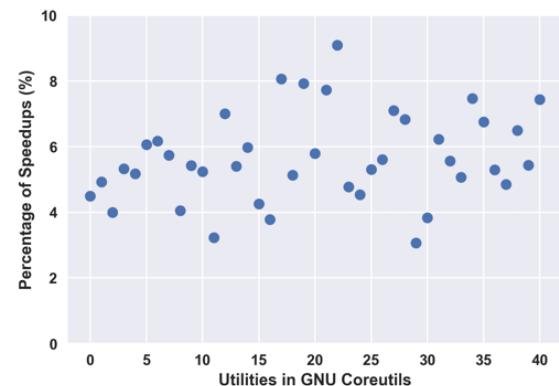
Results for Direction 2: Can we search vulnerable paths first?

- **Vital** outperforms existing search strategies by covering up to **90.03%** unsafe pointers and detecting up to **57.14%** more unique memory errors.
- **Vital** outperforms existing solutions by achieving a speedup of up to **30x** execution time and a reduction of up to **20x** memory consumption.
- Detected an unknown vulnerability (a new CVE)

Evaluation Results & Takeaways

Results for Direction 1: Can we do the path search faster?

- **FastKLEE** can reduce by up to **9.1%** time compared with the state-of-the-art approach (i.e., KLEE)



Results for Direction 2: Can we search vulnerable paths first?

- **Vital** outperforms existing search strategies by covering up to **90.03%** unsafe pointers and detecting up to **57.14%** more unique memory errors.
- **Vital** outperforms existing solutions by achieving a speedup of up to **30x** execution time and a reduction of up to **20x** memory consumption.
- Detected an unknown vulnerability (a new CVE)

• Takeaways

- Combing type inference with symbolic execution could help do a better path exploration

Outline

□ Background

- Software vulnerabilities
- Symbolic execution

□ Motivation

- What limitations prevent current symbolic execution for vulnerability detection?
- Thesis statement and research objects

□ Methodology

- **SymLoc**: A new memory model for symbolic execution
- **FastKLEE and Vital**: Two new path exploration for symbolic execution
- **Cottontail**: A new structured input generation for symbolic execution

□ Future work

- Extension of current solutions
- Combine Program Analysis with GenAI

□ Conclusion and Acknowledgment

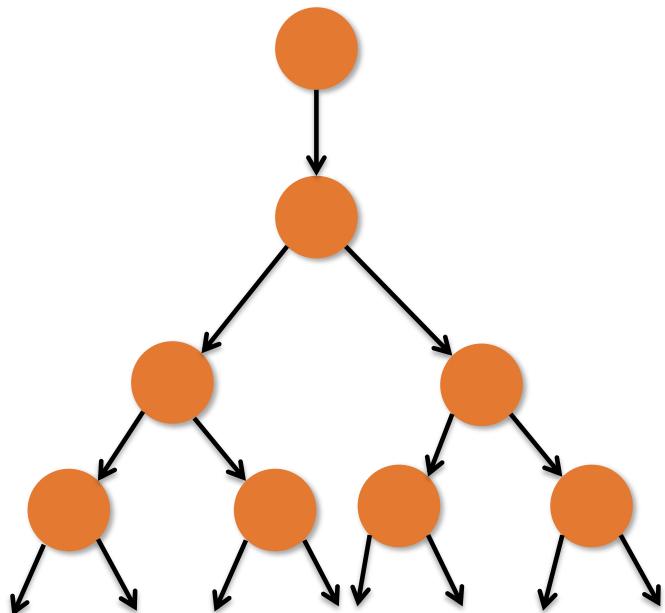
Preliminary: online and offline

- ***Online*** and ***offline*** symbolic execution

Preliminary: online and offline

- **Online** and **offline** symbolic execution

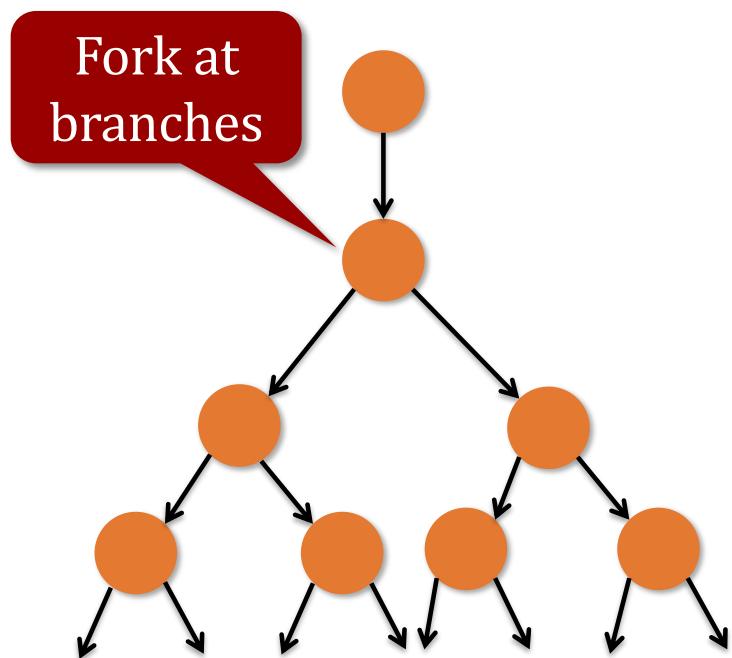
Online (without initial seeds)



Preliminary: online and offline

- **Online** and **offline** symbolic execution

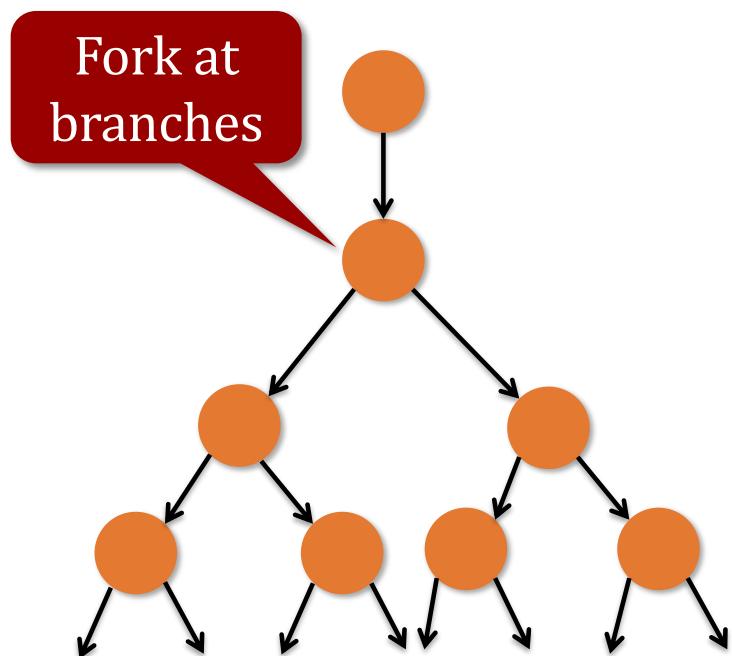
Online (without initial seeds)



Preliminary: online and offline

- **Online** and **offline** symbolic execution

Online (without initial seeds)



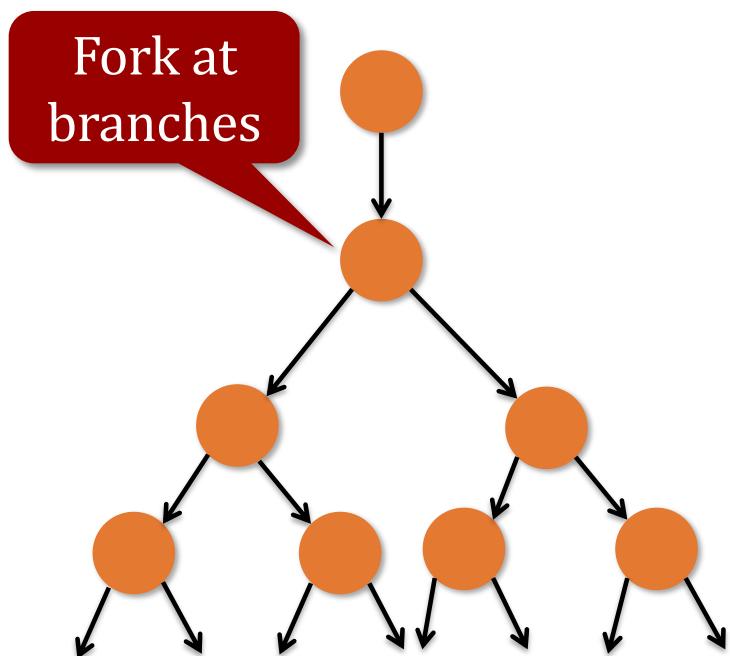
- **Example: KLEE**

Cristian Cadar, Daniel Dunbar, and Dawson Engler. 2008. "KLEE: unassisted and automatic generation of high-coverage tests for complex systems programs". In OSDI.

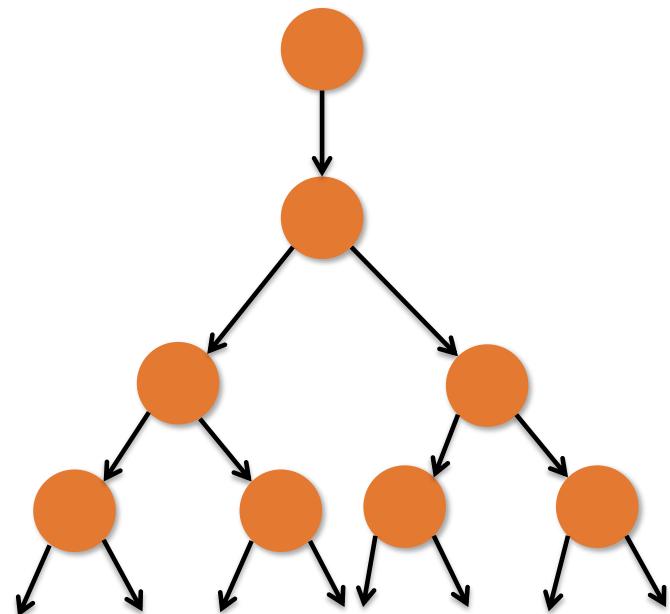
Preliminary: online and offline

- **Online** and **offline** symbolic execution

Online (without initial seeds)



Offline (require initial seeds to setup)



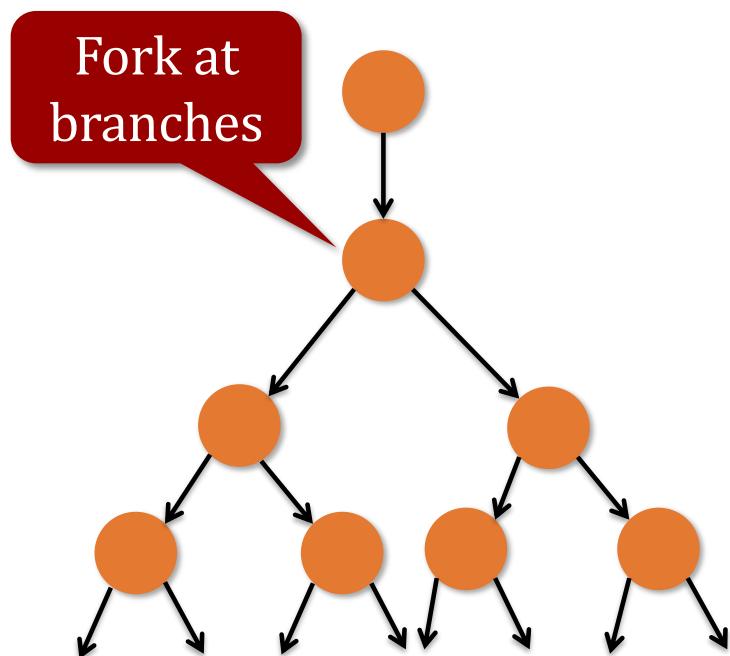
- Example: **KLEE**

Cristian Cadar, Daniel Dunbar, and Dawson Engler. 2008. "KLEE: unassisted and automatic generation of high-coverage tests for complex systems programs". In OSDI.

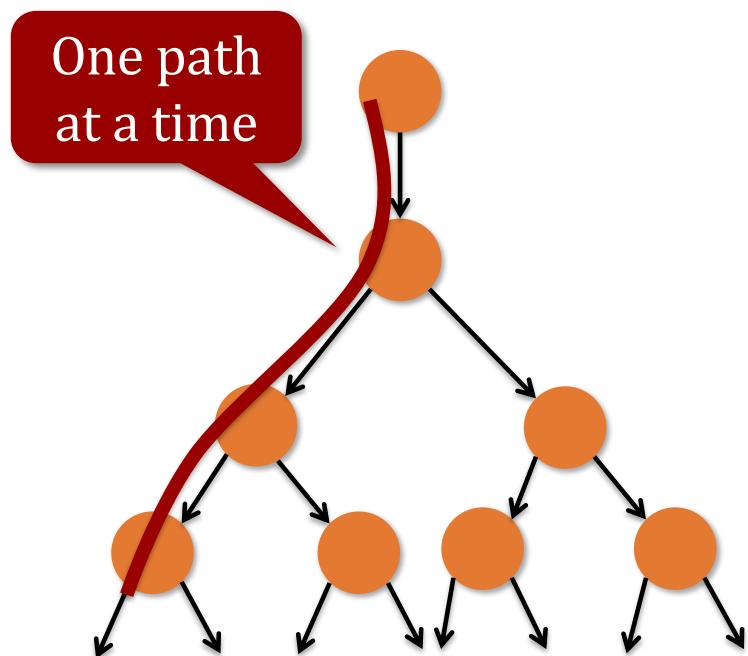
Preliminary: online and offline

- **Online** and **offline** symbolic execution

Online (without initial seeds)



Offline (require initial seeds to setup)



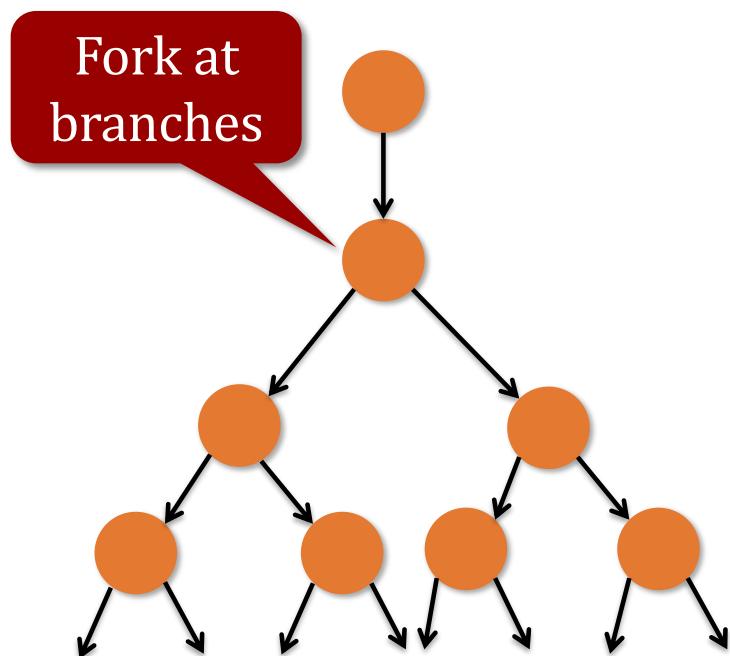
- **Example: KLEE**

Cristian Cadar, Daniel Dunbar, and Dawson Engler. 2008. "KLEE: unassisted and automatic generation of high-coverage tests for complex systems programs". In OSDI.

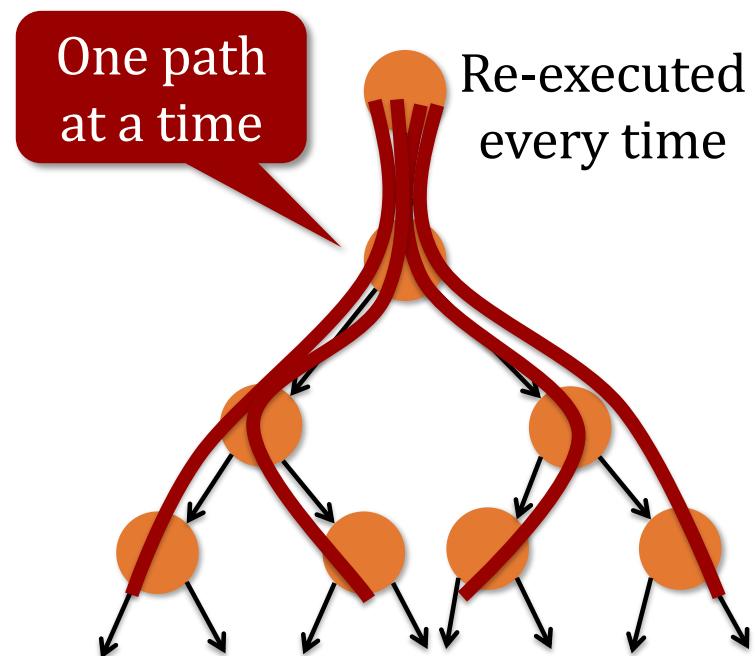
Preliminary: online and offline

- **Online** and **offline** symbolic execution

Online (without initial seeds)



Offline (require initial seeds to setup)



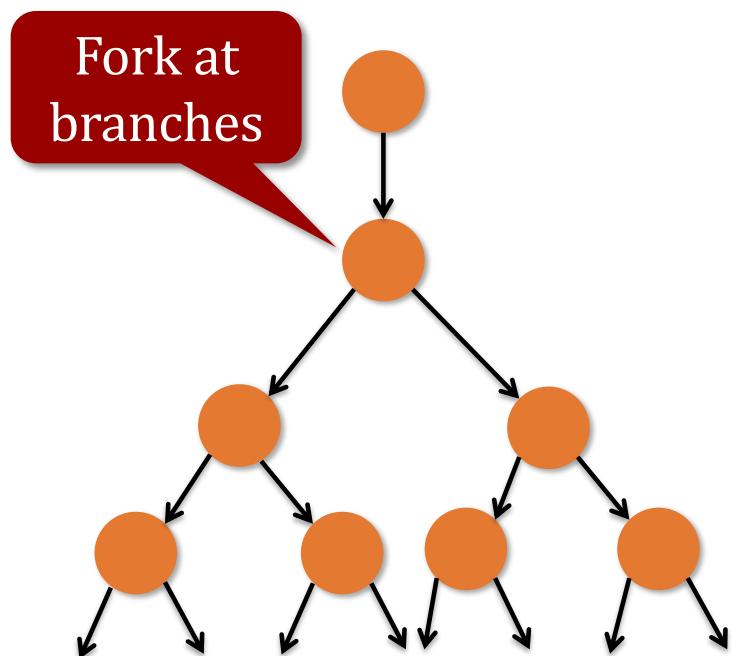
- **Example: KLEE**

Cristian Cadar, Daniel Dunbar, and Dawson Engler. 2008." KLEE: unassisted and automatic generation of high-coverage tests for complex systems programs". In OSDI.

Preliminary: online and offline

- **Online** and **offline** symbolic execution

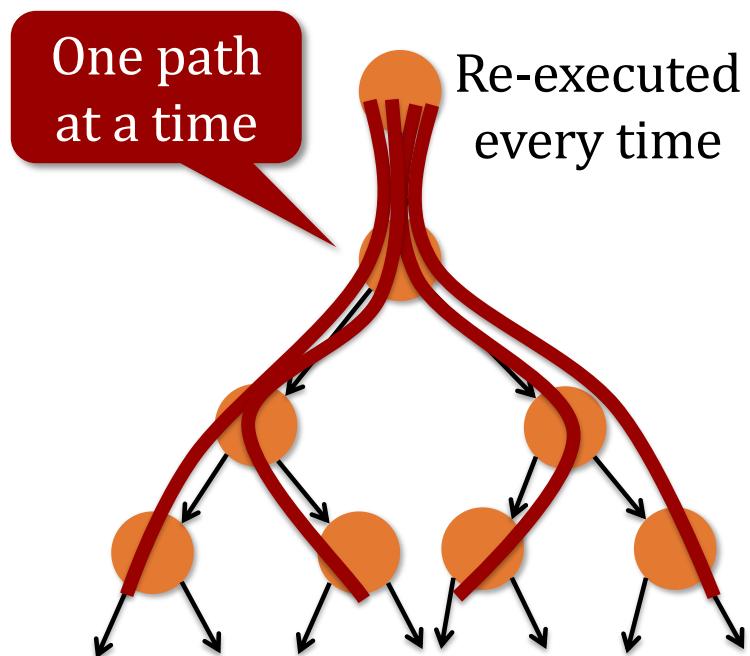
Online (without initial seeds)



- Example: **KLEE**

Cristian Cadar, Daniel Dunbar, and Dawson Engler. 2008. "KLEE: unassisted and automatic generation of high-coverage tests for complex systems programs". In OSDI.

Offline (require initial seeds to setup)



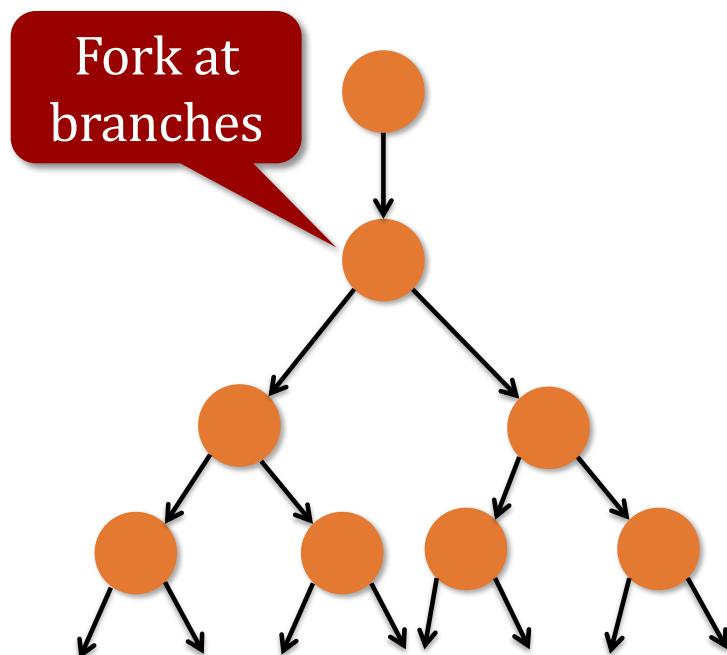
- Example: **SAGE**

Patrice Godefroid, Michael Y. Levin, and David Molnar. 2012. SAGE: whitebox fuzzing for security testing. Commun. ACM.

Preliminary: online and offline

- **Online** and **offline** symbolic execution

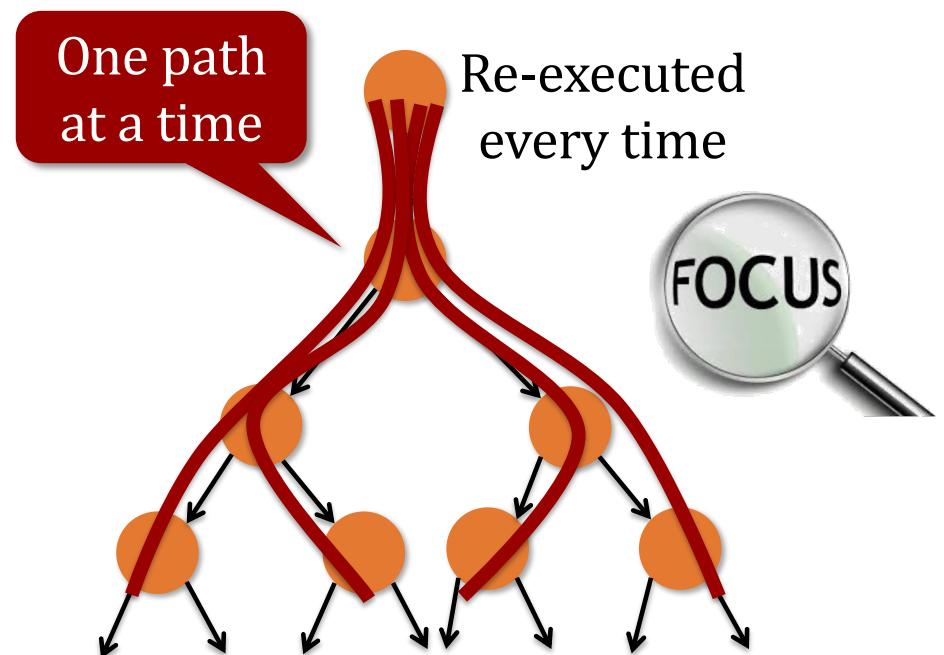
Online (without initial seeds)



- Example: **KLEE**

Cristian Cadar, Daniel Dunbar, and Dawson Engler. 2008. "KLEE: unassisted and automatic generation of high-coverage tests for complex systems programs". In OSDI.

Offline (require initial seeds to setup)



- Example: **SAGE**

Patrice Godefroid, Michael Y. Levin, and David Molnar. 2012. SAGE: whitebox fuzzing for security testing. Commun. ACM.

Preliminary: compilation-based



School of
Computing and
Information Systems

Preliminary: compilation-based

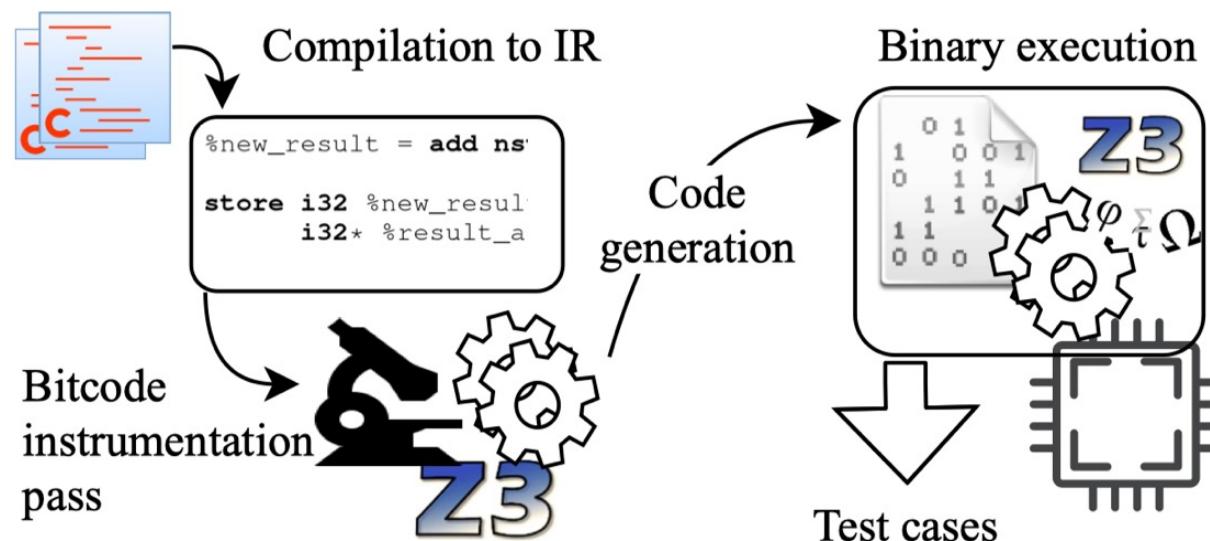
□ Prevalent concolic execution style

- Compilation-based concolic/symbolic execution: SymCC (Usenix Sec'20)
- Potentially (ideally) can be used to test any scale software systems

Preliminary: compilation-based

□ Prevalent concolic execution style

- Compilation-based concolic/symbolic execution: SymCC (Usenix Sec'20)
- Potentially (ideally) can be used to test any scale software systems



Preliminary: compilation-based

□ Prevalent concolic execution style

- Compilation-based concolic/symbolic execution: SymCC (Usenix Sec'20)
- Potentially (ideally) can be used to test any scale software systems

How can we perform concolic execution to generate highly structured test inputs for systematically testing parsing programs?



Challenges



School of
**Computing and
Information Systems**

Challenges

❑ #C1: what to solve?

- Not every path is worth solving (reduce resources/costs)
- Existing path constraint selection strategy is not structure-aware

Challenges

□ #C1: what to solve?

- Not every path is worth solving (reduce resources/costs)
- Existing path constraint selection strategy is not structure-aware

□ #C2: how to solve?

- Need to make sure resulting test cases are both satisfactory to constraints and valid to syntax
- Existing constraint solving only solve for satisfactory, which yields a large number of invalid inputs

Challenges

□ #C1: what to solve?

- Not every path is worth solving (reduce resources/costs)
- Existing path constraint selection strategy is not structure-aware

□ #C2: how to solve?

- Need to make sure resulting test cases are both satisfactory to constraints and valid to syntax
- Existing constraint solving only solve for satisfactory, which yields a large number of invalid inputs

□ #C3: How to acquire new seeds after saturation?

- For a better continuous testing
- Existing measurements (e.g., using gcov) could be time-consuming to get during runtime

Key insights



School of
**Computing and
Information Systems**

Key insights

❑ Insight 1: *unique implementation characteristic*

- Can help distinguish structural program paths
 - Select only interesting path constraints

Key insights

❑ Insight 1: *unique implementation characteristic*

- Can help distinguish structural program paths
 - Select only interesting path constraints



School of
Computing and
Information Systems

```
1 // seed input
2 c = 30000; a = [];
3 for (i = 0; i < 2 * c; i += 1) { a.push(i%c); }
4 a.sort(function (x, y) { return x - y; });
print(a[2 * c - 2]);
```

Seed

Input string

```
1 // Parsing logic /* adapted from jslex.c */
2 static int jsY_isidentifierpart(int c) {
3     return isdigit(c) //"(bvsge #x00000039 (concat #x000000 k!95))"
4     || isalpha(c) || c == '$' || c == '_' || isalphanumeric(c);
5 }
6 static int jsY_lexx(js_State *J){
7     while (1) {
8         //...
9         switch (J->lexchar) {
10            case '(': jsY_next(J); return '(';
11            case ')': jsY_next(J); return ')';
12            case ',': jsY_next(J); return ',';
13            //...
14        }
15    //...
16 }
17 }
```

MuJS implementation

Key insights

❑ Insight 1: *unique implementation characteristic*

- Can help distinguish structural program paths
 - Select only interesting path constraints

❑ Insight 2: strong completion capabilities of LLMs

- Can help solve constraints smartly (more details later)
 - Solve – Complete paradigm (our key contribution)



School of
Computing and
Information Systems

```
1 // seed input
2 c = 30000; a = [];
3 for (i = 0; i < 2 * c; i += 1) { a.push(i%c); }
4 a.sort(function (x, y) { return x - y; });
print(a[2 * c - 2]);
```

Seed

Input string

```
1 // Parsing logic /* adapted from jslex.c */
2 static int jsY_isidentifierpart(int c) {
3     return isdigit(c) //"(bvsge #x00000039 (concat #x000000 k!95))"
4     || isalpha(c) || c == '$' || c == '_' || isalpharune(c);
5 }
6 static int jsY_lexx(js_State *J){
7     while (1) {
8         /**
9         switch (J->lexchar) {
10             case '(': jsY_next(J); return '(';
11             case ')': jsY_next(J); return ')';
12             case ',': jsY_next(J); return ',';
13             /**
14         }
15     /**
16 }
17 }
```

MuJS implementation

Key insights

❑ Insight 1: *unique implementation characteristic*

- Can help distinguish structural program paths
 - Select only interesting path constraints

❑ Insight 2: strong completion capabilities of LLMs

- Can help solve constraints smartly (more details later)
 - Solve – Complete paradigm (our key contribution)



School of
Computing and
Information Systems

```
1 // seed input
2 c = 30000; a = [];
3 for (i = 0; i < 2 * c; i += 1) { a.push(i%c); }
4 a.sort(function (x, y) { return x - y; });
print(a[2 * c - 2]);
```

Seed

Input string

```
1 // Parsing logic /* adapted from jslex.c */
2 static int jsY_isidentifierpart(int c) {
3     return isdigit(c) //"(bvsge #x00000039 (concat #x000000 k!95))"
4     || isalpha(c) || c == '$' || c == '_' || isalpharune(c);
5 }
6 static int jsY_lexx(js_State *J){
7     while (1) {
8         /**
9         switch (J->lexchar) {
10             case '(': jsY_next(J); return '(';
11             case ')': jsY_next(J); return ')';
12             case ',': jsY_next(J); return ',';
13             /**
14         }
15     /**
16 }
17 }
```

MuJS implementation

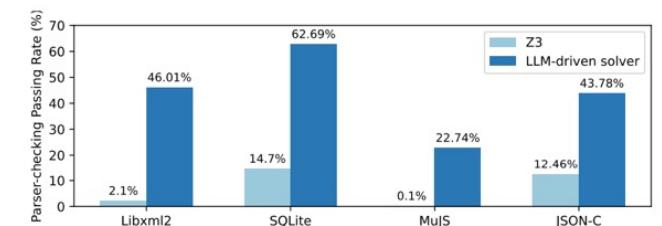


Figure 1: Parser checking passing rate comparison between traditional constraint solver (i.e., Z3) and LLM-driven solver (designed in COTTONTAIL).

Key insights

❑ Insight 1: *unique implementation characteristic*

- Can help distinguish structural program paths
 - Select only interesting path constraints

❑ Insight 2: strong completion capabilities of LLMs

- Can help solve constraints smartly (more details later)
- Solve – Complete paradigm (our key contribution)

❑ Insight 3: knowledgeable resources for seeds

- Can help generate fresh seeds when saturated



School of
Computing and
Information Systems

```
1 // seed input
2 c = 30000; a = [];
3 for (i = 0; i < 2 * c; i += 1) { a.push(i%c); }
4 a.sort(function (x, y) { return x - y; });
print(a[2 * c - 2]);
```

Seed

```
1 // Parsing logic /* adapted from jslex.c */
2 static int jsY_isidentifierpart(int c) {
3     return isdigit(c) //"(bvsge #x00000039 (concat #x000000 k!95))"
4     || isalpha(c) || c == '$' || c == '_' || isalphanumeric(c);
5 }
6 static int jsY_lexx(js_State *J){
7     while (1) {
8         ...
9         switch (J->lexchar) {
10            case '(': jsY_next(J); return '(';
11            case ')': jsY_next(J); return ')';
12            case ',': jsY_next(J); return ',';
13            ...
14        }
15    ...
16 }
17 }
```

MuJS implementation

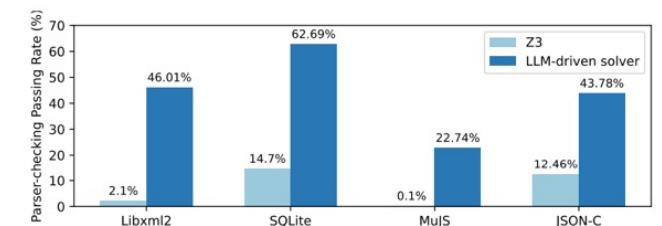
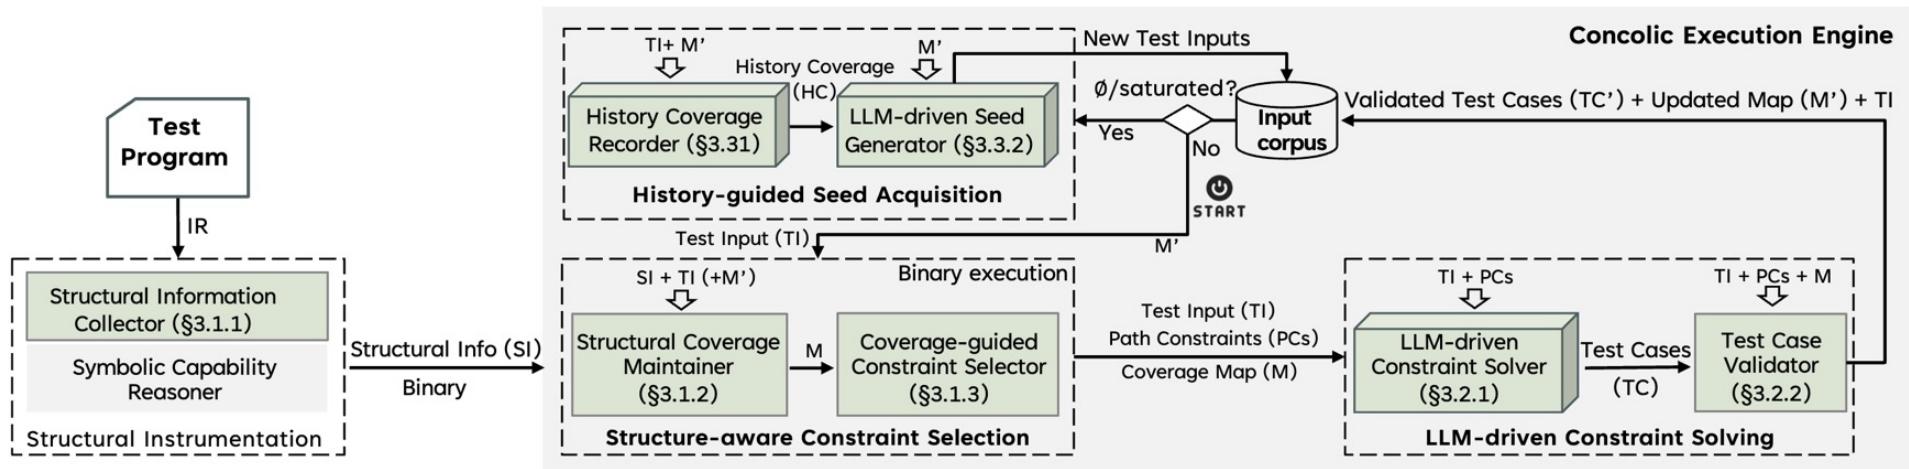
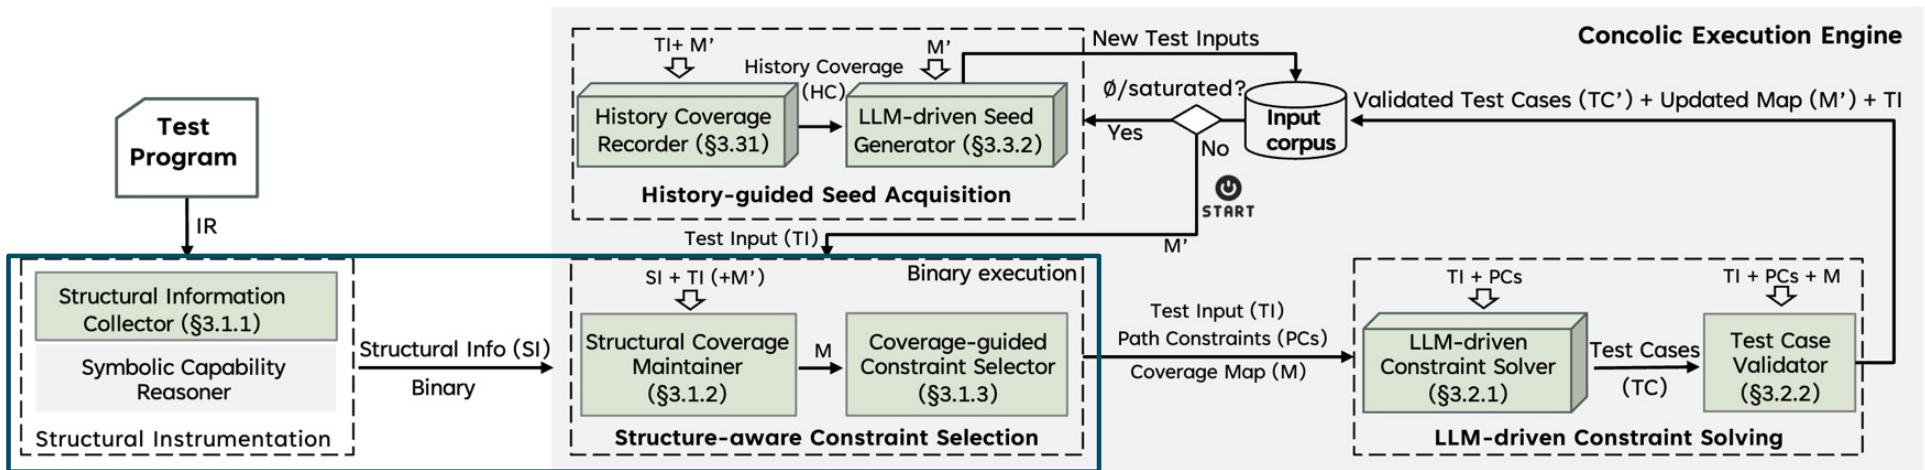


Figure 1: Parser checking passing rate comparison between traditional constraint solver (i.e., Z3) and LLM-driven solver (designed in COTTONTAIL).

Our solution: Cottontail

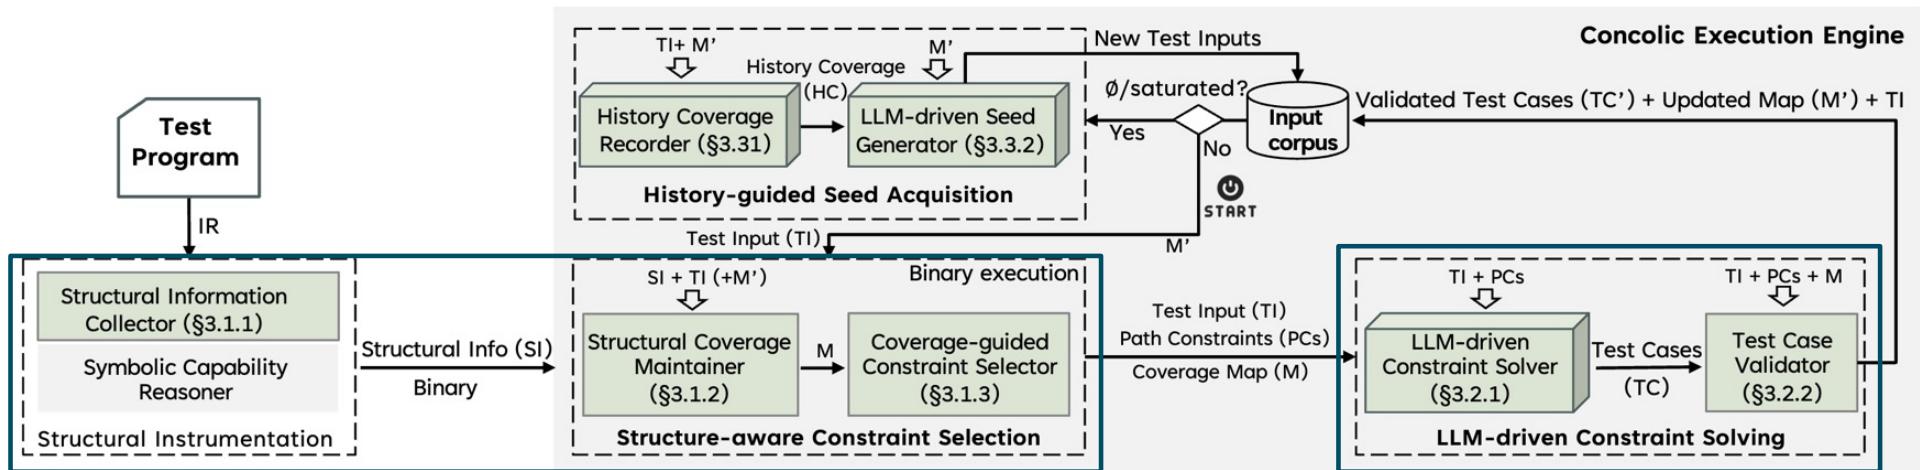


Our solution: Cottontail



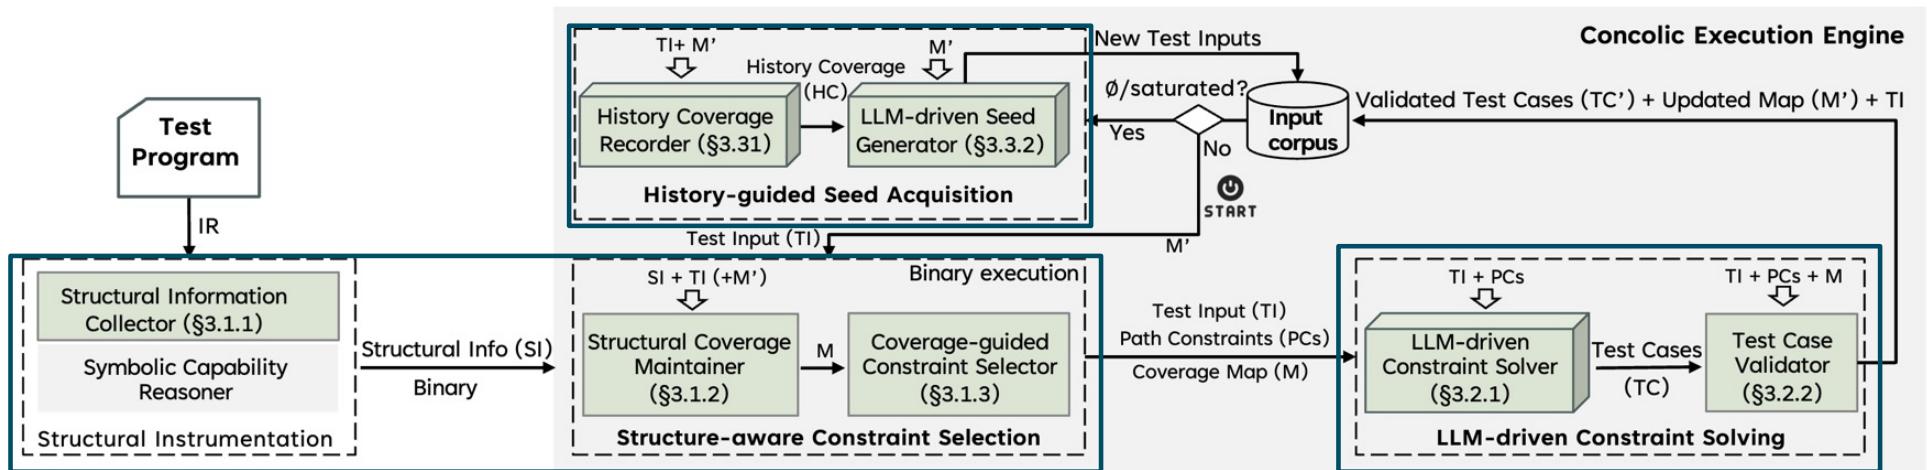
- **Three new techniques**
 - **Structure-aware constraint selection:** select only interesting path constraints (#C1)

Our solution: Cottontail



- **Three new techniques**
 - **Structure-aware constraint selection:** select only interesting path constraints (#C1)
 - **LLM-driven constraint solving:** smart LLM-solving + test case validator (#C2)

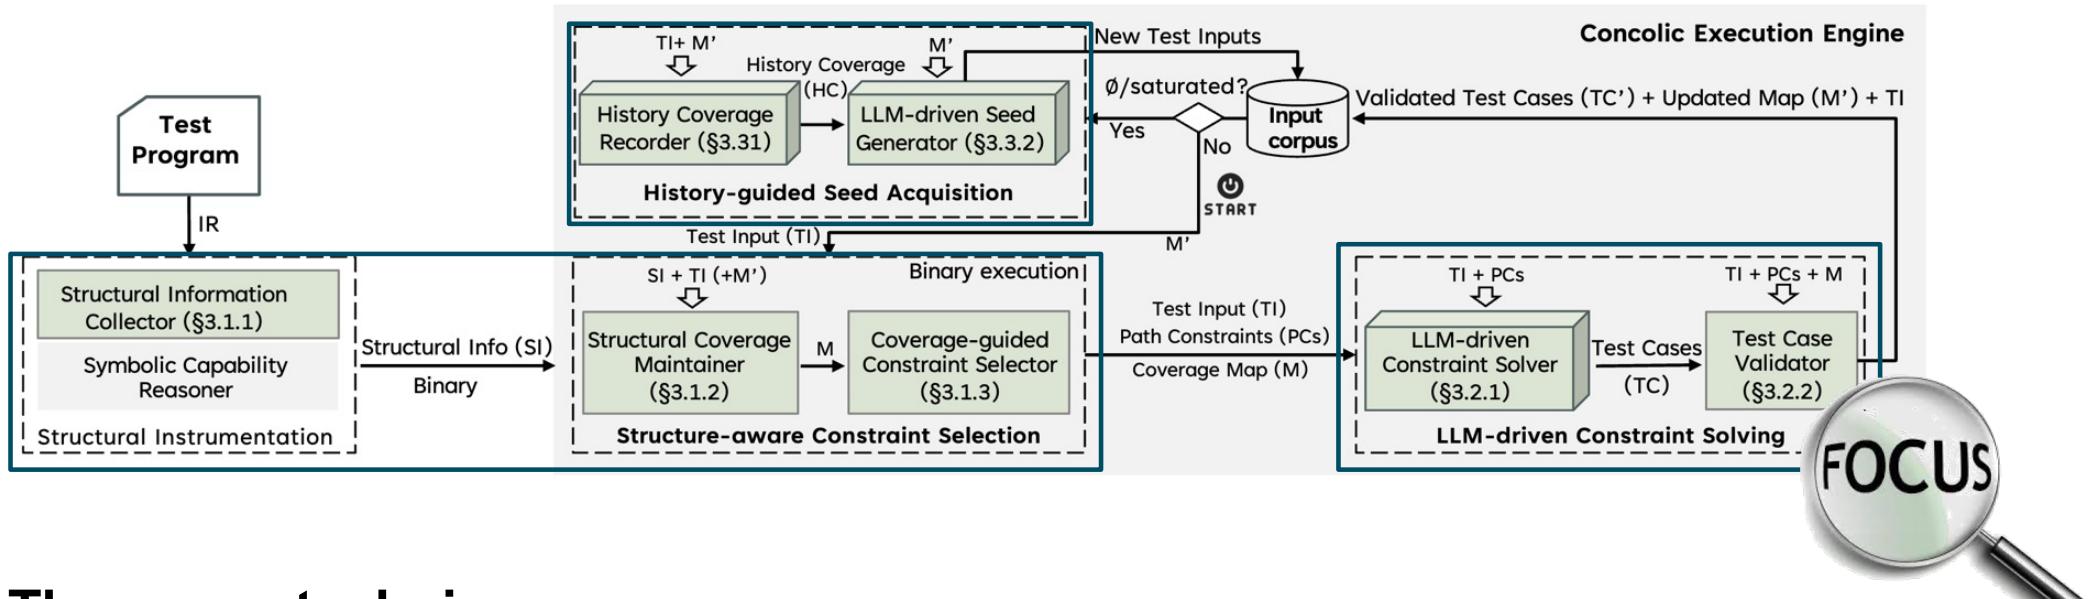
Our solution: Cottontail



- **Three new techniques**

- **Structure-aware constraint selection**: select only interesting path constraints (#C1)
- **LLM-driven constraint solving**: smart LLM-solving + test case validator (#C2)
- **History-guided seed acquisition**: capable seed generation (#C3)

Our solution: Cottontail



- **Three new techniques**

- **Structure-aware constraint selection:** select only interesting path constraints (#C1)
- **LLM-driven constraint solving:** smart LLM-solving + test case validator (#C2)
- **History-guided seed acquisition:** capable seed generation (#C3)

Novelty: Solve-Complete paradigm

</> System Prompt (*Solve-Complete* Paradigm with Chain of Thought)

You are a smart Constraint Solver.



Given a constraint and a marked test input string, please generate new test case step by step:

1. Identify the location of **Constraint Mask** ([k!n]) and *solve* it with the solution of the constraint:
 - The solution is not arbitrarily chosen; instead, make the combined string align with input syntax, such as keywords, defined functions, or variable names to ensure the resulting string remains valid.
2. Identify the location of **Flexible Mask** ([xxx]) and *complete* it with a **flexible size** string:
 - The replacement is not randomly selected: instead, follow the syntax of JavaScript and make the whole string remaining a valid JavaScript code.
3. Ensure final output is enclosed with `

</> User Prompt

Given the path constraint ```bvsge #x00000039 (concat #x000000 k!95)``` and the marked seed input you need to change: ```c = 30000; a = []; for (i = 0; i < 2 * c; i += 1) {a.push(i%c);} a.sort(function (x, y) { r[k!95][xxx] ```, what is the new test input?

</> Response

Let's solve the constraint step by step:

1. The constraint mask is [k!95] and we now solve it based on the given constraint.
 - The constraint suggests the value should be large or equal to ASCII 0x39. Considering the previous string 'r', the solution 'e' should satisfy the requirements, as 'return' is a keyword in JavaScript.
2. The flexible mask is [xxx] and we now complete it with other flexible size strings:
 - Since previous string starting with 're' after solving, complete '[xxx]' with 'turn y;}' print(a[100]);', as 'return' is a keyword, and 'y' and 'a' are defined variables.
3. Final output:

```
```c = 30000; a = []; for (i = 0; i < 2 * c; i += 1) {a.push(i%c);} a.sort(function (x, y) {return y;}  
print(a[100]);````
```



# Novelty: Solve-Complete paradigm

## </> System Prompt (*Solve-Complete* Paradigm with Chain of Thought)

You are a smart Constraint Solver.



Given a constraint and a marked test input string, please generate new test case step by step:

1. Identify the location of **Constraint Mask** ([k!n]) and *solve* it with the solution of the constraint:
  - The solution is not arbitrarily chosen; instead, make the combined string align with input syntax, such as keywords, defined functions, or variable names to ensure the resulting string remains valid.
2. Identify the location of **Flexible Mask** ([xxx]) and *complete* it with a **flexible size** string:
  - The replacement is not randomly selected: instead, follow the syntax of JavaScript and make the whole string remaining a valid JavaScript code.
3. Ensure final output is enclosed with `

## </> User Prompt

Given the path constraint `bvsge #x00000039 (concat #x000000 k!95)` and the marked seed input you need to change: `c = 30000; a = []; for (i = 0; i < 2 \* c; i += 1) {a.push(i%c);} a.sort(function (x, y) { r?`  
{ r[k!95][xxx]}`, what is the new test input?

## </> Response

Let's solve the constraint step by step:

- </>
1. The constraint mask is [k!95] and we now solve it based on the given constraint.
    - The constraint suggests the value should be large or equal to ASCII 0x39. Considering the previous string 'r', the solution 'e' should satisfy the requirements, as 'return' is a keyword in JavaScript.
  2. The flexible mask is [xxx] and we now complete it with other flexible size strings:
    - Since previous string starting with 're' after solving, complete '[xxx]' with 'turn y;}' print(a[100]);' as 'return' is a keyword, and 'y' and 'a' are defined variables.
  3. Final output:

```
```c = 30000; a = []; for (i = 0; i < 2 * c; i += 1) {a.push(i%c);} a.sort(function (x, y) {return y;}  
print(a[100]);```
```

c = 30000; a = []; for (i = 0; i < 2 * c; i += 1)
{a.push(i%c);} a.sort(function (x, y) { r?
`

Novelty: Solve-Complete paradigm

</> System Prompt (*Solve-Complete* Paradigm with Chain of Thought)

You are a smart Constraint Solver.

Given a constraint and a marked test input string, please generate new test case step by step:

1. Identify the location of **Constraint Mask** ([k!n]) and *solve* it with the solution of the constraint:
 - The solution is not arbitrarily chosen; instead, make the combined string align with input syntax, such as keywords, defined functions, or variable names to ensure the resulting string remains valid.
2. Identify the location of **Flexible Mask** ([xxx]) and *complete* it with a **flexible** size string:
 - The replacement is not randomly selected: instead, follow the syntax of JavaScript and make the whole string remaining a valid JavaScript code.
3. Ensure final output is enclosed with ` `.

</> User Prompt

Given the path constraint ``bvsge #x00000039 (concat #x000000 k!95))`` and the marked seed input you need to change: ``c = 30000; a = []; for (i = 0; i < 2 * c; i += 1) {a.push(i%c);} a.sort(function (x, y) { r? [k!95][xxx] }, what is the new test input?

</> Response

Let's solve the constraint step by step:

1. The constraint mask is [k!95] and we now solve it based on the given constraint.
 - The constraint suggests the value should be large or equal to ASCII 0x39. Considering the previous string 'r', the solution 'e' should satisfy the requirements, as 'return' is a keyword in JavaScript.
2. The flexible mask is [xxx] and we now complete it with other flexible size strings:
 - Since previous string starting with 're' after solving, complete '[xxx]' with 'turn y; print(a[100]);', as 'return' is a keyword, and 'y' and 'a' are defined variables.
3. Final output:

```c = 30000; a = []; for (i = 0; i < 2 \* c; i += 1) {a.push(i%c);} a.sort(function (x, y) {return y;};  
print(a[100]);````

c = 30000; a = []; for (i = 0; i < 2 \* c; i += 1)  
{a.push(i%c);} a.sort(function (x, y) { r?

Step 1: syntax-aware *solve* process

c = 30000; a = []; for (i = 0; i < 2 \* c; i += 1)  
{a.push(i%c);} a.sort(function (x, y) {re

# Novelty: Solve-Complete paradigm

</> System Prompt (*Solve-Complete* Paradigm with Chain of Thought)

You are a smart Constraint Solver.



Given a constraint and a marked test input string, please generate new test case step by step:

1. Identify the location of **Constraint Mask** ([k!n]) and *solve* it with the solution of the constraint:
  - The solution is not arbitrarily chosen; instead, make the combined string align with input syntax, such as keywords, defined functions, or variable names to ensure the resulting string remains valid.
2. Identify the location of **Flexible Mask** ([xxx]) and *complete* it with a **flexible** size string:
  - The replacement is not randomly selected: instead, follow the syntax of JavaScript and make the whole string remaining a valid JavaScript code.
3. Ensure final output is enclosed with ` ``.

</> User Prompt

Tradition solver (Z3): r9

Given the path constraint ``bvsge #x00000039 (concat #x000000 k!95))`` and the marked seed input you need to change: ``c = 30000; a = []; for (i = 0; i < 2 \* c; i += 1) {a.push(i%c);} a.sort(function (x, y) { r[k!95][xxx] ``, what is the new test input?

</> Response

Let's solve the constraint step by step:

1. The constraint mask is [k!95] and we now solve it based on the given constraint.
  - The constraint suggests the value should be large or equal to ASCII 0x39. Considering the previous string 'r', the solution 'e' should satisfy the requirements, as 'return' is a keyword in JavaScript.
2. The flexible mask is [xxx] and we now complete it with other flexible size strings:
  - Since previous string starting with 're' after solving, complete '[xxx]' with 'turn y; print(a[100]);' as 'return' is a keyword, and 'y' and 'a' are defined variables.
3. Final output:

```c = 30000; a = []; for (i = 0; i < 2 \* c; i += 1) {a.push(i%c);} a.sort(function (x, y) {return y;};  
print(a[100]);````

c = 30000; a = []; for (i = 0; i < 2 * c; i += 1)
{a.push(i%c);} a.sort(function (x, y) { r?

Step 1: syntax-aware *solve* process

c = 30000; a = []; for (i = 0; i < 2 * c; i += 1)
{a.push(i%c);} a.sort(function (x, y) {re

Novelty: Solve-Complete paradigm



</> System Prompt (*Solve-Complete* Paradigm with Chain of Thought)

You are a smart Constraint Solver.

Given a constraint and a marked test input string, please generate new test case step by step:

1. Identify the location of **Constraint Mask** ([k!n]) and *solve* it with the solution of the constraint:
 - The solution is not arbitrarily chosen; instead, make the combined string align with input syntax, such as keywords, defined functions, or variable names to ensure the resulting string remains valid.
2. Identify the location of **Flexible Mask** ([xxx]) and *complete* it with a **flexible** size string:
 - The replacement is not randomly selected: instead, follow the syntax of JavaScript and make the whole string remaining a valid JavaScript code.
3. Ensure final output is enclosed with ` ``.

</> User Prompt

Tradition solver (Z3): r9

Given the path constraint ``bvsge #x00000039 (concat #x000000 k!95))`` and the marked seed input you need to change: ``c = 30000; a = []; for (i = 0; i < 2 * c; i += 1) {a.push(i%c);} a.sort(function (x, y) { r[k!95][xxx] ``, what is the new test input?

</> Response

Let's solve the constraint step by step:

1. The constraint mask is [k!95] and we now solve it based on the given constraint.
 - The constraint suggests the value should be large or equal to ASCII 0x39. Considering the previous string 'r', the solution 'e' should satisfy the requirements, as 'return' is a keyword in JavaScript.
2. The flexible mask is [xxx] and we now complete it with other flexible size strings:
 - Since previous string starting with 're' after solving, complete '[xxx]' with 'turn y;}' print(a[100]);', as 'return' is a keyword, and 'y' and 'a' are defined variables.
3. Final output:

```
```c = 30000; a = []; for (i = 0; i < 2 * c; i += 1) {a.push(i%c);} a.sort(function (x, y) {return y;}\nprint(a[100]);````
```

c = 30000; a = []; for (i = 0; i < 2 \* c; i += 1)  
{a.push(i%c);} a.sort(function (x, y) { r?

Step 1: syntax-aware *solve* process

c = 30000; a = []; for (i = 0; i < 2 \* c; i += 1)  
{a.push(i%c);} a.sort(function (x, y) {re

Step 2: syntax-aware *complete* process

c = 30000; a = []; for (i = 0; i < 2 \* c; i += 1)  
{a.push(i%c);} a.sort(function (x, y) {return y;}\nprint(a[100]);````

# Novelty: Solve-Complete paradigm



</> System Prompt (*Solve-Complete* Paradigm with Chain of Thought)

You are a smart Constraint Solver.

Given a constraint and a marked test input string, please generate new test case step by step:

1. Identify the location of **Constraint Mask** ([k!n]) and *solve* it with the solution of the constraint:
  - The solution is not arbitrarily chosen; instead, make the combined string align with input syntax, such as keywords, defined functions, or variable names to ensure the resulting string remains valid.
2. Identify the location of **Flexible Mask** ([xxx]) and *complete* it with a **flexible** size string:
  - The replacement is not randomly selected: instead, follow the syntax of JavaScript and make the whole string remaining a valid JavaScript code.
3. Ensure final output is enclosed with ` ``.

</> User Prompt

Tradition solver (Z3): r9

Given the path constraint ``bvsge #x00000039 (concat #x000000 k!95))`` and the marked seed input you need to change: ``c = 30000; a = []; for (i = 0; i < 2 \* c; i += 1) {a.push(i%c);} a.sort(function (x, y) { r[k!95][xxx] ``, what is the new test input?

</> Response

Let's solve the constraint step by step:

1. The constraint mask is [k!95] and we now solve it based on the given constraint.
  - The constraint suggests the value should be large or equal to ASCII 0x39. Considering the previous string 'r', the solution 'e' should satisfy the requirements, as 'return' is a keyword in JavaScript.
2. The flexible mask is [xxx] and we now complete it with other flexible size strings:
  - Since previous string starting with 're' after solving, complete '[xxx]' with 'turn y;` print(a[100]);` as 'return' is a keyword, and 'y' and 'a' are defined variables.
3. Final output:

```c = 30000; a = []; for (i = 0; i < 2 \* c; i += 1) {a.push(i%c);} a.sort(function (x, y) {return y;`  
print(a[100]);``

c = 30000; a = []; for (i = 0; i < 2 * c; i += 1)
{a.push(i%c);} a.sort(function (x, y) { r?

Step 1: syntax-aware *solve* process

c = 30000; a = []; for (i = 0; i < 2 * c; i += 1)
{a.push(i%c);} a.sort(function (x, y) {re

Step 2: syntax-aware *complete* process

c = 30000; a = []; for (i = 0; i < 2 * c; i += 1)
{a.push(i%c);} a.sort(function (x, y) {return y;`
print(a[100]);``

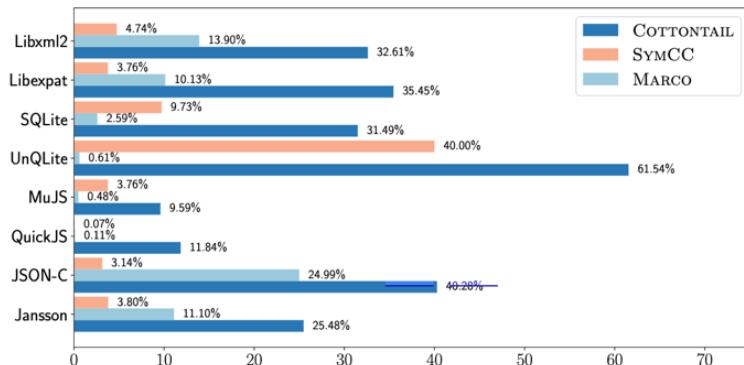
Solve and Complete for not only constraint satisfiability but also syntax validity

Evaluation Results & Takeaways

Evaluation Results & Takeaways

RQ1: How does Cottontail perform compared with baseline approaches?

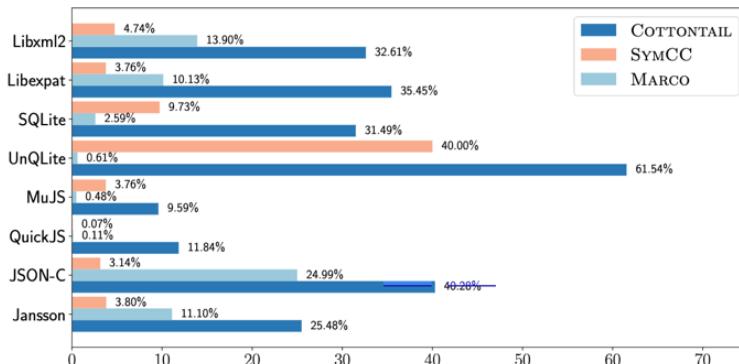
- Cottontail outperforms state-of-the-art approaches (SymCC and Marco) by **14.15%** and **14.31%** in terms of line coverage



Evaluation Results & Takeaways

RQ1: How does Cottontail perform compared with baseline approaches?

- Cottontail outperforms state-of-the-art approaches (SymCC and Marco) by **14.15%** and **14.31%** in terms of line coverage



RQ2: Can Cottontail find new vulnerabilities in practice?

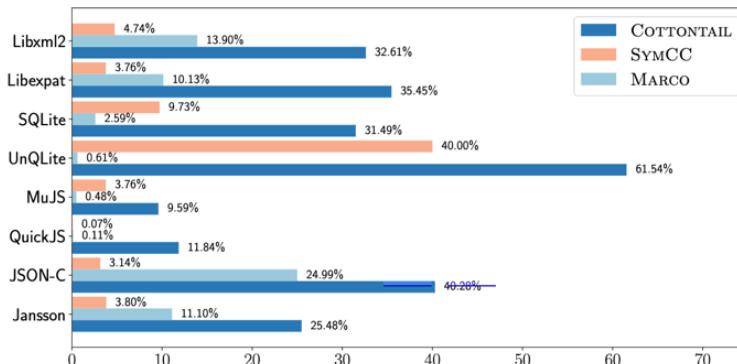
- Cottontail detected **6** new vulnerabilities, showing practical vulnerability detection capability
- Marco can only detect one and SymCC can detect four

| ID | Subject | Description | Status | CVE-Assigned |
|----|---------|-----------------|----------|----------------|
| #1 | MuJS | Memory leak | Fixed | CVE-2024-55061 |
| #2 | MuJS | Heap overflow | Fixed | CVE-2025-26082 |
| #3 | QuickJS | Stack overflow | Fixed | CVE-2024-13903 |
| #4 | QuickJS | Stack overflow | Fixed | CVE-2025-26081 |
| #5 | UnQLite | Global overflow | Reported | CVE-2025-26083 |
| #6 | UnQLite | Heap overflow | Reported | CVE-2025-3791 |

Evaluation Results & Takeaways

RQ1: How does Cottontail perform compared with baseline approaches?

- Cottontail outperforms state-of-the-art approaches (SymCC and Marco) by **14.15%** and **14.31%** in terms of line coverage



RQ2: Can Cottontail find new vulnerabilities in practice?

- Cottontail detected **6** new vulnerabilities, showing practical vulnerability detection capability
- Marco can only detect one and SymCC can detect four

| ID | Subject | Description | Status | CVE-Assigned |
|----|---------|-----------------|----------|----------------|
| #1 | MuJS | Memory leak | Fixed | CVE-2024-55061 |
| #2 | MuJS | Heap overflow | Fixed | CVE-2025-26082 |
| #3 | QuickJS | Stack overflow | Fixed | CVE-2024-13903 |
| #4 | QuickJS | Stack overflow | Fixed | CVE-2025-26081 |
| #5 | UnQLite | Global overflow | Reported | CVE-2025-26083 |
| #6 | UnQLite | Heap overflow | Reported | CVE-2025-3791 |

• Takeaways

- Smart LLM-driven constraint solving could help generate more complex test cases

Retrospection of Motivation Example

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```

Retrospection of Motivation Example

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```



□ SymLoc: new memory modeling

Retrospection of Motivation Example

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```



❑ SymLoc: new memory modeling

❑ FastKLEE: efficient path exploration

❑ Vital: effective path exploration

Retrospection of Motivation Example

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```



❑ SymLoc: new memory modeling

❑ FastKLEE: efficient path exploration

❑ Vital: effective path exploration

❑ Cottontail: smart test input generation

Retrospection of Motivation Example

```
void buggy () {  
  
    // Vuln1: require complete memory modeling  
    void * from = malloc (100);  
    void * to = malloc (100);  
    if (to > from) { ... }  
    if (from > to) {  
        vulnerable_func1();  
    }  
  
    // Vuln2: require handling path explosion  
    ... // deeply nested  
    vulnerable_func2();  
    ...  
  
    // Vuln3: require structured test inputs  
    if (input = “{....}”) {  
        // application logic  
        vulnerable_func3();  
    } else {  
        earlier_termination();  
    }  
}
```



SymLoc: new memory modeling

FastKLEE: efficient path exploration

Vital: effective path exploration

Cottontail: smart test input generation

But not all vulnerabilities can be caught

Outline

□ Background

- Software vulnerabilities
- Symbolic execution

□ Motivation

- What limitations prevent current symbolic execution for vulnerability detection?
- Thesis statement and research objects

□ Methodology

- **SymLoc**: A new memory model for symbolic execution
- **FastKLEE and Vital**: Two new path exploration for symbolic execution
- **Cottontail**: A new structured input generation for symbolic execution

□ Future work

- Extension of current solutions
- Combine program analysis with GenAI

□ Conclusion and Acknowledgment

Future work

□ Extension of Current Solutions



- Extend SymLoc with more realistic memory modeling
- Extend Vital into greybox/hybrid as targets to guide directed fuzzing
- Extend Cottontail to support complex input formats (e.g., binaries)

Future work

□ Extension of Current Solutions



- Extend SymLoc with more realistic memory modeling
- Extend Vital into greybox/hybrid as targets to guide directed fuzzing
- Extend Cottontail to support complex input formats (e.g., binaries)

□ Program analysis + GenAI

Future work

□ Extension of Current Solutions



- Extend SymLoc with more realistic memory modeling
- Extend Vital into greybox/hybrid as targets to guide directed fuzzing
- Extend Cottontail to support complex input formats (e.g., binaries)

□ Program analysis + GenAI

- New testing engine that can be more general (e.g., across different languages)

Future work

□ Extension of Current Solutions



- Extend SymLoc with more realistic memory modeling
- Extend Vital into greybox/hybrid as targets to guide directed fuzzing
- Extend Cottontail to support complex input formats (e.g., binaries)

□ Program analysis + GenAI

- New testing engine that can be more general (e.g., across different languages)
- Combine LLM with formal verification techniques to guarantee the robustness

Future work

□ Extension of Current Solutions



- Extend SymLoc with more realistic memory modeling
- Extend Vital into greybox/hybrid as targets to guide directed fuzzing
- Extend Cottontail to support complex input formats (e.g., binaries)

□ Program analysis + GenAI

- New testing engine that can be more general (e.g., across different languages)
- Combine LLM with formal verification techniques to guarantee the robustness
- Practical hybrid selective symbolic execution for mixed programs
 - New software ecosystem (human written code + AI generated code)

Outline

□ Background

- Software vulnerabilities
- Symbolic execution

□ Motivation

- What limitations prevent current symbolic execution for vulnerability detection?
- Thesis statement and research objects

□ Methodology

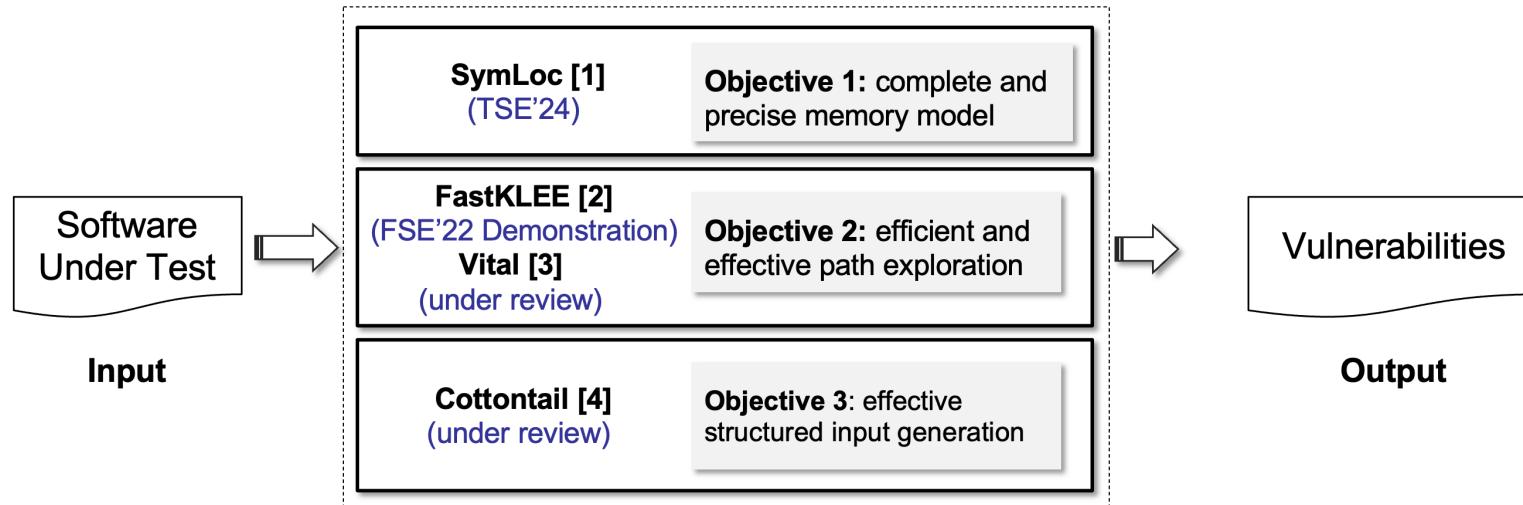
- **SymLoc**: A new memory model for symbolic execution
- **FastKLEE and Vital**: Two new path exploration for symbolic execution
- **Cottontail**: A new structured input generation for symbolic execution

□ Future work

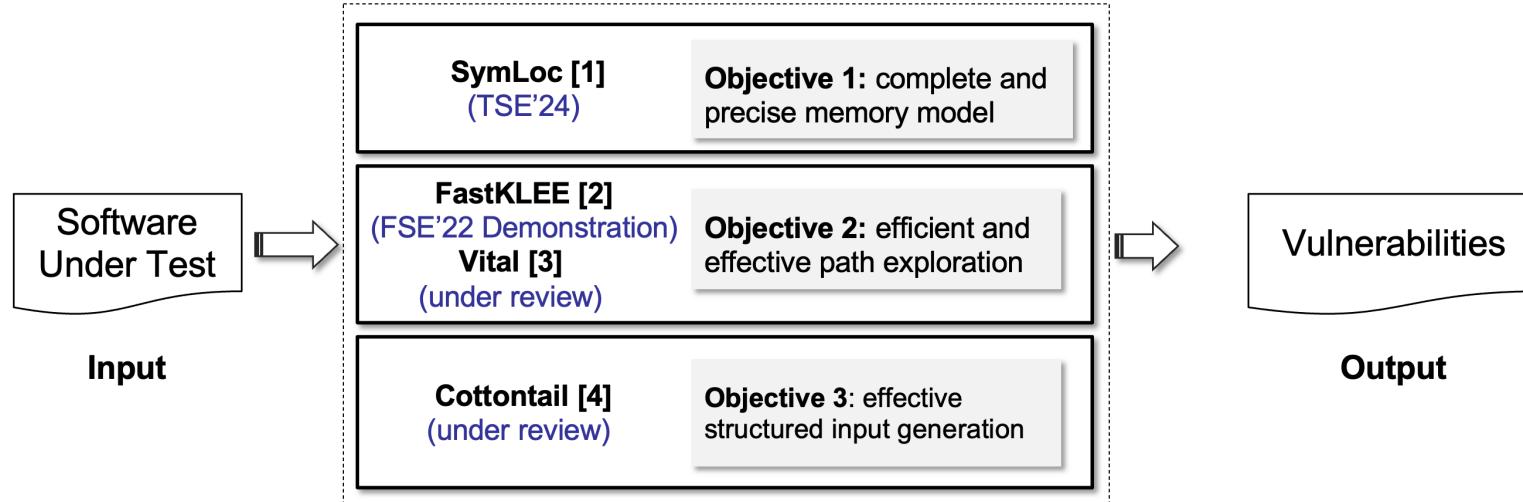
- Extend key ideas to hybrid fuzzing
- Combine Program Analysis with GenAI

□ Conclusion and Acknowledgment

Conclusion



Conclusion



- **Research impact**

- **Four** open-sourced tools to foster further symbolic execution research
- **10+** new memory related vulnerabilities are detected (with 7 new CVE assigned)
 - All of them have been fixed by developers so far

[1] Haixin Tu, Lingxiao Jiang, and et.al, "Concretely Mapped Symbolic Memory Locations for Memory Error Detection" (IEEE TSE).

[2] Haixin Tu, Lingxiao Jiang, and et.al., "FastKLEE: Faster Symbolic Execution via Redundant Bound Checking of Type-Safe Pointers", in Tool Demonstrations Track of FSE 2022.

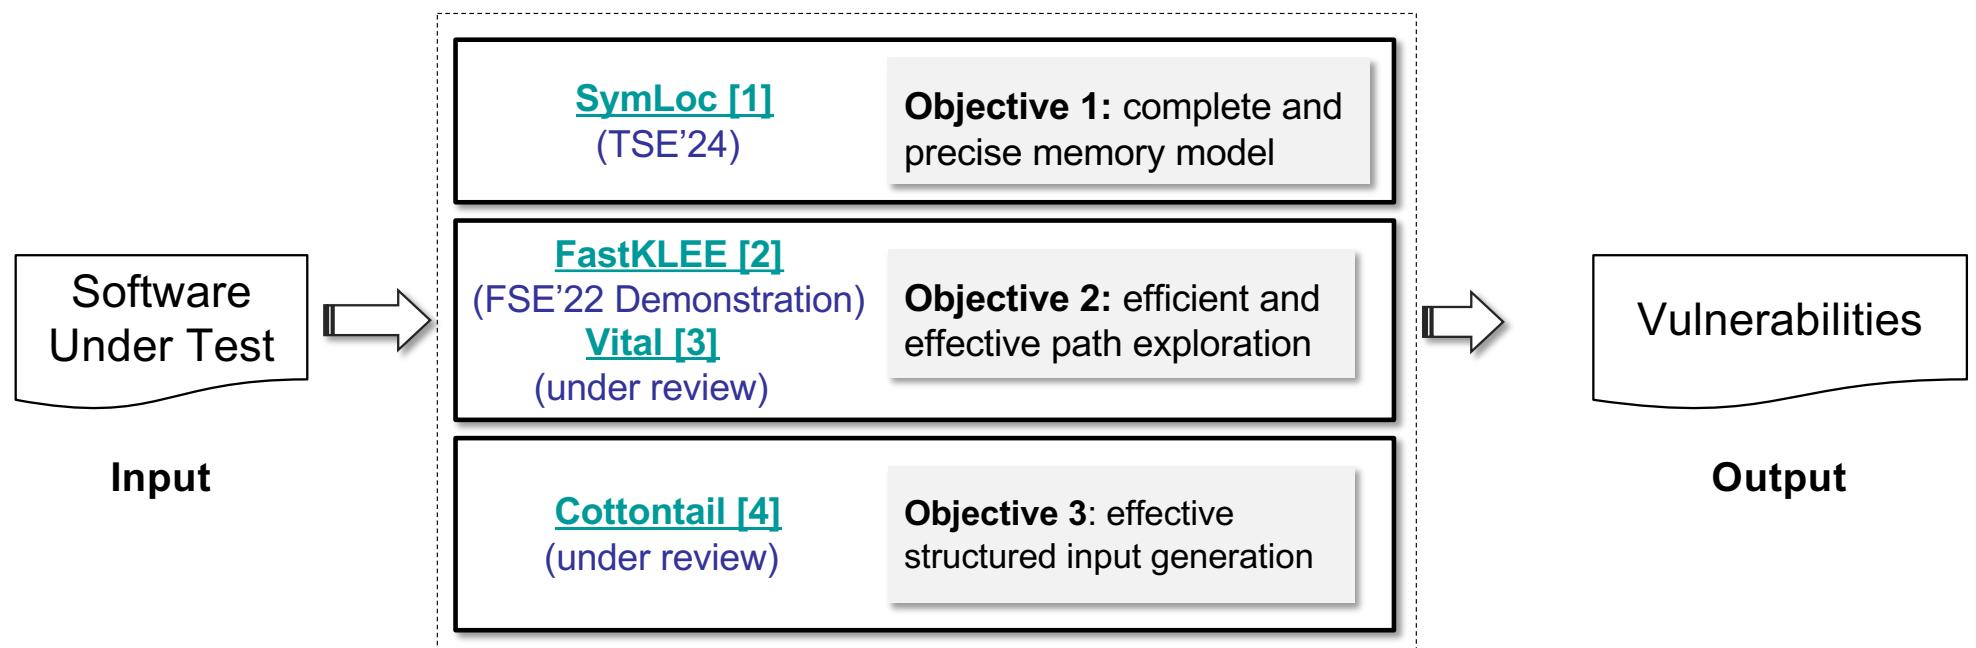
[3] Haixin Tu, Lingxiao Jiang and et.al., "Vital: Vulnerability-Oriented Symbolic Execution via Type-Unsafe Pointer-Guided Monte Carlo Tree Search." arXiv:2408.08772 (2024).

[4] Haixin Tu, Seongmin Lee, and et.al. "Large Language Model-Driven Concolic Execution for Highly Structured Test Input Generation." arXiv:2504.17542 (2025).

Acknowledgement

Thank you & Questions?

Boosting Symbolic Execution for Vulnerability Detection



Dissertation Defense by Haoxin Tu (May 9th, 2024)

References

- [1] Automated Test Generation: “**A Journey from Symbolic Execution to Smart Fuzzing and Beyond**” (Keynote by Koushik Sen)
- [2] Zhide Zhou, Zhilei Ren, Guojun Gao, He Jiang. “**An empirical study of optimization bugs in GCC and LLVM**”. JSS, 2021.
- [3] James C. King. 1976. **Symbolic execution and program testing**. Commun. ACM 19, 7 (July 1976), 385–394.
- [4] Roberto Baldoni, Emilio Coppa, Daniele Cono D’elia, Camil Demetrescu, and Irene Finocchi. 2018. “**A Survey of Symbolic Execution Techniques**”. ACM Computer Survey. 51, 3, Article 50 (July 2018), 39 pages.
- [5] Seo, Hyunmin, and Sunghun Kim. “**How we get there: a context-guided search strategy in concolic testing**.” Proceedings of the 22nd ACM SIGSOFT International Symposium on Foundations of Software Engineering. ACM, 2014.
- [6] Cristian Cadar, Daniel Dunbar, and Dawson Engler. 2008. “**KLEE: unassisted and automatic generation of high-coverage tests for complex systems programs**”. In Proceedings of the 8th USENIX conference on Operating systems design and implementation (OSDI’08). USENIX Association, USA, 209–224.
- [7] C. Cadar and K. Sen, “**Symbolic execution for software testing: three decades later**,” Commun. ACM, vol. 56, no. 2, pp. 82–90, 2013.
- [8] V. Chipounov, V. Kuznetsov, and G. Cadea, “**S2E: a platform for in-vivo multi-path analysis of software systems**,” in Proceedings of the sixteenth international conference on Architectural support for programming languages and operating systems, New York, NY, USA, Mar. 2011, pp. 265–278.
- [9] S. Poeplau and A. Francillon, “**SymQEMU: Compilation-based symbolic execution for binaries**,” presented at the in Proceedings of the 2021 Network and Distributed System Security Symposium, 2021.
- [10] Y. Shoshitaishvili et al., “**SOK: (State of) The Art of War: Offensive Techniques in Binary Analysis**,” in 2016 IEEE Symposium on Security and Privacy (SP), May 2016, pp. 138–157.
- [11] S. Poeplau and A. Francillon, “**Symbolic execution with SymCC: Don’t interpret, compile!**,” in 29th USENIX Security Symposium, 2020, pp. 181–198.
- [12] David Trabish, Timotej Kapus, Noam Rinetzky, and Cristian Cadar. 2020. “**Past-sensitive pointer analysis for symbolic execution**”. In Proceedings of the 28th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE 2020). New York, NY, USA, 197–208.